

QUANT REFERENCE SERIES

Delta One Desks and Securities Lending

A practitioner's reference

TOPIC OUTLINE

SEC-LENDING ECONOMICS • BORROW-RATE STACK
CROSS-BORDER DIVIDEND ARBITRAGE
PUT-STEAL & EARLY-EXERCISE INTEREST ARBITRAGE
SPX SYNTHETIC BOXES • CONVERTIBLE-BOND ARB
INDEX-REBALANCE FLOWS • ETF AP MECHANICS
HARD-TO-BORROW CASE STUDIES • TRS / ARCHEGOS
RISK MANAGEMENT • MICROSTRUCTURE & TAX FOOT-GUNS

14 chapters ■ 60,000+ words ■ 30+ worked examples

Companion to the Quant Course — theory first, execution second

Contents

0.1	Preface — What This Guide Is	5
0.2	Table of Contents	5
1	Chapter 1 — What a Delta One Desk Is	6
1.1	1.1 Definition and Origin	6
1.2	1.2 The Three Revenue Streams	6
1.3	1.3 A Worked Example: How a \$200m Index Swap Turns Into \$4.6m Revenue	7
1.4	1.4 The “Whose Balance Sheet” Question	8
1.5	1.5 An ASCII Map of a Modern Delta One Floor	8
1.6	1.6 The Cost of Capital Calculation	8
1.7	1.7 A Day on the Desk — The Human Texture	9
1.8	1.8 A Full Week of Trades — A Composite Example	11
1.9	1.9 References	13
2	Chapter 2 — Securities Lending Fundamentals	13
2.1	2.1 The Chain From Beneficial Owner to Short Seller	13
2.2	2.2 Cash vs Non-Cash Collateral	14
2.3	2.3 The 2008 AIG Lesson on Reinvestment Risk	14
2.4	2.4 General Collateral vs Specials	15
2.5	2.5 Recall Risk	15
2.6	2.6 The Loan Lifecycle: A Concrete Example	15
2.7	2.7 The Tri-Party Repo and Sec-Lending Plumbing	16
2.8	2.8 The Lending-Indemnification Question	16
2.9	2.9 Worked Example: The Lifecycle of a Single Loan in Detail	17
2.10	2.10 The Loan Board — Inside the Sec Lending Pod	18
2.11	2.11 The Mid-Sized Lender’s Calculus	20
2.12	2.12 References	21
3	Chapter 3 — The Borrow-Rate Stack	21
3.1	3.1 From Risk-Free Rate to Rebate	21
3.2	3.2 A Worked Numerical Waterfall	21
3.3	3.3 The “Specials Concentration” of Sec Lending Revenue	22
3.4	3.4 The Term Structure of Borrow	22
3.5	3.5 The “On Special” Threshold and Why It Matters	23
3.6	3.6 The Borrow-Rate Forecasting Problem	23
3.7	3.7 Implied Borrow From Options	24
3.8	3.8 Why Specials Exist — The Intuition	24
3.9	3.9 The Term Structure — Why It Slopes the Way It Does	25
3.10	3.10 Three Worked Examples Across Regimes	25
3.11	3.11 The Retail-Visible Surface	27
3.12	3.12 References	27
4	Chapter 4 — Dividend Trades	28
4.1	4.1 The Cross-Border Tax Wedge	28
4.2	4.2 The Yield-Enhancement Trade	28
4.3	4.3 Section 871(m) — The Regulatory Counter	28
4.4	4.4 Cum-Ex (Germany 2007–2012) — How Far Wrong It Can Go	29
4.5	4.5 Options-Based Dividend Trades — The American Call Early Exercise	29
4.6	4.5b The Retail-Visible Signature of Institutional Dividend Wedge	30
4.7	4.6 Worked Numerical Example: A Dividend Capture Trade	30
4.8	4.7 Dividend Forecasting and the Implied Dividend Curve	31
4.9	4.8 Special Dividends and Option Adjustments	31
4.10	4.9 The Dividend Future and the SX5E Dividend Index	32
4.11	4.10 The Dividend Forecast vs Realization Spread	32
4.12	4.11 Worked Example: The Pre-Ex-Date Term Repo	32
4.13	4.12 Dividend-Adjusted Volatility and the Implied Dividend Volatility Surface	33

4.14	4.12b The Day in the Life of the Dividend Desk	33
4.15	4.13 Cum-Cum vs Cum-Ex — The Two Faces of Dividend Stripping	34
4.16	4.14 Australian Franking Credits — A Separate Tax Architecture	35
4.17	4.15 EU Treaty Shopping and the ATAD Counter	37
4.18	4.16 The 2024 EU FASTER Directive	38
4.19	4.17 The Tax-vs-Structural Split — Retail Accessibility Matrix	38
4.20	4.18 Conversion Arbitrage Around Ex-Date	40
4.21	4.19 References	40
5	Chapter 5 — Put-Steal and Early-Exercise Interest Arbitrage	41
5.1	5.1 The Setup — American vs European Puts	41
5.2	5.1b Deriving NII From First Principles	42
5.3	5.1c A Brief History of the Trade	43
5.4	5.2 Why Retail Doesn’t Exercise	43
5.5	5.2b The Behavioral Anatomy of Non-Exercise	44
5.6	5.2c TastyTrade and Robinhood UX Traps — Specific Examples	44
5.7	5.3 The Short Stock Interest Arbitrage Game	45
5.8	5.4 Numerical Examples Across Rate Regimes	46
5.9	5.4b Side-by-Side Trade Variants — Put vs Call, Single-Leg vs Spread	47
5.10	5.5 The Repo Spec Connection	47
5.11	5.5b Why the Wedge Doesn’t Get Arbitraged Away	48
5.12	5.6 The Bull Put Spread Implementation	48
5.13	5.7 When the Trade Breaks	49
5.14	5.8 The Mathematics of the Early Exercise Boundary	49
5.15	5.9 Multi-Period NII Analysis	49
5.16	5.10 Cross-Strike Comparison Table	50
5.17	5.11 The “Pin and Crush” Variant	50
5.18	5.12 Failure Modes — Detailed Analysis	50
5.19	5.13 Actionable Rules of Thumb for Retail	51
5.20	5.14 Walking Through <code>strategies/put_steal.py</code>	52
5.21	5.15 References	54
5.22	5.16 A Closing Note on Chapter 5	54
6	Chapter 6 — SPX Boxes and Synthetic Lending	54
6.1	6.1 The Four-Leg Structure	55
6.2	6.2 SPX vs SPY — The Critical Distinction	55
6.3	6.3 The Box Loan: Synthetic Borrowing	55
6.4	6.4 Margin Treatment of Boxes	56
6.5	6.5 The 2018 Robinhood “Box” Incident	57
6.6	6.6 Pin Risk and Box Mechanics Near Expiry	57
6.7	6.7 Rolling Boxes	57
6.8	6.8 Why Retail Brokers Treat Boxes Differently	58
6.9	6.9 The Tax Treatment of Boxes — A Closer Look	58
6.10	6.10 Why the Box Implied Rate Is Below SOFR	58
6.11	6.11 Box Liquidity and Bid/Ask	59
6.12	6.12 The “Stress” Box and What Happens in March 2020	59
6.13	6.13 The Retail Box-As-Loan in Practice — A Detailed Walkthrough	60
6.14	6.14 The Risks of Box Lending vs Box Borrowing	61
6.15	6.15 A 2-Year Simulation With Bid-Ask Drag	62
6.16	6.16 The Failure Case — When Boxes Go Wrong	64
6.17	6.17 References	65
7	Chapter 7 — Convertible Bond Arbitrage	65
7.1	7.1 The Convertible Bond Structure	65
7.2	7.2 The Classical Convert Arb Trade	65
7.3	7.3 Where Securities Lending Enters	65
7.4	7.4 The Death Spiral Convertible	66

7.5	7.5 Worked Example	66
7.6	7.6 The Convertible Issuance Cycle	66
7.7	7.7 The Death-Spiral Convertible Worked Example	67
7.8	7.8 Modern Convertible Markets	68
7.9	7.9 The Historical Arc of Convert Arb — 1990s to Present	68
7.10	7.10 References	69
8	Chapter 8 — Index-Rebalance and Corporate-Action Microstructure	70
8.1	8.1 The Index Rebalance Trade	70
8.2	8.2 The Russell Reconstitution	70
8.3	8.3 Spin-Offs and When-Issued Markets	70
8.4	8.4 M&A Risk Arb From the Desk’s Perspective	70
8.5	8.5 Tender-Offer Arb	71
8.6	8.6 Corporate Action Worked Example: A Spin-Off	71
8.7	8.7 The S&P 500 Add Trade — Detailed Mechanics	71
8.8	8.8 The “Index Fund Front-Running” Question	72
8.9	8.9 Spin-Offs — The When-Issued Market in Detail	72
8.10	8.10 The Tesla S&P 500 Inclusion (December 2020)	73
8.11	8.11 The AT&T / Discovery Spin-Off (April 2022)	74
8.12	8.12 An M&A Risk-Arb Walkthrough — Microsoft / Activision Blizzard (2022-2023)	74
8.13	8.13 References	75
9	Chapter 9 — ETF Creation/Redemption Arbitrage	75
9.1	9.1 The Creation/Redemption Mechanism	75
9.2	9.2 Cash vs Basket Creates	76
9.3	9.3 Worked Example: VOO Creation	76
9.4	9.4 Stress Period Dislocations: March 2020	76
9.5	9.5 The “Cash Drag” and ETF Tracking Error	76
9.6	9.6 The “Heartbeat” Trade — Tax Efficiency in ETFs	76
9.7	9.7 ETF Share Class Structures and Vanguard’s Patent	77
9.8	9.8 The “Hidden Inventory” of Big-Basket ETFs	77
9.9	9.9 The 4:00pm \$50m ETF Redemption — A Detailed Walkthrough	77
9.10	9.10 References	78
10	Chapter 10 — Hard-to-Borrow and Squeezes From the Desk’s Perspective	79
10.1	10.1 The Other Side of the Squeeze	79
10.2	10.2 GME 2021 Case Study From the Prime Broker Side	79
10.3	10.3 The “Locate Theater”	79
10.4	10.4 The Buy-In Risk	80
10.5	10.5 Synthetic Short via Options — The Workaround	80
10.6	10.6 The “Locate Pool” and Why It’s Not Infinite	80
10.7	10.7 The “Threshold Securities List”	81
10.8	10.8 GME 2021 — A Full Prime Broker Narrative	81
10.9	10.9 References	82
11	Chapter 11 — Total-Return Swaps, Equity Swaps, and the Archegos Case	82
11.1	11.1 The Total-Return Swap Mechanism	82
11.2	11.2 Section 871(m) Detail	83
11.3	11.3 The Archegos Capital Case Study	83
11.4	11.4 What Banks Should Have Done	83
11.5	11.5 Post-Archegos Industry Changes	84
11.6	11.6 Detailed Anatomy of an Equity Swap — A Sample Trade	84
11.7	11.7 The Variance-Swap Cousin	85
11.8	11.8 Section 871(m) Compliance Detail	85
11.9	11.9 The Archegos Blow-Up — Week-By-Week	85
11.10	11.10 References	87
12	Chapter 11.5 — Historical Case Studies of Delta-Desk and Sec-Lending Episodes	88

12.1	11.5.1	The Volkswagen Squeeze (October 2008)	88
12.2	11.5.2	The Long-Term Capital Management Equity Hedging (1998)	89
12.3	11.5.3	The Bear Stearns Hedge Funds Collapse (June 2007)	89
12.4	11.5.4	The Flash Crash (May 6, 2010)	90
12.5	11.5.5	The 2014 Russian Stock Collapse and the Sberbank Borrow Spike	90
12.6	11.5.6	The 2024 Yen Carry Trade Unwind (August 2024)	90
12.7	11.5.7	The Bear Stearns Hedge Funds Collapse — The Auditor Narrative (June 2007)	91
13		Chapter 12 — Risk Management on a Delta One Desk	92
13.1	12.1	The Core Risks	92
13.2	12.2	The Central Risk Book	92
13.3	12.3	Stress Testing	92
13.4	12.4	Margin Calls and the 24-Hour Cycle	92
13.5	12.5	The CSA Hierarchy	93
13.6	12.6	The Day-of-Failure Playbook	93
13.7	12.7	The CCAR / DFAST Stress Tests	94
13.8	12.8	The Internal Risk Limits	94
13.9	12.9	Operational Risk — The Quiet Killer	94
13.10	12.10	The Day in the Life of the Central Risk Book	94
13.11	12.11	References	96
14		Chapter 13 — Microstructure, Tax, and the Retail Lending Program	96
14.1	13.1	W-9 vs W-8BEN — The Retail Tax Wall	96
14.2	13.2	Qualified vs Ordinary Dividends	96
14.3	13.3	IRS Section 1058	96
14.4	13.4	Per-Share Lending Fee Math	97
14.5	13.5	The Robinhood Stock Lending Program	97
14.6	13.6	Schwab and Fidelity Programs	97
14.7	13.7	The Wash-Sale Rule and Sec Lending	97
14.8	13.8	The 60-Day Holding Period for Qualified Dividends	98
14.9	13.9	The Qualified Dividend and Substitute-Payment Distinction	98
14.10	13.10	The Substitute Dividend Wedge for Foreign Holders	98
14.11	13.11	The Retail Brokerage Lending Program Comparison Matrix	98
14.12	13.12	The Voting-Rights Tradeoff	99
14.13	13.13	Retail Tax Foot-Guns — Specific Examples	99
14.14	13.14	References	100
15		Chapter 14 — Connecting Back to This Repo	100
15.1	14.1	Strategy Status Summary	100
15.2	14.2	Put Steal — Confirmed Math	101
15.3	14.3	Dividend Arb — Why Inactive	101
15.4	14.4	Volatility Arbitrage and Sec-Lending Dislocation	102
15.5	14.5	What Is NOT Retail-Executable	102
15.6	14.6	What IS Retail-Executable	102
15.7	14.6b	A Repository Tour for the Curious Reader	102
15.8	14.7	The Big Picture for the Retail Quant	103
15.9		Appendix A — Glossary	103
15.10		Appendix B — Symbols and Notation	104
15.11		Appendix C — Worked Examples Compendium	104
15.12		Appendix D — Quick-Reference Formulas	107
15.13		Appendix E — Common Trade Comparison Table	108
15.14		Appendix F — A Note on Sources	108
15.15	14.8	A Frequently Asked Question — “Why Do The Strategies in This Repo Look Different From Institutional Versions?”	108
15.16	14.9	What’s Missing — Future Work	109
15.17	14.10	The Final Word — How to Read Industry News After This Guide	109
15.18	A	Final Mental Model — The Retail Quant’s Synthesis	110

0.0.1 How Investment Banks Print Money From Funding, Borrow, Dividends, and Boxes — and What of It Reaches Retail

0.1 Preface — What This Guide Is

This is a flagship reference. The reader who finishes it should be able to:

1. Read a delta-one term sheet without confusion.
2. Trace a single share of AAPL from a CalPERS pension trust through a Goldman Sachs prime broker through a Citadel short-sale through to its return on recall — and price every fee charged at each hop.
3. Construct a four-leg SPX box on paper, value it under put-call parity, and explain why the funding rate it implies is a function of options-market supply and demand for collateral.
4. Compute the Net Interest Income on a deep-ITM American put and back out the dollar value the desk earns when retail fails to exercise.
5. Quote, with order-of-magnitude correctness, what the desk makes on a corporate hard-to-borrow at a 22% fee versus a \$50bn AAPL inventory at 12 bps general collateral — and explain why both businesses coexist on the same trading floor.

It is unusual material because most retail traders never see inside the financing book. The retail world is “buy a stock or buy an option.” The institutional world is “buy a stock, lend it overnight at 25 bps, finance the inventory at SOFR + 8 bps, hedge the dividend exposure with a swap, and roll the package quarterly.” The economics of the latter dwarf the former. Understanding it makes the retail trader vastly more capable — not because all of these trades are retail-executable (most are not), but because the prices retail sees are *set* by these flows.

A second reason this material matters: a small handful of these trades — the put-steal arb in particular, certain conversion/reversal mispricings, and SPX box financing — *are* retail-executable and have been executed at scale by sophisticated retail since 2019. Where a retail edge exists, this guide identifies it explicitly.

The guide spans fourteen chapters. Read in order; each builds on its predecessor. Citations to academic literature are inline.

0.2 Table of Contents

1. What a Delta One Desk Is
 2. Securities Lending Fundamentals
 3. The Borrow-Rate Stack
 4. Dividend Trades
 5. Put-Steal and Early-Exercise Interest Arbitrage
 6. SPX Boxes and Synthetic Lending
 7. Convertible Bond Arbitrage
 8. Index-Rebalance and Corporate-Action Microstructure
 9. ETF Creation/Redemption Arbitrage
 10. Hard-to-Borrow and the Squeeze From the Desk’s Perspective
 11. Total-Return Swaps, Equity Swaps, and the Archegos Case
 12. Risk Management on a Delta One Desk
 13. Microstructure, Tax, and the Retail Lending Program
 14. Connecting Back to This Repo
-

1 Chapter 1 — What a Delta One Desk Is

1.1 1.1 Definition and Origin

A Delta One desk is a sell-side trading group that warehouses, hedges, and finances products whose payoff is approximately linear in the underlying price — *delta of one*, hence the name. The label is a contrast with the equity-derivatives desk next to it, which warehouses *non-linear* exposures (gamma, vega, vanna, vomma).

The product list is short on paper, sprawling in practice:

Product	What it is	Who buys it
Index futures (ES, NQ, FESX)	Exchange-traded delta-one contracts	Macro funds, hedgers, prop firms
ETFs (SPY, QQQ, EFA, EEM)	Listed funds tracking baskets	Retail, asset allocators, model funds
Equity total-return swaps	Bilateral swap on a stock or basket	Hedge funds wanting leverage / discretion
Custom thematic baskets	Bespoke collections (AI names, EV supply chain, etc.)	Macro funds, family offices
Quanto and composite swaps	Swap on a foreign index in non-foreign currency	Cross-border allocators
Single-stock futures (where legal)	Exchange-traded forward contracts	Tax/arb shops
Certificates and notes	OTC structured wrappers replicating exposure	European retail, private banks
ADRs and GDRs	Bank-issued depositary receipts	Cross-border retail and institutions

The desk does not exist to “be long” or “be short” any of these products in a directional sense. It exists to *intermediate*. A pension fund wants exposure to the MSCI Emerging Markets Index. A bank’s delta-one desk sells the pension a total-return swap referencing MSCI EM. The desk now has a short delta exposure to MSCI EM. To hedge, it buys the underlying basket of EM stocks. The desk now has a flat economic position — but it has *picked up the financing leg* and the *securities-lending revenue* on every name in that basket. The pension paid the desk a financing spread; the desk also lends each EM name to short sellers at whatever the borrow market will bear. The economic substance: the desk is renting its balance sheet and its securities-lending franchise, *not* taking market risk.

This is the “fund-the-trade” mental model. Every delta-one product is a packaging of three economic services:

Delta One product price = Spot ± Hedge cost ± Financing ± Borrow value ± Tax wedge ± Margin

The desk’s gross margin is the sum of frictions it can extract at each step. The desk’s risk is operational and counterparty, not directional.

1.2 1.2 The Three Revenue Streams

A canonical delta-one desk runs three intertwined P&L streams:

1. **Financing.** The desk advances cash to a client (long swap, long basket via repo, etc.) and is repaid SOFR + spread. On a \$1bn client book at SOFR + 35 bps, financing P&L runs at roughly \$3.5m/year before funding cost.
2. **Securities lending.** The desk’s inventory is loaned out to short sellers via the prime brokerage matched-book. On a \$50bn long inventory, even a 12 bps blended rate generates \$60m/year. On the special names — the 5–10% of inventory at >100 bps — the same desk can earn \$80m/year more. This is the largest stream by absolute dollars at a top-three prime broker.
3. **Execution and rebalancing.** Index swaps reset; ETF baskets are created and redeemed; quanto hedges are rolled. Each touch is a chance to capture bid/ask. On a top-tier desk, this is a steadier \$30–50m/year.

A typical mid-size delta-one desk at a Tier 1 bank generates \$200–400m gross revenue. At Goldman, JPM, Morgan Stanley pre-Archehos, the number was closer to \$700m–\$1bn annually for global delta-one.

1.3 1.3 A Worked Example: How a \$200m Index Swap Turns Into \$4.6m Revenue

Consider a hedge fund “Silvermark” that wants to express a tactical view: long the Russell 2000 for six months. Silvermark could buy IWM ETF — but ETF tracking error and rebalancing costs annoy them, and they want leverage with daily reset.

The trade:

- **Notional:** \$200m on the Russell 2000 index
- **Tenor:** 6 months
- **Reset:** Daily, financing accrues on actual notional
- **Funding leg:** Silvermark pays SOFR + 60 bps to bank
- **Performance leg:** Bank pays Silvermark the index total return
- **Initial margin:** 10% (\$20m) posted in T-bills

The bank’s delta-one desk now has the *economic short* of the Russell 2000 on this swap. To hedge, the desk buys the Russell 2000 basket — 2,000 names, weighted appropriately. Suppose the desk uses a 200-name optimized sample tracking the index to within 5 bps tracking error.

The desk’s books:

Asset side: Long \$200m basket (200 names sampling Russell 2000)
Liability: Short \$200m via swap with Silvermark
Net delta: ~0 (perfect to within tracking error)

Now the revenue streams:

Financing income. Silvermark pays SOFR + 60 bps on \$200m. SOFR averages 4.50% over the period. The desk receives $\$200\text{m} \times 5.10\% \times 0.5 = \5.10m for six months.

Funding cost. The desk’s treasury charges an internal funding rate of SOFR + 25 bps to fund the long basket position. Cost: $\$200\text{m} \times 4.75\% \times 0.5 = \4.75m .

Net financing P&L: $\$5.10\text{m} - \$4.75\text{m} = \$0.35\text{m}$

Securities-lending revenue. The desk’s matched-book lends out the 200-name basket. The blended SLB fee on Russell 2000 names is roughly 35 bps (some general collateral at 8 bps, some specials at 200+ bps). Revenue: $\$200\text{m} \times 0.35\% \times 0.5 = \0.35m . Of this the agent split keeps 70% to the desk: $\$0.245\text{m}$.

Dividend tax wedge. The basket pays approximately 1.4% in dividends over 6 months: $\$200\text{m} \times 1.4\% \times 0.5 = \1.40m gross. The swap reset pays the *full* gross dividend back to Silvermark. But the desk receives only the qualified-dividend portion if it lends the underlying out across the record date — it captures the *substitute payment* difference. On a \$1.40m dividend stream, this net wedge is approximately $30 \text{ bps} \times \$200\text{m} \times 0.5 = \0.30m .

Reset and rebalancing. Russell rebalances annually and quarterly, each name has a corporate-action drift, and Silvermark may add or reduce notional. Bid/ask captures over six months: roughly \$0.20m.

Total revenue from this single trade:

Financing net:	\$0.35m
Sec-lending share:	\$0.245m
Dividend wedge:	\$0.30m
Execution friction:	\$0.20m
Total:	\$1.10m on \$200m notional, six months
	110 bps annualized "all-in"

The bank reports this to Silvermark as “we charged you 60 bps over SOFR.” The actual all-in margin to the bank is roughly 110 bps. The 50 bps differential is the value of the desk’s franchise — its ability to monetize the trade across financing, lending, and dividend channels simultaneously.

This is the economic engine of the entire industry. Multiply by 30–50 active swap clients per Tier 1 bank, plus the ETF AP business and the dividend trades that follow, and the total franchise revenue runs into the billions.

- Operational cost (technology, ops, legal)
- Allocation of bank overhead

Cost of regulatory capital example for a \$1bn long basket:

Equity RWA at 100% × \$1bn = \$1bn RWA

Required Tier 1 capital × 8% = \$80m

Cost of equity (post-tax) at 11% = \$8.8m/year

Translating: each dollar of basket inventory costs the bank roughly 88 bps/year in regulatory capital alone, on top of funding cost. A trade earning SOFR + 30 bps is therefore *negative-NPV* if it consumes a full dollar of RWA per dollar of notional, even though the P&L line shows positive.

Modern desks therefore prioritize:

1. Trades that net within the CRB (low marginal RWA).
2. Trades hedged with cash, not derivatives (lower RWA than swaps, which carry counterparty RWA).
3. Trades on collateral-eligible names (lower LCR consumption).
4. Trades that monetize *non-balance-sheet* economics (sec lending, dividend wedge, execution skill).

The capital constraint is what shifted the industry from “balance-sheet-as-product” (pre-2008) to “balance-sheet-as-friction” (post-Basel III). It is also why the lending and dividend franchises — both of which earn revenue with relatively little balance sheet — became so disproportionately important.

1.7 1.7 A Day on the Desk — The Human Texture

The previous sections describe what a delta-one desk *does*. The following sketches what it *feels like* to sit on one — the sights, sounds, hand-offs, and rituals that make up a working day. This matters because the people on the desk are not abstract optimizers; they are humans operating in a particular built environment, and the trades they execute are shaped as much by ergonomic geography as by mathematics.

Consider a Tuesday in mid-September on the Goldman Sachs Equities Floor at 200 West Street, lower Manhattan. The trading floor occupies the entire fourth floor, an open-plan rectangle the size of two football fields with low ceilings and a single uninterrupted band of windows facing the Hudson. The lighting is fluorescent and slightly green. The air-conditioning is set aggressively — the rule of thumb is that the floor should run at 19°C/66°F because the screens give off heat and traders complain when they can feel it.

Pre-market begins at 06:15 ET. The first arrivals are the European-shift coverage — traders who came in to manage the overnight Asian session and stayed through London open. At 06:30, the morning meeting starts in the glass-walled conference room at the floor’s east end. The agenda runs to a precise rhythm: macro overnight (the macro economist on the line from London takes three minutes), Asia recap (Hong Kong-based equity strategist takes two minutes), what the desk is sitting on (the desk head, two minutes), what’s expected from clients today (the head of equity sales, three minutes), risk flags (the risk manager, one minute, sometimes longer if there’s an issue). Total: fifteen minutes. The meeting is held standing — this was Hank Paulson’s rule from his time as CEO, and it has stuck. Anyone who wants a chair stays in their seat and listens by speaker.

By 07:00 the desk is at full strength: roughly 40 traders, structurers, and operations staff for the global delta-one franchise. They are clustered by product. Index swaps to the north end (closest to the cash desk for hedge sourcing). ETF AP and creates in the middle. Single-stock futures and quanto on the south side. Sec lending matched-book — six traders — sits at the sec lending pod, slightly elevated on a raised platform so the senior trader can see the whole floor. Treasury and funding sits adjacent to sec lending; the two functions report into a single “Financing” co-head structure that has become standard at the major banks since 2015.

Each trader has between four and eight monitors. The standard arrangement: the leftmost two screens display the order management system (Fidessa, REDI, or the bank’s proprietary equivalent) and the desk’s risk dashboard. The middle two show market data — Bloomberg in the upper left, the bank’s internal pricing engine in the lower right. The rightmost monitors carry chat (Bloomberg messaging on top, internal Symphony or similar at the bottom), email, and a fourth ad-hoc workspace for ad-hoc analysis. The senior traders sit nearest the aisle so they can stand and talk to the desk head without leaning over.

[Stylized seating - Goldman Equities Delta One pod]

Aisle Aisle Aisle Aisle Aisle

```

      |           |           |           |           |
+---+ +---+ +---+ +---+ +---+
| A |---| A |---| B |---| B |---| C | <-- Junior traders
+---+ +---+ +---+ +---+ +---+ "screens facing aisle"
| A |---| A |---| B |---| B |---| C | <-- Mid-level
+---+ +---+ +---+ +---+ +---+
| A |---| A |---| B |---| B |---| C |
+---+ +---+ +---+ +---+ +---+

                                <-- Sr trader (raised seat)
                                <-- Head of desk (standing)
                                <-- Risk manager (one row back)

```

A = Index swaps team B = Custom baskets C = ETF AP

The hierarchy is precise. A junior trader (years 1-3) executes orders, monitors P&L, and handles pre-trade checks. A mid-level trader (3-7 years) prices client trades, takes inbound calls from sales, and warehouses small risk. A senior trader (7+ years) owns a relationship book, prices the largest tickets, and supervises the juniors. The desk head (typically MD-level, 10-15+ years) takes ultimate responsibility, interfaces with risk and senior management, and steps in only on the largest or most contentious trades. Above the desk head is the global head of equity financing or equity derivatives, who reports to the head of the equities division, who reports to the CEO. Five layers between a junior trader and the firm CEO.

09:30 ET. The opening bell. The cash equities desk to the south erupts into the controlled chaos of the morning auction. The delta-one desk is quieter — most of its trades will price after 10:00 ET when the morning volatility calms. The first hour is mostly inbound calls. A hedge fund client on speaker says: “We want \$300m on the FTSE 100, six-month tenor, can you indicate?” The trader handling that account (typically a mid-level) puts the call on hold, looks at his pricing engine, runs a check with the sec-lending pod (“how is FTSE 100 GC borrow today?”), runs a check with treasury (“what’s our internal funding for £-denominated?”), and quotes back: “SOFR + 32 bps, \$300m, six-month, 12% IM.” The client says “show me 35,” the trader says “I can do 33,” the client says “done at 33.” The trade is booked. Fifteen seconds, \$300m of notional, \$1m of expected revenue. This happens dozens of times a day.

Through the morning, the Russell index rebalance team is monitoring announcements (Russell publishes preliminary rebal lists each May, with execution at the June close). Today is a quiet rebalance day, but the team has roughly \$400m in pre-positioned long inventory and \$200m short across various Russell 1000 / 2000 boundary names. The P&L on these positions ticks daily — small losses today on a couple, small gains on others, with the major payoff coming on Recon Friday in late June.

12:00 ET. Lunch is desk-side. The standard sandwich and cookies tray rolls in from the bank’s catering operation. Most traders eat at their desks, one hand on the mouse. The senior traders sometimes go to the cafeteria upstairs for ten minutes, mostly to talk with peers from other desks. The tradition of long client lunches at expensive restaurants — once a hallmark of bulge-bracket Wall Street — has all but vanished post-Volcker. Compliance restrictions on entertainment, plus the relentless screen-bound nature of the modern workflow, mean that most client interaction is now over chat and the occasional voice call.

14:00 ET. The afternoon’s main event for the financing pod is the daily Tri-Party Recon meeting. Run jointly with the bank’s operations team and a tri-party agent (BNY or JPM), this 30-minute call confirms the day’s collateral movements, identifies any breaks (where the bank’s books and the agent’s books disagree), and queues up tomorrow’s expected cash and securities movements. Most days nothing dramatic happens — a few small breaks are reconciled, settlement times confirmed. On stressful days (March 2020, August 2024), the meeting can run two hours and involve real-time decisions about which collateral to substitute.

15:30 ET. The closing-auction window. ETF AP teams are at their busiest — index funds rebalance into the close, and ETFs that are off NAV by even a basis point are an arbitrage opportunity. The AP team’s screens are scrolling rapidly with creation/redemption indications. A junior on the AP team says into a chat: “VOO trading 2bps over NAV, want to create?” The senior nods. They put together a creation order: 50,000 VOO shares, basket sourced from internal inventory plus tactical buying in the underlying names. The trade ticks at 16:00. Net P&L: \$1,800. The junior’s hands shake slightly — it’s only his second AP trade.

16:00 ET. Close. The desk doesn’t actually go quiet — there is always end-of-day reconciliation, position keeping, and

prep for the Asian session. But the intensity drops. Senior traders begin their post-close routine: review the day's P&L, check tomorrow's flow expectations, write a brief end-of-day note to the desk head. Mid-levels review their trades for any operational issues. Juniors triple-check that all morning trades are properly booked and reconciled.

The P&L is reviewed at three levels. First, the *trader-level* P&L: each trader's individual book, marked to the close. Second, the *desk-level* P&L: the aggregate across all traders, after intra-desk netting. Third, the *enterprise-level* P&L flowing into the firm's books and ultimately the income statement. Each level has a different audience. Trader P&L drives bonuses; desk P&L drives the desk head's compensation and senior management's perception of the franchise; enterprise P&L is what shareholders see.

Compensation on a delta-one desk follows a typical bulge-bracket structure. The base salary for a junior trader is around \$175,000-\$225,000. Total comp at year one (base + bonus) typically lands at \$250,000-\$350,000 if the desk has a normal year. Mid-level traders make \$400,000-\$800,000. Senior traders pull \$1m-\$3m. The desk head's compensation, at a top-three franchise, can be \$5m-\$10m+ in a strong year, with the global head of equity financing in the \$10m-\$25m range. Compensation is heavily back-end loaded with deferred stock, ostensibly to align long-term incentives but also because the bank wants to keep the trader from defecting to a hedge fund. A typical comp package might be 50% cash, 50% deferred stock vesting over three years, with a clawback clause for trades that go bad after the bonus is paid.

The clawback is theoretical but not zero. The most famous invocation: in 2014, multiple Goldman traders had bonuses clawed back after a trader on the structured products desk was implicated in a mortgage-related settlement. On the delta-one desk, where individual trades rarely produce huge wins or losses, clawback is essentially never invoked. The bigger risk to comp is *desk-level* underperformance — if delta-one revenue falls 30% year-over-year due to industry compression, the entire desk's bonus pool shrinks 30%, regardless of any individual trader's contribution.

The reporting lines and inter-desk coordination matter as much as the trading itself. The delta-one desk talks constantly to:

- **Equity research** for fundamental views on names the desk holds in inventory. A research downgrade on a name with \$500m of desk inventory is a coordination event — the desk may need to reduce inventory or hedge.
- **Equity sales-trading** for client flow context. A salesperson saying “Big Hedge Fund X is asking about Russell shorts” tells the desk to expect borrow demand and start sourcing locates.
- **The cash equities desk** for execution. The delta-one desk doesn't typically execute its own cash trades; it routes them to the cash desk's execution algorithms.
- **The equity derivatives desk** for vol-related products. A hedge on a long basket via SPX puts requires the derivatives desk to provide the puts, which the delta-one desk pays for from its P&L.
- **Risk and central risk book** for aggregate exposure. The CRB views all desks' positions and routes hedges optimally.
- **Legal and compliance** for ISDA negotiations, 871(m) compliance, and stress-test responses.
- **Treasury** for funding and balance sheet usage.
- **Operations** for trade settlement, custody, and reconciliation.

In a typical week, a single trader will probably interact with all of the above at least once. The trader's job is, in some sense, to be the *node* through which client flow becomes a fully-integrated firmwide trade. The technical execution is mostly handled by systems; the human is needed for judgment, relationship management, and the constant navigation of edge cases that the systems do not handle gracefully.

Lubin (2019) describes the modern derivatives trading floor as a “low-emotion, high-information environment” — the affect is almost office-like, despite the dollar amounts flowing through. Foreshadowing this guide's later chapters: the *crisis* moments (Archegos in March 2021, August 2024 yen unwind, March 2020 COVID) involve a different floor entirely — voices raised, executives walking through the rows, weekend conference calls. But on a normal Tuesday in September, the floor's sound is mostly the clicks of mouse buttons and the occasional voice into a Bloomberg-side handset. It looks more like a quiet open-plan office than a Hollywood depiction of Wall Street.

1.8 1.8 A Full Week of Trades — A Composite Example

To make the week-of-a-desk concrete, here is a composite of the kinds of trades a Tier 1 desk would execute over a normal week in late 2024 / early 2025. Names and sizes are realistic but the specific clients are fictionalized.

Monday.

- 09:45 ET: Inbound call from a London-based macro fund. They want \$400m short MSCI Emerging Markets via a 6-month total return swap. The desk quotes SOFR + 75 bps (the EM borrow blended cost is structurally higher than DM). Trade done at SOFR + 73 bps after the client pushes back. Desk hedges by buying \$400m of an EM index basket through the cash desk.
- 11:20 ET: A US pension fund (via its agent lender) lets the desk know that 4m shares of TSLA they had been lending will be recalled effective T+1. The desk's prime brokerage clients are short TSLA against this borrow; the recall must be passed through. The desk's PB team contacts each of the affected hedge fund clients and arranges either re-borrow at the new market rate (TSLA borrow has spiked from 30 bps to 80 bps over two weeks due to a convert issuance) or forced cover.
- 14:30 ET: A new ETF launch. The desk's ETF AP team participates in seeding a new \$2bn thematic ETF on AI-related semiconductors. The seed is cash-create — desk delivers \$2bn of cash, receives \$2bn of fresh ETF shares, then arbitrages the shares against the underlying basket as the ETF begins trading.
- 16:10 ET: End-of-day P&L review. The desk's day P&L is +\$3.2m, which is normal-good. The big contributor: the EM swap from the morning, which already shows a small mark-to-market gain because EM has rallied 30 bps since the trade was struck.

Tuesday.

- 09:00 ET: The pre-open meeting flags an IPO going effective Friday. The desk's structured-products team will market-make in the new shares; the AP team will structure an ETF inclusion if the new shares qualify; the sec lending team will need to source borrow as soon as the lockup expires (180 days out). All three teams coordinate on inventory sourcing.
- 10:45 ET: A sovereign wealth fund client requests a custom basket: 50 names representing the European energy transition. The structurer takes the request, builds the basket spec (with weights, rebalancing rules, currency hedging), prices it (financing rate, basket fee, sec lending share), and quotes back. The fund accepts at SOFR + 110 bps for a 3-year structured note referencing the basket. Notional: \$750m. Expected revenue: \$4-6m over three years.
- 13:00 ET: A small-cap biotech announces a positive Phase 2 result. The desk's HTB book on this name jumps from 12% to 35% borrow rate within an hour. The desk's hedge fund clients shorting this name are facing a 23-percentage-point hike in carry cost; several call to recall their shorts.
- 15:30 ET: The Russell rebal team adjusts a small position based on an unexpected M&A announcement (one of the names on the upgrade list is being acquired, which changes the rebal calculus).

Wednesday.

- 09:15 ET: A weekly tri-party recon issue: \$40m of equities posted as collateral against a swap have been reclassified as "non-eligible" by the agent due to a corporate action (a pending spin-off makes the shares temporarily ineligible). The desk substitutes other collateral by mid-day.
- 11:00 ET: Earnings season is in full swing. The desk's ETF inventory in single-name-heavy ETFs (like XLK with concentrated MSFT and AAPL exposure) is being marked daily through earnings. The desk has hedged via index futures, but the basis between XLK and the hedge is moving, costing a few bps daily.
- 14:00 ET: The desk head meets with the global head of equity financing for the monthly review. Topics: balance sheet usage trending up (currently \$52bn of long inventory vs the \$50bn target); HTB exposure within limits but rising (3.8% of inventory, vs 5% cap); client concentration acceptable (largest client is 12% of swap book, below the 15% threshold). The meeting includes a discussion of whether to add a senior trader to the EM team (yes, recruiting will start).

Thursday.

- 10:00 ET: Quarterly reset day for the desk's structured note book. Notes referencing baskets reset their financing legs based on the prior 3-month average. The desk computes the resets, books any small P&L from prefunded hedges, and passes the new pay rates to clients.
- 13:30 ET: A hedge fund client requests an unusual structure: synthetic short Tesla via long puts, short calls, hedged with a delta-one swap on the broad market to neutralize beta. The trade is bespoke and requires the structuring team, the equity derivatives desk, and the delta-one desk to price it together. After two hours of back-and-forth, the trade is quoted; the client passes (the price was higher than they wanted), and the trade dies. The desk has spent two senior people-hours on a non-trade — a normal cost of doing business.
- 16:30 ET: The desk's monthly compensation pre-discussion. The senior trader on the EM book is being approached by a hedge fund offering 50% more total comp. The desk head's role: try to retain by emphasizing the franchise stability, the deferred-stock value, and the medium-term promotion path. Sometimes this works;

sometimes the trader leaves anyway.

Friday.

- 09:30 ET: Index rebal day for a smaller index. Modest flow; the desk’s pre-positions perform well, capturing about \$800k in basis-point gains.
- 11:00 ET: The IPO lockup-expiration team prepares for next quarter’s wave of expirations. Lockup expirations create predictable selling pressure 6 months after IPO; the desk prepositions short via swap structures with hedge fund clients.
- 14:00 ET: Weekly P&L attribution. Total weekly P&L: +\$28m. Decomposition: financing income \$11m, sec lending \$9m, dividend wedge \$3m, execution \$5m. The desk head emails the global head a one-paragraph summary. The week is below the long-run average (\$30-35m typical) but within normal variance.
- 16:00 ET: The senior traders go for drinks at Brookfield Place across the street. The juniors stay an extra hour to clean up positions and write end-of-week notes. By 18:30 the floor is mostly empty, save for a few traders monitoring overnight Asian markets via remote chat.

This is the texture of the work: relentless, multi-channel, highly leveraged on relationships and judgment, and surprisingly office-like in its hour-by-hour rhythm. The dollar amounts are vast — a typical week sees \$5-15bn of new notional booked — but the actual decisions are fast, mostly mechanical, and embedded in a culture of risk discipline that has been refined over forty years.


1.9 1.9 References

- Choudhry, M. (2010), *The Repo Handbook*, Wiley.
- Adrian, T. and Shin, H.S. (2010), “Liquidity and Leverage”, *Journal of Financial Intermediation*.
- Gorton, G. and Metrick, A. (2012), “Securitized Banking and the Run on Repo”, *Journal of Financial Economics*.
- Duffie, D. (2018), “Post-Crisis Bank Regulations and Financial Market Liquidity”, *Banca d’Italia Lectures*.
- Lubin, R. (2019), “Inside the Modern Equity Trading Floor”, *Risk Magazine*, June.
- Cetorelli, N. and Peristiani, S. (2012), “The Role of Banks in Asset Securitization”, *FRBNY Economic Policy Review*.
- Goldman Sachs (2023), Annual Report — Equities Segment Disclosure.

2 Chapter 2 — Securities Lending Fundamentals

2.1 2.1 The Chain From Beneficial Owner to Short Seller

A share of AAPL doesn’t just appear in a short seller’s account. There is a chain:



Every link extracts a fee. The top-of-stack beneficial owner — the pension plan that actually owns the share — earns somewhere between 25% and 70% of the all-in fee charged to the bottom-of-stack short seller, with the residual split among the agent lender and the prime broker.

Concretely on a \$100m AAPL loan at a 25 bps annual fee:

Total annual fee charged to short seller:	\$250,000	
Prime broker keeps (markup):	\$50,000	(20%)
Agent lender keeps (split):	\$60,000	(24%)
Beneficial owner receives:	\$140,000	(56%)

For specials at 200 bps the absolute dollars are 8x bigger but the percentage splits are roughly the same. This is the bread and butter of the agency lending business — high-volume low-margin loans on the long tail of GC names, with episodic large gains on specials.

2.2 2.2 Cash vs Non-Cash Collateral

When a short seller borrows a stock, they post collateral. In the US, the dominant form is **cash collateral**, typically posted at 102% of the loan value (European markets often use 105%). The lender invests the cash collateral in money-market instruments and pays back to the borrower a *rebate rate*.

[Cash collateral path - US standard]

Borrower posts \$102 cash for every \$100 of stock borrowed.
 Lender invests \$102 in repo, T-bills, money funds.
 Lender pays borrower a "rebate" rate (e.g., SOFR - 25 bps).
 Lender keeps the spread between investment yield and rebate.
 The "borrow fee" is implicit: SOFR investment yield - rebate paid.

In Europe and parts of Asia, **non-cash collateral** — sovereign bonds, equities, corporates — is more common. Here there is no investment yield; the lender simply charges an explicit fee:

[Non-cash collateral path - European standard]

Borrower posts \$105 of government bonds.
 Lender holds bonds, no reinvestment.
 Borrower pays explicit fee (e.g., 25 bps annualized).
 This is called a "fee-based" loan.

US institutional money over the past decade has migrated partially toward fee-based markets even when collateral is cash, because of the **reinvestment risk** the 2008 crisis exposed.

2.3 2.3 The 2008 AIG Lesson on Reinvestment Risk

AIG's securities-lending program, in 2007–2008, took cash collateral from short sellers and invested it in *subprime mortgage-backed securities* yielding far above SOFR. When MBS prices collapsed, AIG had to return cash collateral to short sellers at par. The losses ran into tens of billions and were a material contributor to the AIG bailout.

The post-2008 rule of thumb:

[Reinvestment guidelines - modern industry standard]

Cash collateral may only be invested in:

- Tri-party repo against US Treasuries / Agencies
- Short-dated T-bills (<3 months)
- Government money market funds (Rule 2a-7 government)

Forbidden:

- Corporate paper, commercial paper > A-1
- Asset-backed securities of any kind
- Long-duration anything

Beneficial owners now sign "investment guidelines" that explicitly cap reinvestment risk. This costs them 2–5 bps of yield versus the pre-2008 era but eliminates the AIG-style tail.

2.4 2.4 General Collateral vs Specials

The market segments by *availability*:

General Collateral (GC). A name where supply (lending inventory) far exceeds demand (short interest). Most large-cap, S&P 500 names. The borrow rate sits very near SOFR. Examples: AAPL, MSFT, JPM, KO. Borrow rate: 8–25 bps annualized. The desk earns the spread between SOFR investment of cash collateral and the rebate paid back.

Special. A name where short demand exceeds easy supply, driving the borrow rate well above SOFR-equivalent. Example: BBY in Q1 2021 traded at 80% per annum. AMC at certain points exceeded 100%. Special rates can persist for weeks or for years (borrow rates on small-cap biotech with structural short interest can sit at 30–60% indefinitely).

A typical universe distribution at a Tier 1 prime broker:

Tier	% of names	% of balance	Rate range
Deep GC	70%	80%	8–25 bps
Light GC	18%	12%	25–100 bps
Warm	8%	5%	100 bps – 5%
Special	3%	2%	5% – 30%
Hot/HTB	1%	1%	30% – 200%+

Despite hot/HTB being only 1% of names, on a \$50bn book the 50 bps in fees these generate is $50 \text{ bps} \times \$500\text{m} \times \sim 50\%$ blended fee = roughly \$125m/year. The long tail of specials is the single most profitable line item on a major prime broker's P&L per dollar of balance sheet.

2.5 2.5 Recall Risk

A loan is callable. Either side can recall:

- **Lender recall.** Beneficial owner sells the stock or the agent reallocates. The borrower must return shares within standard settlement (T+1 in the US after May 2024).
- **Borrower return.** Short seller closes their position; returns the borrow.

For a hedge fund running a long-running short, a recall is a real risk. If the short can't be re-borrowed, the hedge fund must cover, often at a worse price. Prime brokers generally promise *best efforts* re-borrowing but no guarantee.

The recall risk is asymmetric:

- For a vanilla GC short (e.g., short JPM), recall risk is essentially zero.
- For a hard-to-borrow short (e.g., short BBY in 2020), recall risk is *the* dominant risk.

The classic recall scenario: a short fund borrows BBY at 25% per annum. The position works, the short rises, the desk is paid handsomely. Then the desk's beneficial owner sells the underlying. The desk recalls. The short fund must cover; the act of covering creates a short squeeze; the position that was profitable yesterday is now a 40% loss. This is the GME/AMC story from one specific angle.

2.6 2.6 The Loan Lifecycle: A Concrete Example

A short seller wishes to short 100,000 GameStop (GME) shares. GME is special. The lifecycle:

Day 0 — Locate. The short fund's broker (the Prime) checks its locate inventory. The Prime's trading desk reaches out to its agent lender contacts. A locate is provided at an indicative borrow rate.

Locate quote: 100,000 GME at indicative 25% per annum
 Subject to allocation at start of day
 Subject to recall on availability change

Day 1 — Trade. The short sells 100,000 GME at \$35 = \$3.5m proceeds. The Prime borrows the shares from the agent.

Cash collateral posted: \$3.57m (102%)
 Loan term: Open (until recalled or returned)
 Daily fee accrual: $\$3.5m \times 25\% / 360 = \$2,431/\text{day}$

Of which:

Borrower (short fund) pays:	\$2,431
Prime broker takes ~25% markup:	\$608
Agent ~25% take of remaining:	\$456
Beneficial owner receives:	\$1,367

Day 30 — Position holds. GME has moved to \$32. The short is up \$300k mark-to-market. The borrow has accrued $\$2,431 \times 30 = \$72,930$ in fees, of which the short paid \$72,930 minus rebate. Assume the rebate is SOFR — 25% = roughly —20.5% net (negative rebate because the borrow is so deep). The cash collateral earns SOFR 5.30%, but the rebate paid back is —19.7%, so the *spread* the lender keeps is roughly 25%. The short paid \$72,930 net.

Day 45 — Recall. The beneficial owner (a pension fund) decides to liquidate GME from its index portfolio. The agent recalls the 100,000 shares from the Prime. The Prime must return shares within T+1.

The Prime's options: 1. Re-borrow elsewhere — possible if locate inventory exists. 2. Force the short fund to cover (if no re-borrow available).

Suppose only 60,000 can be re-borrowed at 35%. The Prime tells the short fund: “You have 30,000 to cover by tomorrow.” The short covers 30,000 at the bid, reducing its position size. The fund's exposure profile changes mid-trade.

This recall friction is exactly why some institutional shorts pay enormous “term borrow” premiums — fees of 5%, 10%, even 20% of notional — to lock in non-recallable borrow for 30, 60, 90 days. Term borrow is its own market.

2.7 The Tri-Party Repo and Sec-Lending Plumbing

In the US, the daily settlement infrastructure for collateralized lending runs through *tri-party agents* — Bank of New York Mellon (BNY) and JP Morgan are the two. The tri-party agent stands between the borrower and lender, holding collateral in segregated accounts, marking it to market daily, and managing substitutions.

A typical day in tri-party sec lending:

```
07:00 ET  Tri-party agent unwinds previous day's positions.
          Cash returns to lenders, securities return to borrowers.
07:30 ET  Lenders post their availability to the matched-book.
08:30 ET  Borrowers (PB clients of hedge funds) request locates.
09:00 ET  Allocations finalized - locate confirmations sent.
09:30 ET  Trading begins. New shorts use the morning's locates.
15:30 ET  End-of-day mark-to-market on collateral.
          Variation margin movements computed.
17:00 ET  Tri-party agent re-collateralizes positions for the next day.
```

The friction at each step:

- Locate availability is *intraday* — a locate confirmed at 09:00 may not be available at 14:00 if a beneficial owner sells.
- Variation margin is daily — a sharp price move triggers same-day collateral movement.
- Re-collateralization can fail if the borrower doesn't have eligible collateral on hand.

For a trader, the practical implication is that a “borrowed” stock is borrowed *for one business day at a time*. The continued availability is implicit — if the locate disappears, the position must close. This is the recall risk that retail short sellers (especially via Robinhood's HTB list) experience as “your short was closed for you.”

2.8 The Lending-Indemnification Question

A central legal question in agency lending: who bears credit risk if the borrower defaults?

Pre-2008, indemnification was standard — the agent lender contractually guaranteed the loan against borrower default. Post-Lehman, indemnification became expensive (the agent had to hold capital against the indemnification commitment), and many programs went non-indemnified.

Today the market is split:

- **Indemnified programs.** Common at smaller pension and insurance lenders. The agent absorbs losses if the borrower defaults. The agent's fee share is higher (50-60%) to compensate.
- **Non-indemnified programs.** Common at large sovereign and pension-mega-lenders. The lender bears default risk. The lender's fee share is higher (65-75%).

For the desk on the borrowing side, this affects who the desk talks to. Sovereign-wealth lenders with non-indemnified programs are the cheapest source of borrow but also the strictest about collateral eligibility. Smaller indemnified programs are more flexible but slightly more expensive.

2.9 2.9 Worked Example: The Lifecycle of a Single Loan in Detail

A short fund "ShortAlpha" wants to short 50,000 shares of a small-cap medical device company "MEDX" trading at \$42. MEDX is special — it's been in the news for an FDA delay.

Step 1 — Locate request (09:15 ET).

ShortAlpha calls Morgan Stanley PB: "I want to short 50,000 MEDX."

MS PB checks its locate inventory: - Internal MS clients holding MEDX: 8,000 shares available. - Goldman agency lending program: 12,000 shares available at indicative 18% per annum. - BNY agency lending: 15,000 shares at 22%. - State Street agency: 0 available. - Externally located: 15,000 shares at 25% via brokerage chain.

Total located: 50,000 at blended rate.

Step 2 — Locate confirmation (09:30 ET).

MS PB confirms: "50,000 MEDX located at indicative 22.5% per annum, subject to recall, non-term."

ShortAlpha accepts. The locate is a *promise*, not an actual loan yet. The actual borrow happens when the short sale settles.

Step 3 — Trade execution (10:00 ET).

ShortAlpha sells 50,000 MEDX at \$42 = \$2.1m proceeds.

Step 4 — T+1 settlement (next business day).

The short sale settles. The 50,000 borrowed shares are delivered to the buyer. ShortAlpha's account shows: - Short position: -50,000 MEDX - Cash collateral posted to MS PB: $102\% \times \$2.1\text{m} = \2.142m - Net cash: \$2.1m credit (from short sale) - \$2.142m (collateral) = -\$42k margin requirement.

Step 5 — Daily fee accrual.

Daily borrow accrual on MEDX at 22.5% per annum:

$$\$2.1\text{m} \times 22.5\% / 360 = \$1,312.50/\text{day}$$

Of which:

ShortAlpha pays:	\$1,312.50
MS PB markup (~25%):	\$328.13 (PB revenue)
Agent share (~25% of remainder):	\$246.10
Beneficial owners (split across multiple):	\$738.27

Step 6 — Variation margin (Day 7).

MEDX rallies to \$45 on positive news. ShortAlpha is down \$150k mark-to-market. MS PB issues a variation margin call: - Required collateral: $102\% \times (\$45 \times 50,000) = \2.295m - Already posted: \$2.142m - VM call: \$153,000

ShortAlpha must post \$153k by next morning or face buy-in.

Step 7 — Recall scenario (Day 14).

Goldman’s pension client sells their 12,000 MEDX position. Goldman recalls those shares from MS PB. MS PB has T+1 to source replacement.

MS PB checks the market: - BNY can re-lend 8,000 more at 28% per annum. - State Street has 2,000 at 35%. - 2,000 cannot be re-borrowed.

MS PB tells ShortAlpha: “We need to close 2,000 of your 50,000 MEDX short. The remaining 48,000 will be at a blended 24% per annum starting tomorrow.”

ShortAlpha covers 2,000 at the bid (\$46) — \$92k cost on a position they didn’t want to close. The forced cover is a \$4k loss on those 2,000 shares versus the original short price.

Step 8 — Final close (Day 30).

MEDX has fallen to \$40 on FDA approval delay. ShortAlpha closes 48,000 shares at \$40.

Total P&L for ShortAlpha:

Original 2,000 closed at \$46: $-2000 \times (\$46 - \$42) = -\$8,000$
 Remaining 48,000 closed at \$40: $-48000 \times (\$40 - \$42) = +\$96,000$
 Gross trade P&L: +\$88,000

Borrow cost over 30 days:

Days 1-7 at 22.5% on \$2.1m: $\$1,312.50 \times 7 = \$9,188$
 Days 8-13 at 22.5% on \$2.295m: $\$1,434.38 \times 6 = \$8,606$
 Days 14-30 at 24% on \$1.92m: $\$1,280 \times 17 = \$21,760$
 Total borrow: -\$39,554

Net P&L: +\$48,446

On notional capital deployed (~\$153k VM peak): 31.7% return in 30 days

But on average notional (\$2.1m): 2.3% return

This is the full lifecycle. The key features:

- Borrow cost was 45% of gross alpha — the rate matters enormously.
- Forced cover during recall destroyed \$8k of expected return.
- Variation margin calls required sourcing \$153k cash mid-trade.
- The “true” return on a fully-funded basis (including margin and proceeds) was ~2.3% on the gross notional, not the headline 31.7%.

For a multi-billion-dollar fund running thousands of such positions, the operational complexity is the entire game. The PB earns its fees by handling all of this on behalf of the client.

2.10 The Loan Board — Inside the Sec Lending Pod

The numerical mechanics of stock loans, while elegant on paper, are mediated in practice by a piece of software called the *loan board*. Every major prime broker has one. At Goldman it’s a proprietary system known internally by an unflattering acronym; at Morgan Stanley it’s “Liquidity Tools”; at JPMorgan it’s part of the broader “Athena” platform. State Street, BNY Mellon, and Brown Brothers Harriman — the three largest agency lenders — each run their own. Vendors like EquiLend, FIS Astec, and Pirum sit in the middle and stitch the proprietary systems together for cross-firm matching.

The screen view of the loan board is unforgiving in its density. A typical layout, reconstructed without bank-specific UI choices:

[Stylised loan board view, single trader's screen]

TICKER	SI%	DTC	AVAIL	RATE_OPEN	RATE_1WK	RATE_3M	TARGET	PNL_DAY
AAPL	0.7	0.5	47.2M	8 bps	10 bps	14 bps	hold	+\$22k
MSFT	0.4	0.3	89.1M	7 bps	8 bps	11 bps	hold	+\$18k
TSLA	3.2	2.1	12.4M	85 bps	110 bps	180 bps	INCREASE	+\$140k
GME	18.4	4.7	0.2M	2200 bps	2800 bps	4000bps	RECALL?	+\$840k

AMC		14.2		3.9		0.4M		1100 bps		1450 bps		1800bps		hold		+\$320k
MSTR		9.1		1.4		1.8M		450 bps		580 bps		750 bps		hold		+\$180k
...		...		...		...		...		...		...		...		...

SI% is short interest as a percentage of float. DTC is days to cover (short interest divided by 30-day average daily volume). AVAIL is shares available to lend in the desk's pool *right now*. RATE columns are the desk's own quoted rates at different tenors. TARGET is a human-set flag for actions to take when conditions shift. PNL_DAY is the day's P&L on the existing book of loans on this name.

The senior trader on the loan board pod scans this every morning. Names with red-flag combinations — high SI, low availability, climbing rate — are the focus. A typical morning scan might surface fifty names worth a closer look out of two thousand on the screen. The trader picks five to ten and digs in: who is borrowing, who is lending, what corporate actions are pending, what news is on the tape.

The phone rings constantly. A call from a hedge fund client's PB rep: "Hey, I see GME locate at 22%, can we go term to lock in three months?" The trader checks the term curve, talks to the agent lender contacts, gets a price back. "I can do three months at 28%." The client: "Done." Five seconds. \$4m of additional revenue is now booked over the term, assuming the position holds.

A different call, twenty minutes later: "Goldman's lending us 50,000 BBBY at 80%, but they're flagging a possible recall on Friday — what's our backup?" This is the kind of question that defines the trader's day. The answer requires checking three or four other lenders, estimating the probability of recall, and quoting an alternative borrow at, say, 95% with a different agent. The hedge fund accepts or rejects; either way, the trader has spent five minutes and built the relationship one call deeper.

The recall call itself — when a beneficial owner pulls a loan — has a particular tone. The agent lender's representative on the line is usually polite but firm: "We need our 100,000 GME back by tomorrow. Our pension client is liquidating." The desk has approximately T+1 to source replacement, force a cover, or accept the unwind. The desk's response involves calling every other lender, every internal source, every tri-party agent. If shares can be re-borrowed, the position rolls to the new lender at the new rate (which is almost always higher). If not, the hedge fund client is told: "You have until 16:00 tomorrow to cover or we will buy in." This call goes to the client; the client either covers gracefully or panics. In a typical year, a Tier 1 desk handles thousands of recall events; perhaps a hundred trigger a forced buy-in.

The Tri-Party reconciliation meeting itself, which occurs daily across the major banks, is best understood as a controlled chaos session. The bank's operations team and the tri-party agent (BNY or JPM) reconcile the day's collateral movements line by line. The meeting starts at 14:00 ET sharp. Half a dozen people on each side. Every break — every disagreement between the bank's ledger and the agent's — is investigated. Most breaks are simple: a settlement that didn't process on time, a corporate action that the agent and the bank classified differently, an FX timing mismatch on a euro-denominated position. Each break has a "look-back" requirement: you must reconcile within 24 hours or escalate. Persistent breaks become operational risk events that get reported up the chain.

The tone of the recon call is procedural. Nobody yells. The senior operations person from the agent reads the day's mismatches in a flat tone. The bank's senior ops person responds. They jointly assign each break to a team member to investigate. The meeting ends in 30 minutes on a calm day, two hours on a stressful one (March 2020, August 2024). On stressful days, the meeting includes risk management observers and sometimes a CFO-level escalation if the breaks involve collateral that has fallen below required haircuts.

When an indemnification triggers — which is to say, when a borrower defaults and the agent must make the lender whole — the procedure is precise but rare. The most recent industry-wide trigger event was the September 2008 Lehman Brothers collapse. Lehman was a major borrower in the agency lending pools at State Street, Northern Trust, and BNY. When Lehman entered bankruptcy on September 15, those agents had to make their pension and mutual fund lender clients whole on the borrowed positions that Lehman had not yet returned. The total indemnification claim across the industry was estimated at \$5-8bn. Of that, the agents recovered substantial amounts from Lehman's bankruptcy estate over the following years — but the *immediate* cash outlay (or rather, the repurchasing of equivalent positions in the open market) was material. State Street took a \$200m+ charge directly attributable to indemnification on Lehman-related lending.

Post-Lehman, the indemnification machinery was redesigned. Most major agents now require borrowers to post collateral marked-to-market more frequently (intraday at Tier 1 borrowers), require larger collateral haircuts, and limit indemnified exposure to specific sub-pools rather than the entire program. The cost of indemnification —

measured as the reserve required to fund a Lehman-style event — is now reflected in the agent’s fee share, which is typically 5-8 bps higher on indemnified programs than on non-indemnified ones.

The day-to-day routine on the sec lending pod has not been fundamentally re-architected since the 1990s. The screen densities are higher. The matching algorithms are smarter. The post-trade ops are more automated. But the core human workflow — locate request, quote, accept, book, monitor, recall, replace — is unchanged. A senior sec lending trader from 1995 would, after a brief tour of the screen layout, find themselves perfectly capable of running the modern desk. The product is conceptually static. What has changed is the speed, the data quality, and the regulatory load.

2.11 2.11 The Mid-Sized Lender’s Calculus

A useful counterpoint to the Tier 1 narrative: how does a *smaller* beneficial owner — say, a \$5bn university endowment or a \$20bn state pension — make the choice to enter or exit a securities lending program? The decision is not trivial, and the way it gets made illustrates the fee structures and risk tradeoffs from the lender’s perspective rather than the borrower’s.

Consider the State of Michigan Retirement Systems (SMRS), with roughly \$90bn AUM as of 2024. SMRS holds large positions in US large-cap equities and US Treasuries. Their securities lending program, run by State Street, generates roughly \$40-60m annually in fee income — a meaningful but not enormous contribution to overall returns. The decision to be in the program rests on three considerations:

1. Fee income vs. risk borne

- Indemnified: lower fee share (~50%), agent absorbs default risk
- Non-indemnified: higher fee share (~65-70%), SMRS absorbs default risk

Decision: SMRS chose non-indemnified in 2018, betting on State Street's selection of borrowers and their tri-party collateral quality.

2. Reinvestment risk

- Pre-2008: cash collateral could be invested in commercial paper, corporate bonds. Higher yield, higher risk.
- Post-2008: limited to Treasury repo and government MMFs. Lower yield, near-zero risk.

Decision: SMRS limits reinvestment to government repo only. Yield is ~10 bps below the borrow rebate, captured by State Street as a fee.

3. Voting rights

- Lent shares lose voting rights during the loan period.
- Active equity managers (typically external) may want to vote on specific proxy contests.

Decision: SMRS has a recall policy: shares are recalled across contested proxy votes when the position represents >0.5% of the issuer or when an active manager flags a vote of substantive importance.

The trustee committee of SMRS reviews the lending program annually. The 2023 review concluded the program added 4-6 bps annually to total returns net of risks borne — a modest contribution but worthwhile given the operational simplicity (State Street handles everything). The committee has occasionally debated exiting the program, particularly after the 2008 AIG-style episode, but has consistently elected to remain in.

Smaller pensions and endowments make similar calculations. The threshold below which a lending program is uneconomic is roughly \$500m AUM — below this size, the agent’s setup costs exceed the expected fee income. Above this, almost every institutional pool participates in some form of lending program, with the specific structure (indemnified vs not, cash vs non-cash collateral, fee split) negotiated based on the lender’s size and bargaining power.

For ultra-large lenders — the \$500bn+ pension systems (CalPERS, ABP, Norges Bank Investment Management) — the negotiation looks more like a *vendor selection* than a typical lending arrangement. These lenders auction

their lending business across multiple agents, with explicit performance benchmarks and fee renegotiations every 3-5 years. The agent’s fee share for these mega-lenders is typically 30-40% (lower than for smaller pensions because the bargaining power is reversed).

The flow of money up and down the chain — from beneficial owner to short seller — therefore does not simply split along the textbook 25/25/50 lines. It is a continuous negotiation, with the larger players keeping a higher share. The \$1.9bn in industry SLB revenue from the top 2% of names is split such that, roughly, the largest five beneficial owners (BlackRock as a pool manager, Vanguard, the largest US pensions) capture 60-70% of the fee value on their own holdings, while the long tail of smaller pensions and mutual funds take what they can get.

2.12 2.12 References

- D’Avolio, G. (2002), “The Market for Borrowing Stock”, *Journal of Financial Economics*, 66(2-3): 271-306. The foundational paper on the equity lending market.
- Saffi, P.A.C. and Sigurdsson, K. (2011), “Price Efficiency and Short Selling”, *Review of Financial Studies* 24(3): 821-852.
- Kolasinski, A.C., Reed, A.V., and Ringgenberg, M.C. (2013), “A Multiple Lender Approach to Understanding Supply and Search in the Equity Lending Market”, *Journal of Finance* 68(2): 559-595.
- Krishnamurthy, A., Nagel, S., and Orlov, D. (2014), “Sizing Up Repo”, *Journal of Finance* 69(6): 2381-2417.
- Aggarwal, R., Saffi, P.A.C., and Sturgess, J. (2015), “The Role of Institutional Investors in Voting: Evidence from the Securities Lending Market”, *Journal of Finance* 70(5): 2309-2346.
- Risk Management Association (2024), *Securities Lending Quarterly Aggregate Report Q4*.
- IHS Markit (2023), *Securities Finance Year in Review*.
- State Street (2023), *Agency Lending Risk and Practice Disclosure*.

3 Chapter 3 — The Borrow-Rate Stack

3.1 3.1 From Risk-Free Rate to Rebate

The borrow rate paid by a short seller in cash-collateral US markets is *not* directly quoted. What is quoted is the **rebate** — the rate the lender pays back to the borrower on the cash collateral. The borrow fee is implicit:

$$\text{borrow_fee} = \text{Risk-Free Rate} - \text{Rebate}$$

When SOFR is 5.30% and the rebate is 5.05%, the borrow fee is 25 bps — a typical GC name. When SOFR is 5.30% and the rebate is 4.30%, the borrow fee is 100 bps — a warm name. When SOFR is 5.30% and the rebate is *negative* (e.g., -20%), the borrow fee is 25.30% — deep special territory.

A negative rebate means: the short seller hands over \$102 in cash collateral and receives back, effectively, *less than nothing* — the lender is keeping not only the spread but additional fee directly from the short. This is the “I pay you to hold my collateral and lend me your stock” arrangement.

3.2 3.2 A Worked Numerical Waterfall

Consider a \$1m short position in a hard-to-borrow biotech, “DRUG”, at a quoted **borrow fee of 8%** annualized. SOFR is 5.30%.

[Day 0 - Trade entry]

Short sale proceeds:	\$1,000,000	(at the broker)
Cash collateral required:	\$1,020,000	(102% standard)

The cash collateral comes from:

Short proceeds:	\$1,000,000
Margin requirement (fund cash):	\$20,000
Total posted:	\$1,020,000

[Daily accrual - for 30 days, assuming flat price]

Cash collateral invested at SOFR: $\$1,020,000 \times 5.30\% / 360 = \$150.17/\text{day}$
 Rebate paid to borrower: $\text{SOFR} - \text{borrow_fee} = 5.30\% - 8.00\% = -2.70\%$
 $\$1,020,000 \times (-2.70\%) / 360 = -\$76.50/\text{day}$

This is debited from the short fund: $+\$76.50/\text{day}$ debit
 Plus the short does not earn the SOFR yield on its own short proceeds
 (the lender keeps it): $\$150.17/\text{day}$ opportunity cost

Net cost to short fund per day: $\$76.50 + \$150.17 = \$226.67/\text{day}$ (8.27% annualized, on \$1m)

[30 days later]

Total cost to short fund: $\$6,800$
 Total revenue to lending stack: $\$6,800$
 Prime broker share (~25%): $\$1,700$
 Agent share (~25% of residual): $\$1,275$
 Beneficial owner (~50%): $\$3,825$

Now consider what happens when the borrow fee is **22%** — DRUG goes deep special after a positive Phase 3 readout that crushed shorts.

[Borrow fee = 22%]

Daily cost to short fund (on \$1m):
 $\$1\text{m} \times 22\% / 360 = \$611.11/\text{day}$
 Annualized cost: 22% of notional

For a short fund expecting the stock to fall 10% over 90 days:

Expected gross gain: $10\% \times \$1\text{m} = \$100,000$
 Borrow cost over 90: $\$611 \times 90 = \$55,000$
 Net expected gain: $\$45,000$
 Effective borrow drag: 55% of gross alpha

The borrow rate alone reduces a high-conviction short's expected return by more than half. This explains why the short side of a market with 22%+ borrow rates is dominated by very short-term, very high-conviction names — long-running shorts can't justify the carry cost.

3.3 3.3 The “Specials Concentration” of Sec Lending Revenue

Industry data from the major lending platforms (IHS Markit, S&P Global Securities Finance) consistently shows:

% of names by fee bucket	Share of total industry SLB revenue
Bottom 90% (GC <50 bps)	~30%
Next 8% (warm 50-500 bps)	~25%
Top 2% (specials 500+ bps)	~45%

On a \$5 trillion industry-wide loan balance generating roughly \$12bn in annual fees, the specials concentration means roughly \$5.4bn comes from the 2% of names with rates above 500 bps. This is why every major prime broker has an entire “specials trading” subdesk whose only job is to find, source, and price specials. They are the highest-margin product on the floor.

3.4 3.4 The Term Structure of Borrow

Borrow rates have a term structure. A loan can be open (callable any day), 1-week, 1-month, 3-month, or term (locked for the life of a position). Term borrow trades at a premium because the lender is foregoing optionality:

Tenor	GC name (bps)	Light GC (bps)	Warm (bps)	Special (bps)
Open	12	50	250	1500
1-week	14	60	280	1700
1-month	18	80	350	2000
3-month	25	120	500	2800
6-month	35	180	700	3500

The slope is steeper for warmer names because the lender bears more risk that the rate moves further special during the term. On a 6-month term borrow of a name currently at 1500 bps (15%), the 6-month price of 3500 bps (35%) implies a 20% per annum *option premium* for non-recall protection.

3.5 The “On Special” Threshold and Why It Matters

A name is colloquially “on special” when its borrow fee exceeds a critical threshold relative to risk-free. The standard market threshold is **borrow fee > SOFR** — i.e., the rebate goes negative. At this point:

1. The borrower pays *more* than the cash collateral earns.
2. The lender keeps *more* than just the spread.
3. The position becomes meaningfully expensive on a multi-month carry basis.

Below the threshold, borrow is cheap (a few bps) and the trade economics are dominated by the directional view. Above the threshold, borrow becomes a *first-order* cost.

Concretely: - SOFR = 5.30%, borrow fee = 25 bps: rebate = 5.05%, lender keeps the spread = 25 bps. Borrow cost on \$1m = \$2,500/year. Trivial. - SOFR = 5.30%, borrow fee = 5.30%: rebate = 0%. Lender keeps full SOFR yield. Borrow cost = \$53k/year. Material. - SOFR = 5.30%, borrow fee = 15%: rebate = -9.70%. Borrower pays 9.70% out of pocket. Borrow cost = \$150k/year. Trade-defining.

The threshold for “structural alpha matters” depends on the trade’s expected return. A short with 30% expected gain over 6 months can absorb 5% borrow cost; a short with 5% expected gain over 6 months cannot.

3.6 The Borrow-Rate Forecasting Problem

Predicting borrow rate moves is its own discipline. Major drivers:

1. **Short interest changes.** Demand-side. As more shorts pile in, the rate rises.
2. **Lending supply changes.** Supply-side. As beneficial owners enter or exit lending programs, supply shifts.
3. **Corporate actions.** Pending mergers, dividends, or rights offerings affect demand.
4. **Idiosyncratic news.** Earnings, lawsuits, FDA decisions can spike demand.
5. **Macro factors.** During risk-off periods, broad short interest rises; during risk-on, it falls.

A typical professional rate-forecasting workflow:

Daily inputs:

- Short interest (from FINRA, biweekly)
- Failed-to-deliver volume (SEC, weekly)
- Days-to-cover (short interest / avg daily volume)
- Borrow inventory (proprietary; major PB systems track this)
- Lending supply (proprietary; agency lender data)

Daily outputs:

- 1-week forecast borrow rate
- 1-month forecast
- Recall probability index
- Squeeze probability index (combination of high SI, low supply, high vol)

The major data providers (S&P Global Securities Finance, IHS Markit / S&P Global, EquiLend, FIS Astec) sell professional borrow-forecasting data at roughly \$50k-200k/year per major dataset. This is institutional infrastructure; retail does not have access to comparable data.

3.7 3.7 Implied Borrow From Options

A useful technique: extract the implied borrow rate from option prices via put-call parity. For non-dividend-paying stock:

Put-call parity:

$$C - P = S \times e^{-q \times T} - K \times e^{-r \times T}$$

where q is the implicit borrow rate (the cost of synthetic short).

Solving:

$$q = -(1/T) \times \ln((C - P + K \times e^{-r \times T}) / S)$$

For a stock $S = \$50$, $K = \$50$ ATM, $T = 30$ days, $r = 5\%$, observed $C = \$1.10$, $P = \$1.30$:

$$C - P = \$1.10 - \$1.30 = -\$0.20$$

$$K \times e^{-0.05 \times 30/365} = \$50 \times 0.9959 = \$49.795$$

$$S \times e^{-q \times 30/365} = \$50 \times e^{-q \times 0.0822}$$

$$C - P = S \times e^{-q \times T} - K \times e^{-r \times T}$$

$$-0.20 = 50 \times e^{-0.0822q} - 49.795$$

$$50 \times e^{-0.0822q} = 49.595$$

$$e^{-0.0822q} = 0.99190$$

$$-0.0822q = \ln(0.99190) = -0.00813$$

$$q = 0.0989$$

Implied borrow rate 9.89% per annum.

This is “expensive” — the options market is pricing the underlying as having ~10% per annum borrow cost. Compare to the actual borrow market: if PB is quoting the locate at 8.5%, the *option-implied* borrow is *higher* than the actual borrow. Translation: option prices may be slightly mispriced (puts too rich relative to calls), or there may be expected stress that the borrow market hasn’t yet priced in.

This put-call parity-implied borrow is a clean signal. Sophisticated retail can compute it from any option chain. When the implied borrow exceeds the actual borrow by >100 bps, there is a structural skew premium that may be capturable.

3.8 3.8 Why Specials Exist — The Intuition

Section 3.3 noted that 2% of names produce 45% of industry SLB revenue. That stark concentration deserves its own explanation, because it is one of the most counterintuitive features of equity finance for newcomers. Why would the lending market not arbitrage this away? Why don’t beneficial owners supply more shares to specials, driving down the rate? Why don’t more shorts cover when the carry hits 50%, neutralizing the demand?

The answer is that specials persist because of three structural frictions that are invisible to outsiders.

The first friction is *supply rigidity*. The lending pool for any given stock is fixed in the short run. It consists of the shares held by beneficial owners who have signed up to lend (in pension trusts, mutual fund custody accounts, sovereign wealth holdings, and so on). New supply does not appear when the rate rises, because the holders who *would* lend are already lending — every share they own that is in a lending-eligible account is already in the pool. The marginal supply response to a rising fee is therefore tiny: maybe a few percent of additional shares come online if an asset manager extends a sub-fund into lending, but most lendable shares are already there. The supply curve is nearly vertical.

The second friction is *demand inelasticity*. Short sellers face a high carrying cost (the borrow fee) but face an even higher cost from being forced to cover at the wrong moment. A hedge fund that has analyzed Camber Energy or Bed Bath & Beyond and concluded the equity is a zero will continue paying 80% borrow rates because their forecast return is 100%, not because they want to. The demand curve is therefore not a smooth function of the rate; it is closer to vertical for a group of high-conviction shorts and flat for a group of opportunistic shorts. The rate clears at the intersection, which can be anywhere depending on the conviction distribution.

The third friction is *information asymmetry about future supply*. A short seller borrowing at 50% today does not know whether the lending pool will be recalled tomorrow. The risk of recall is itself a service that the lender provides, and the lender must be paid for accepting it. A high rate is partly a payment for the *option* the lender is implicitly writing — the option to recall when conditions change. This is why term borrow trades at a premium over open borrow: term locks the option, and the locking cost is real economic value.

Together, these frictions mean that a special name can stay special for a long time. The borrow rate on AMC sat above 10% for most of 2021 even after the meme-rally peaked, simply because short interest remained high (hedge funds with conviction) and supply remained limited (passive index holders not motivated to expand lending into a controversial name). The market does not converge to a “fair” rate quickly because the participants are not playing the same game.

3.9 The Term Structure — Why It Slopes the Way It Does

Section 3.4 presented a numerical term structure table without explaining the slope. The slope deserves attention because it tells the reader what the market believes about the future.

For a *general collateral* name, the term structure is gently upward-sloping: open at 12 bps, three-month at 25 bps, six-month at 35 bps. The slope reflects the lender’s option-locking premium and the modest expected drift in rates over the period. Roughly: the lender is foregoing the option to redeploy at a higher rate if conditions tighten, and the borrower is paying for term protection. The slope is similar to the slope on the GC repo curve and is driven by the same forces (Fed expectations, balance-sheet seasonality, regulatory month-end effects).

For a *warm* name (250 bps open), the slope is steeper: term six months at 700 bps. This is roughly 3x the open rate. Why? Because warm names have a wider distribution of possible future rates. The name might cool to GC (50 bps) in three months, or it might get squeezed to 1500 bps. The lender pricing six-month term has to be compensated for the chance that the rate runs to 1500 bps but they’re locked at the term price. The implied volatility of warm-name rates is much higher than GC-name rates, and the term structure encodes this through a steeper slope.

For *specials* (1500 bps open going to 3500 bps six-month term), the slope reflects something more dramatic: the *recall optionality* the lender is writing. On a name like GME at 25% per annum, the lender’s loan is callable at any time. If the name suddenly becomes worth 50% per annum, the lender wants to recall and re-lend. By writing a six-month term at 35%, the lender accepts a fee that is *between* the current rate and the worst-case future rate, weighted by the probability of getting there. The slope is steep because the upper tail (a 100%+ rate event) is a real possibility.

A useful way for the retail reader to internalize this: the term structure of borrow rates is, in essence, a *forward curve* for the rebate, just like SOFR has its own forward curve. A retail trader looking at a term-borrow curve quoted to them by IBKR is implicitly seeing the market’s belief about the future path of rates on that name. A flat term structure means the market expects the rate to stay roughly where it is; a steep term structure means the market expects the rate to rise (or, equivalently, that the lender is demanding compensation for the risk that it does).

3.10 Three Worked Examples Across Regimes

The single-regime example in Section 3.2 is useful but does not show how the borrow stack behaves across different macro environments. Here are three side-by-side regimes.

Regime 1: ZIRP, late 2020.

Federal Funds rate: 0-0.25%. SOFR: ~0.05%. Equity markets: V-shape recovery from March COVID lows. VIX averaged 25-30 through Q4 2020.

Borrow stack on a typical hedge fund’s book:

General collateral (e.g. AAPL, MSFT, JPM): `borrow_fee` 5-15 bps
On a \$1bn GC short: cost ~\$50k-\$150k/year. Negligible.

Warm names (mid-cap with modest short interest):
`borrow_fee` 50-200 bps
On a \$100m short: cost ~\$500k-\$2m/year. Material but tolerable.

Specials (post-IP0 new issues, hot tech IPOs of 2020):
`borrow_fee` 500-2000 bps

Examples: PLUG, BLNK, FSLY, NIO during certain windows.
On a \$50m short at 15%: cost ~\$7.5m/year. Trade-defining.

Hot/HTB (deep meme stocks, pre-GME wave):

borrow_fee 2000-5000 bps

Limited names. Example: NKLA from June-Sept 2020 ran at ~25%.

The desk's revenue mix in this regime: thin GC margins (because SOFR is near zero, so the spread the desk captures on cash collateral is also near zero), augmented by relatively healthy warm-and-special revenue. Specials concentration was high because the "stay home" stocks (PTON, ZM, DOCU) plus the "EV / hydrogen / SPAC" universe generated sustained short interest. A typical Tier 1 desk reported sec-lending revenue ~\$80-120m/quarter through 2020 — modest because GC was thin, but with episodic spikes from specials.

Regime 2: Peak rates, mid-2024.

Federal Funds rate: 5.25-5.50%. SOFR: ~5.30%. Equity markets: AI rally, mega-cap concentration. VIX averaged 12-15 through the summer until the August unwind.

Borrow stack:

General collateral:

borrow_fee 8-25 bps

On \$1bn GC short: cost \$80k-\$250k/year. Material but expected.

But notably, the desk's revenue from GC is now larger:

$\$1\text{bn collateral} \times (5.30\% \text{ SOFR yield} - 5.05\% \text{ rebate}) = \2.5m/year
captured by the desk on this single \$1bn GC loan. The cash collateral reinvestment yield is the dominant revenue source.

Warm: 50-300 bps

On \$100m short: cost \$500k-\$3m/year.

Specials: 500-2500 bps

Heavy concentration in 2024: TSLA (convert-arb-driven), MSTR (BTC-related), specific small-cap meme revivals.

On \$50m at 15%: cost \$7.5m/year.

Hot/HTB: 2500-6000 bps

Selected names only. Pre-bankruptcy small caps; certain NRG names pre-data-center repricing.

Desk revenue in this regime is *high* across the board. The GC base is now economically significant (because SOFR is high, the cash-collateral reinvestment yield is high, and the spread to the rebate is meaningful). Specials remain a major contributor. Industry-wide SLB revenue in 2024 was estimated at \$14-15bn — the highest level since the pre-Lehman peak. A Tier 1 desk in 2024 reported quarterly SLB revenue in the \$300-450m range.

Regime 3: Crisis, October 2008.

Federal Funds: cut from 2% to 1.5% on October 8, then to 1% on October 29. SOFR predecessor (open repo): volatile, 0.5-3% intraday range. Equity markets: catastrophic decline. VIX peaked at 89 on October 24.

Borrow stack:

General collateral:

borrow_fee 25-60 bps (volatile)

The "GC" in this regime was not really GC - every name was harder to borrow because lenders were uncertain about counterparty risk.

Warm: 200-600 bps (or unavailable)

Many beneficial owners had stopped lending entirely after Lehman. The pool shrank by an estimated 30% in the first two weeks of October. Borrow rates on warm names spiked even though demand was falling - supply was falling faster.

Specials: 1500-4000 bps (or unavailable)

Many specials were on the SEC's emergency short-sale ban list (875 financial names from September 19 to October 8). Borrow on these was technically irrelevant during the ban; afterwards it spiked.

Hot/HTB: simply unavailable

Major financials (LEH pre-bankruptcy, AIG, BSC, GS, MS) had shorts banned outright by SEC emergency order.

The 2008 regime illustrates that *during* a crisis, the borrow market does not behave as a smooth function of demand; it can collapse entirely as supply withdraws. Beneficial owners scared by the AIG reinvestment debacle pulled their lending programs. The agent lenders had margin calls and capital pressures. The price did not clear; the market simply stopped functioning for several names.

The specific damage from this dislocation: short positions that were entered before the crisis at modest borrow rates suddenly faced either skyrocketing rates or *forced cover* because the lender couldn't be replaced. Several quantitative long-short funds reported massive losses in October 2008 not from the directional move but from the *inability to maintain shorts* — they were forced to flatten the short side just as the long side was crashing, double-amplifying the loss. The “August 2007 quant crisis” had been a foreshadowing of this; October 2008 was the full instance.

3.11 3.11 The Retail-Visible Surface

The retail trader does not see the loan board. What they see is the broker's HTB list and the consequences in option pricing. Connecting the institutional plumbing to the retail surface:

When IBKR's HTB list shows AMC at 24.6% on a Tuesday morning, the rate the retail user sees is the *retail rebate-equivalent* — i.e., the cost the user will pay to short. That number is connected to the institutional rate by a margin: the broker captures something between 50 bps and several percentage points depending on the name, the size of the short, and the broker's cost of sourcing. On AMC at 24.6% retail, the institutional rate the broker is paying might be 22-23% — the broker captures the spread.

When a retail trader sees a deep ITM put on a HTB name priced unusually high relative to the call, that pricing is the institutional borrow rate showing up in put-call parity. A retail user fluent in this connection can look at the put-call spread on a name like BBY in 2022 and *back out* the institutional borrow rate without ever calling a broker. The Section 3.7 method (implied borrow from put-call parity) is exactly this.

When a retail trader's short position is closed without warning by the broker, the cause is almost always a recall at the agent-lender level. The chain is: agent lender → prime broker → retail broker → retail user. The retail broker's communication to the user is typically a short notification that the position has been closed; the underlying driver was a beneficial owner three layers up making a portfolio decision.

Understanding this chain matters for retail because it informs which trades to attempt. Shorting a low-fee name at a discount broker is trivial. Shorting a special at any retail broker carries real recall risk and high carry cost. The retail trader who has internalized the chapter's content will price both correctly.

3.12 3.12 References

- Engle, R., Han, R. and Xu, J. (2008), “Repo and the Cost of Equity Lending”, working paper, NYU Stern.
- Cohen, L., Diether, K.B., and Malloy, C.J. (2007), “Supply and Demand Shifts in the Shorting Market”, *Journal of Finance* 62(5): 2061-2096.
- Battalio, R., Engel, S., and Mehran, H. (2017), “The Cost of Short Selling”, working paper.
- Geczy, C.C., Musto, D.K., and Reed, A.V. (2002), “Stocks Are Special Too: An Analysis of the Equity Lending Market”, *Journal of Financial Economics* 66(2-3): 241-269.
- Mitchell, M., Pulvino, T., and Stafford, E. (2002), “Limited Arbitrage in Equity Markets”, *Journal of Finance* 57(2): 551-584.
- Boehmer, E. and Wu, J. (2013), “Short Selling and the Price Discovery Process”, *Review of Financial Studies* 26(2): 287-322.

4 Chapter 4 — Dividend Trades

This chapter is the longest in the guide because dividend trades are the highest-revenue activity on most delta-one desks and are among the least understood outside the institutional world.

4.1 4.1 The Cross-Border Tax Wedge

The fundamental driver of every dividend trade is **cross-border withholding tax**. When a US corporation pays a dividend, it must withhold tax on dividends paid to foreign holders. The standard withholding rate is 30%, reduced by treaty to typically 15% (UK, Germany, France, etc.) and 0% in some specific cases (e.g., qualifying pension funds in certain jurisdictions, Section 892 sovereigns).

A simple example. Apple pays a \$0.25 dividend.

Holder type	Tax withheld	Net received
US resident (no withholding)	\$0.00	\$0.25
UK pension (treaty 0%)	\$0.00	\$0.25
UK individual (treaty 15%)	\$0.0375	\$0.2125
Random foreign holder	\$0.075	\$0.175

On a \$1bn position, the difference between paying 0% and 30% withholding is $\$0.075/\text{share} \times \text{shares} \times \text{dividend rate}$. For Apple alone, that is meaningful. For a multinational holding \$50bn of US equities, the full-year wedge is hundreds of millions.

4.2 4.2 The Yield-Enhancement Trade

The classical dividend trade exploits the wedge: temporarily transfer ownership across the record date from a high-tax holder to a low-tax holder.

The mechanism, in three flavors:

Flavor A: Stock Loan over the Record Date. Foreign holder (high tax) lends the stock to a domestic short seller across the record date. The domestic short seller sells the stock to a US holder. On the record date:

- US holder receives the dividend at 0% withholding (\$0.25 gross).
- Domestic short seller owes the lender a “manufactured dividend” — a substitute payment.
- Substitute payment is *not* a dividend; it is treated as ordinary interest income.
- Critically, in the pre-871(m) world, the substitute payment was paid net of *no* withholding by the borrower — but it was also taxed differently in the lender’s hands.

The arbitrage: the foreign lender historically received the substitute payment without the 30% withholding. The economic dividend was effectively shifted from a 30% tax cost to a 0% tax cost, with the desk capturing the wedge.

Flavor B: The Dividend Repo (DivRepo). The foreign holder enters a repo: sells the stock for cash, agrees to buy back at a fixed forward price after the dividend. The repo counterparty (a US tax-favored entity) holds the stock across the record date and receives the dividend at low tax. The forward repurchase price reflects the dividend, so the foreign holder economically gets back 100% of the gross dividend less only the *desk’s spread*.

Flavor C: Synthetic Forward via Total Return Swap. The foreign holder enters a total-return swap referencing the same underlying stocks. The swap pays the holder the stock return (including dividends) net of a financing rate. The desk hedges by holding the actual stock at a low-tax US entity. Section 871(m), enacted in 2010 and tightened in 2017, was specifically designed to shut down this version.

4.3 4.3 Section 871(m) — The Regulatory Counter

Section 871(m) of the Internal Revenue Code, with regulations finalized in 2017, attempts to disable the dividend-arb wedge by:

1. Treating “dividend equivalents” paid on certain swaps and notional contracts as if they were actual dividends — subject to withholding.
2. Defining “dividend equivalents” broadly: any payment made under a notional principal contract or specified equity-linked instrument referencing a US security that is contingent on or determined by reference to a US-source dividend.

3. Carving out “qualified indices” (broad indices) from the dividend-equivalent treatment.

The practical effect: synthetic equity exposure ($\delta = 1.0$) attracts 30% withholding on dividends as if it were direct ownership. The arbitrage is closed *for delta-one swaps*. But a non-delta-one product (e.g., a deep ITM call option, where the delta is approaching 1.0 but is technically not 1.0) sat in a gray zone for years. The 2017 regulations addressed this with a “delta of 0.8 or higher” rule, but boundary cases remain.

The single largest dividend-arb business model — synthetic ownership at low tax — was therefore largely wound down between 2010 and 2017 in the US. It still exists in pockets of European, Asian, and Latin American markets.

4.4 4.4 Cum-Ex (Germany 2007–2012) — How Far Wrong It Can Go

Cum-cum trades feel almost mundane next to their criminal cousin cum-ex, but the volume that flowed through them in the decade before 2016 was vastly larger and the mental model is more illuminating. Imagine a Dutch pension fund holding €200m in Volkswagen ordinary shares. Every June and October, VW pays a dividend, and every June and October, the German tax authority withholds 26.375% — a hard, cash, no-recovery hit on the Dutch fund. The Dutch fund could file a treaty refund claim, but in 2010 those claims took 18 months and recovered, on average, 80% of the withholding net of accountant fees. So the Dutch fund’s economist looked at the calendar and called Frankfurt: “How would you feel about borrowing my VW shares for two weeks across the record date?” The trader on the other end of the line ran the math in his head — receive the dividend at the bank’s domestic rate, manufacture a substitute payment back to the Dutch fund net of a small spread, hold the position fully collateralized so the balance sheet cost is zero, charge the bank’s compliance team nothing because the trade is in conformity with the BaFin guidance of the era. He quoted a fee. The fund accepted. Multiplied across every German equity dividend paid to every non-treaty foreign holder, the same call took place tens of thousands of times per year. This is what the Frankfurt clearing system looked like, in human terms, before 2016 closed it down.

The infamous **cum-ex** scheme exploited a peculiarity of the German settlement system. Briefly: shares were traded across the dividend record date using short selling and rapid settlement timing, such that *two parties simultaneously claimed the same dividend tax credit* with the German tax authority. The tax authority paid both, effectively refunding tax that had never been collected.

Cum-ex was not “aggressive tax planning” — it was a multi-billion-euro fraud, with tax credits issued for tax that was never paid. Germany has prosecuted bankers, asset managers, and law-firm partners criminally; the total revenue loss to the German treasury is estimated at €10bn–€55bn depending on the timeframe. Several major banks (Maple Bank, Warburg) collapsed or paid massive settlements.

Cum-ex is *not* a trade to study to replicate. It is an example of how the dividend-arb machine, pushed past the legal line, becomes the largest tax fraud in European history.

The closely related “cum-cum” (different mechanism, same intent — temporary transfer to avoid withholding) was widely practiced as recently as 2015 and is now largely shut down by the German legislature.

4.5 4.5 Options-Based Dividend Trades — The American Call Early Exercise

A pure-options dividend play exists in the US and *is* retail-relevant. American-style calls on a stock about to pay a dividend may be optimally exercised early, just before the ex-date. The early-exercise condition:

Optimal early exercise of American call before ex-date if:

$$D > \text{Strike} \times r \times T_{\text{remaining}} + \text{put_value}$$

where D = upcoming dividend, $T_{\text{remaining}}$ is time to expiry from ex-date.

Expressed more cleanly: a deep ITM call should be exercised the day before the ex-date if the dividend exceeds the option’s remaining time value plus the deferral cost on the strike. Most retail call holders fail to exercise — exactly mirroring the put-non-exercise problem from Chapter 5.

The desk’s harvest:

- Desk writes deep ITM calls on dividend-paying stocks days before the ex-date.
- Retail does not exercise.
- Stock goes ex-dividend; call retains its intrinsic value but is now worth less by approximately the dividend amount (because the stock dropped by the dividend amount).
- Desk has effectively captured the dividend.

Per share, the captured value is roughly the dividend minus the (small) excess time value retail holders should not have gifted. On the highest-dividend names (T, MO, KMI, REITs), the per-share capture is on the order of 1–5 cents per quarter — small in isolation but enormous on a \$10bn options book.

4.6 4.5b The Retail-Visible Signature of Institutional Dividend Wedge

The institutional flows described in 4.1 through 4.5 leave a fingerprint on retail-accessible option markets that the perceptive retail trader can read. This subsection describes the fingerprint and how to interpret it.

A retail trader who pulls up the option chain for AT&T (T) — a high-dividend, high-foreign-ownership US name — three days before a dividend ex-date will observe several patterns that look weird without context. First, the deep ITM puts will be priced richer than naive Black-Scholes implies. The market is anticipating that the put holder may exercise early to capture the dividend’s incremental cash, and is pricing in that small probability. Second, the at-the-money calls will be priced cheaper than the equivalent at-the-money puts, again because the market is anticipating the price drop on ex-date and embedding the dividend forecast into the call. Third, the deep ITM calls will be unusually rich — these are the calls that institutional desks may exercise early to capture the dividend, and the market prices in that exercise risk.

These signatures are not accidents. They are the *direct fingerprint* of the institutional dividend wedge flowing through option pricing. Every dollar of the foreign holder’s dividend tax differential, every share-loan-and-recapture across the record date, every conversion arb run by a delta-one desk — all of these flows price into the option chain through put-call parity adjustments. The option market is, in this sense, an auction where the institutional dividend trade is constantly being priced.

The retail trader’s actionable insight: the option pricing on dividend-paying stocks pre-ex-date is *structurally* tilted by these flows. A retail trader buying or selling options on T or VZ in the week before a dividend is not trading vanilla Black-Scholes; they are trading the residual of an institutional flow that happens around them. Pricing options correctly requires acknowledging this — typically by computing the implied dividend (Section 4.7) and comparing it to the announced dividend.

Concretely: if the option chain for T trading at \$20.10 implies a dividend of \$0.32 (via put-call parity) but the announced dividend is \$0.28, the market is *overpricing* the dividend. This often happens when there is rumored special-dividend speculation. A retail trader can position long calls / short puts to bet that the realized dividend will be the announced \$0.28, capturing the 4-cent overpriced wedge. Conversely, if the implied dividend is below the announced, the market may be pricing in a possible cut — a position to short calls / long puts may capture the realized \$0.28 if no cut materializes.

This trade is small per-share but reliable across a large basket of dividend names. A retail user sweeping ~30 dividend-paying names with weekly options around every ex-date can run this systematically with a modest capital base (\$100k+) and earn 5-15% annualized on the deployed capital. The trade is variant of the dividend-arbitrage book run by every Tier 1 desk, expressed in a retail-executable wrapper.

The signature in put-call parity is not the only retail-visible footprint. The **borrow rate on a dividend-paying name shifts predictably around ex-dates**. This shift is, again, a fingerprint of institutional flow. Pre-ex-date, foreign holders are arranging stock loans and term repos to neutralize the withholding hit. The increased lending demand around the ex-date manifests as a temporarily lower borrow rate (more supply hits the market). After the ex-date, the supply withdraws and the rate normalizes. A retail short seller who shorts a dividend name three days before ex-date and covers two days after will systematically pay a *lower-than-normal* rate during the holding window. This is small (5-15 bps in dollars on a typical name) but real, and it represents the retail short benefiting from the institutional cum-cum residual flow that survives in 2024.

4.7 4.6 Worked Numerical Example: A Dividend Capture Trade

Stock XYZ, \$50.00, paying a \$0.75 quarterly dividend. The ex-date is 5 days away. The 30-day call at \$40 strike trades at \$10.20 (with \$0.20 of time value).

Holder analysis (long the call):

- Hold through ex-date: stock drops to \$49.25. Call value falls to \$9.45 (intrinsic) + small time value remaining \$9.50. Loss: $\$10.20 - \$9.50 = \$0.70$.

- Exercise: pay \$40, receive stock worth \$50, sell stock day-of-ex-date at \$50, receive \$0.75 dividend. Gain: $(\$50 - \$40) - \$10.20 + \$0.75 = \$0.55$. Better by $\$0.55 + \$0.70 = \$1.25$. **Should exercise.**

Retail behavior: Most retail does *not* exercise. Empirically (Pool, Stoll, Whaley 2008), about 15-20% of optimal early-exercise opportunities on calls are missed.

Desk's trade: - Desk sells (writes) 1000 contracts of the \$40 call at \$10.20 = \$1.02m credit. - Buys 100,000 shares at \$50.00 = \$5.00m to hedge (delta = 1.0). - Net out-of-pocket: $\$5.00m - \$1.02m = \$3.98m$. - On ex-date with 80% of retail not exercising: - 80,000 shares stay assigned to desk: receive $\$0.75 \times 80,000 = \$60,000$ dividend. - 20,000 shares are exercised against desk: lose $\$0.75 \times 20,000 = \$15,000$ of dividend equivalent (substitute payment). - Net dividend captured: $\$60,000 - (\$15,000 \text{ manufactured dividend tax wedge}) = \$45,000$. - After ex-date, desk closes the option (decayed by approximately the dividend) for ~\$9.45, profit of $\$0.75/\text{contract} \times 1000 = \$75,000$ minus original credit and hedge.

Realized P&L estimate: roughly +\$45,000 to \$75,000 on the trade for one quarter on a single name. Repeated across the entire dividend-paying universe four times per year, this is the dividend-arb business at work.

4.8 4.7 Dividend Forecasting and the Implied Dividend Curve

Options markets price the expected dividend stream into option prices via put-call parity:

$$C - P = S - K \times e^{-rT} - \text{PV}(\text{Dividends to expiry})$$

Solving for PV(Dividends):

$$\text{PV}(\text{Dividends}) = S - K \times e^{-rT} - (C - P)$$

This is a *clean* read on the market's dividend forecast embedded in options prices. Comparing this to the company's announced dividend stream tells the desk:

- **Implied premium:** market expects a higher dividend than company guidance → options are cheap to buy a dividend bull spread.
- **Implied discount:** market expects a cut → options are cheap to bet on dividend cuts.

A famous 2016 example: GE option prices implied a 25% dividend cut by mid-2017. The company maintained its dividend until December 2017, when the cut came in larger than implied. Desks who positioned along the implied dividend curve made money on both sides.

For SPY/SPX, the implied dividend curve is published daily by major dealers and is one of the cleanest macro signals available — it shows the market's expectation of S&P 500 aggregate dividends 1, 2, 3, ..., 10 years forward.

4.9 4.8 Special Dividends and Option Adjustments

A *special* (one-time) dividend triggers an OCC option adjustment. The OCC publishes a memorandum specifying the strike adjustment for all listed options. The standard rule:

- Special dividend > \$0.125 per share triggers strike adjustment.
- Adjusted strike = Old strike – Special dividend amount.
- Number of contracts and contract multiplier may also adjust.

This catches retail constantly. A trader holds 10 contracts of XYZ \$50 calls. The company announces a \$5 special dividend. On the ex-date:

- The strike adjusts from \$50 to \$45.
- The “deliverable” remains 100 shares per contract.
- Stock drops by \$5 to a new equilibrium price reflecting the dividend.
- The option's intrinsic value is *roughly preserved* (small frictions).

Retail traders frequently misread the adjustment as “the broker stole my profit” — they see the new strike at \$45 but don't realize the deliverable is unchanged. Confusion creates exit liquidity for the desk.

4.10 4.9 The Dividend Future and the SX5E Dividend Index

European markets have something the US lacks: an exchange-listed dividend futures contract. The Eurex SX5E (Euro Stoxx 50) Dividend Future contracts allow direct, exchange-traded exposure to *next year's expected dividends* on the Euro Stoxx 50 index.

The contract specifications: - Each contract represents the cash dividends paid by Euro Stoxx 50 constituents during a specific calendar year. - Multiplier: €100 per index point. - Expiries: December of years 1 through 10 forward.

The implied dividend curve from the futures market is one of the cleanest macro signals in finance. It shows the market's expectation of European corporate dividend resilience over the next decade. During the 2020 COVID crash, this curve collapsed dramatically as European companies cut dividends en masse (BCE, Banco Santander, etc.). The 2020 curve traded as low as 60% below the 2019 curve's level for the same forward year.

A delta-one desk uses these contracts in three ways:

1. **Hedge.** A bank holding a long basket of Euro Stoxx 50 names with a known dividend exposure can sell the corresponding dividend future to lock in the cash dividend value.
2. **Trade dividend sentiment.** A fund believing European banks will resume dividends faster than implied can buy the future.
3. **Spread trades.** Buy the 2-year, sell the 1-year, betting on the recovery slope.

The US has no equivalent exchange-listed instrument. The implied dividend curve in the US is constructed from option prices via put-call parity (see Section 4.7). This is a less liquid, less transparent way to express dividend views.

4.11 4.10 The Dividend Forecast vs Realization Spread

Empirically, on the SX5E, the implied dividend (1-year forward) has *consistently traded below* the realized dividend over rolling windows since 2003. The structural reason: dividend futures are typically sold by *natural hedgers* (delta-one desks hedging long basket positions and long convert positions). Buyers are *speculators* taking the other side. The natural sell pressure depresses the implied dividend below fair.

The spread historically averaged 5-15% — meaning a buyer of the 1-year dividend future earned that much in expected value over the actual realized dividend. This is a *risk premium*, not a free lunch — buyers bear the risk of a dividend collapse (as happened in 2020). But across long windows, the buy-and-hold dividend-future trade has been profitable.

This trade is not retail-executable. The Eurex SX5E dividend futures require a futures account with European clearing access; minimums are typically \$100k+ and the contract sizes (€10k+ per contract) make small-scale execution awkward.

4.12 4.11 Worked Example: The Pre-Ex-Date Term Repo

A Tier 1 desk's structured-product team is asked by a London-based client (a non-treaty foreign holder of \$200m in MSFT, S&P 500 weighting) for a "dividend yield enhancement" trade. MSFT is about to pay a \$0.83 quarterly dividend. The client would lose 30% to US withholding (\$0.249/share = \$498k on the position).

The trade:

Trade structure: 1-week term repo across the ex-date.

Day 0 (3 days before ex-date):

Client (foreign holder) repo'd MSFT shares to bank.

Bank pays client $\$200\text{m} \times (\text{MSFT price} - \$0.83 \times 0.30)$ cash today.

Repurchase price (1 week forward): same notional, agreed.

Effectively: bank "borrows" the shares from the client for 1 week, posting cash collateral.

Day 0 to ex-date:

Bank holds the shares. On record date (Day +2), bank receives the

\$0.83 dividend gross at the bank's preferred tax structure.

Day 0 + 1 week:

Client repurchases shares at the agreed forward price.

The forward price is set such that the client effectively receives:

Net dividend received via repo wedge = approx $\$0.83 \times (1 - 0.10) = \0.747

(where 10% is the desk's spread / friction)

Versus the alternative of holding through ex-date:

Net dividend with 30% withholding = $\$0.83 \times 0.70 = \0.581

Improvement to client: $\$0.747 - \$0.581 = \$0.166/\text{share}$

= \$33,200 on \$200m / \$50 share = 4M shares

(Calculation simplified - exact wedge depends
on tax treaty position.)

Bank's economics:

Captures the $\$0.83 - \$0.747 = \$0.083/\text{share spread} \times 4\text{m} = \332k .

Less hedging cost (delta hedge of price risk during the week): ~\$50k.

Net: ~\$280k for one week of work.

Plus: reuses bank balance sheet for the 1-week period at zero cost
since the trade is fully collateralized.

This is one of countless variations on the dividend wedge trade. Across the major banks, billions in similar trades flow through annually for foreign clients.

The retail user *cannot* access this trade structure. The ISDA agreements, the settlement infrastructure, and the relationships with foreign tax-advantaged holders do not exist at retail. The closest retail analog is the dividend capture with put-spread hedge (Chapter 4.5/4.6 above), which captures only a fraction of the institutional wedge.

4.13 4.12 Dividend-Adjusted Volatility and the Implied Dividend Volatility Surface

Sophisticated dealers also model the *volatility* of dividends, not just their expected level. A dividend future has its own implied volatility — the market's expectation of how dividend forecasts will evolve. During 2020, the realized volatility of 1-year-forward Euro Stoxx 50 dividends approached 40% — many times historical norms.

For a holder of long-dated equity options, dividend volatility is a hidden risk. A 5-year LEAP on AAPL is sensitive not just to AAPL stock vol but also to AAPL dividend vol — if AAPL cuts its dividend during a recession, the long call's value is *higher* than the deterministic-dividend pricing implies, while a long put's value is *lower*. Dealer pricing of long-dated options accounts for this; retail option pricing models typically do not.

4.14 4.12b The Day in the Life of the Dividend Desk

To make the abstract concrete, here is what an actual trading day looks like on a major bank's dividend-arb pod during a busy ex-dividend week. The setting: late April 2024, the heart of US dividend-paying earnings season. Dozens of large-caps go ex-dividend within a single calendar week.

Tuesday morning, 06:30 ET. The dividend pod's daily morning meeting at JPMorgan's Chase Tower in Manhattan covers the week's calendar. There are 47 ex-dividends across the S&P 500 this week, of which 18 are >\$0.50/share quarterly dividends — the names that move the dividend-arb book the most. Apple is paying \$0.24, Microsoft \$0.75, JPMorgan \$1.15, Pfizer \$0.42. The pod's structurer has prepared a forecast of expected client demand: the foreign-holder client base alone is estimated to want \$4.2bn of yield enhancement this week, principally on the high-dividend names.

09:00 ET. The first call comes in from a Singapore-based sovereign wealth fund. They hold \$300m of Microsoft and want a yield enhancement structure for the upcoming MSFT dividend. The structurer prices a five-day term repo wedged across the record date, pricing in the foreign-tax-credit differential between Singapore (15% treaty rate to US) and a US-domiciled tax-favorable counterparty. Quote: bank captures 8 bps of the gross dividend, client receives

net 92% of gross vs the 85% they'd receive with direct ownership through the ex-date. Trade booked. The structurer hands off to the financing pod for hedging (delta hedge of MSFT price exposure during the 5-day window) and to operations for tri-party setup.

11:30 ET. A different call: a Brazilian asset manager (\$2bn AUM in US large-caps) wants a similar structure on JPM and Pfizer. But Brazil's treaty position with the US is 15% on dividends to Brazilian corporate holders, and the trade is sensitive to whether the ultimate beneficial owner is a corporate or an individual. The structurer asks for documentation; the back-and-forth takes two hours. By 14:30 the trade is booked at a slightly worse rate than the Singapore one because the documentation friction is real.

13:00 ET. Internal call with the bank's tax department. The 871(m) classification of an unusual structure (a partial-delta swap referencing a basket including PFE and GILD with a derivative kicker) is being audited. The tax team has questions about whether the swap qualifies as 871(m)-exempt or not. This is the kind of low-level compliance work that occupies a substantial fraction of the dividend pod's time. Eventually the tax team confirms the classification; the trade is approved.

15:30 ET. Pre-close, the structurer reviews the day's bookings: \$850m of new yield enhancement structures booked, weighted-average wedge captured by the bank at 9.2 bps of gross dividend. Translation: roughly \$0.85m of revenue from this single day's structuring activity. The week's total is on pace for \$4-5m, which is normal for a busy ex-dividend week.

The dividend trade flow is not glamorous. There are no screaming traders, no dramatic risk events. It is patient, methodical, document-heavy work that generates steady billion-dollar P&L lines for the major banks. The clients are sophisticated. The trades are largely standardized. The compliance load is enormous. But the franchise survives because the underlying tax differentials persist, and as long as they do, the institutional flow will continue.

4.15 4.13 Cum-Cum vs Cum-Ex — The Two Faces of Dividend Stripping

Sections 4.1-4.4 introduced the cross-border tax wedge and named cum-ex as one historical extreme. Cum-cum is the related but legally distinct cousin that operated at vastly larger scale and is still partially with us. The two trades sound similar — both turn on whether a holder is “cum dividend” (entitled to it) or “ex dividend” (not entitled). They differ in mechanism, scale, and legality.

4.15.1 4.13.1 The naming convention

cum dividend = with the right to the upcoming dividend
ex dividend = after the dividend has been paid (no longer entitled)

A trade settled cum-cum means: bought cum-dividend, sold cum-dividend (i.e., a stock-loan-like round trip across the record date with both parties entitled). A trade settled cum-ex means: bought cum-dividend, sold ex-dividend (the seller's claim was extinguished but the buyer's was created — and in fraudulent variants, a *third* party also claimed entitlement based on the same underlying share).

4.15.2 4.13.2 Cum-cum: the legal-if-aggressive transfer

Mechanism: a non-resident holder (paying, say, 30% German withholding tax on a Frankfurt-listed dividend) “parks” the stock with a German resident holder (who pays a far lower or zero domestic rate after foreign-tax-credit offsets) over the record date. The two parties book a stock loan or a repo. The German resident receives the dividend at the favorable rate; net economics are split with the non-resident via fees on the loan.

Non-resident holder (NL pension fund)

Stock loan T-3 to T+5 across record date

German resident counterparty (DE bank trading book)

Dividend received at near-zero effective rate
(foreign tax credit offset internally)

Non-resident receives manufactured dividend = $0.93 \times$ gross dividend
(vs $0.74 \times$ gross if held directly through ex-date - the DE WHT is 26.375%)

Wedge captured: ~19 percentage points of the gross dividend
Split: typically 60% to non-resident, 40% to bank

Pre-2016 this trade ran across the entire EU and Asian-Pacific cross-border equity book. Estimated annual loss to German tax authorities from cum-cum alone: €5-7 billion through the early-2010s.

4.15.3 4.13.3 What killed cum-cum (mostly)

The German Investment Tax Reform Act of 2016 (Investmentsteuerreformgesetz) imposed a 45-day-around-record-date holding period: any stock-loan return executed within ± 45 days of an ex-date no longer qualified for the foreign-tax-credit offset that made the trade economic. Section 36a of the German Income Tax Act codified this for direct holders. Together they reduced cum-cum volume by 90%+ but did not extinguish it entirely — large institutional holders with multi-quarter horizons still execute the trade legally.

Other jurisdictions followed: the Netherlands introduced a similar minimum-holding-period rule in 2018; Italy and Belgium added analogous provisions in 2020. The OECD’s Multilateral Instrument (MLI) Article 8 explicitly targets “transactions structured to avoid dividend withholding.”

4.15.4 4.13.4 Cum-ex: the criminal escalation

The fraud structure layers a *short sale across the record date* on top of the cum-cum mechanic, exploiting the T+2 settlement cycle.

- Step 1: Short seller A sells stock cum-dividend on T-1.
- Step 2: Long buyer B buys stock cum-dividend from A.
- Step 3: Long buyer B receives a dividend credit at record date.
- Step 4: Short seller A locates and delivers (borrowed) stock to B post-record.
- Step 5: The depositary issues a tax certificate to B.
- Step 6: A separately claims the original dividend from the lender, who also held a tax certificate.
- Step 7: TWO tax certificates have been issued for ONE dividend payment.
Both holders claim a refund from the German tax authority.

In the most aggressive variants, three or four parties claimed the same dividend through chained short-sale-and-stock-loan structures. The fraud was enabled by a quirk of German tax law: pre-2012, the depositary auto-issued tax certificates to whoever held the stock at any point during the settlement window, with no central reconciliation.

Estimated fraud volume in Germany alone: €36-55 billion across 2007-2012. Pan-EU including Denmark, France, Belgium, Italy, Austria: €150 billion+. Active prosecutions continue — the *Hanno Berger* trial (Bonn 2024) and *Sanjay Shah* extradition (Copenhagen 2023) are landmark cases. Settlements with major banks (Warburg, M.M. Warburg, MüKi, plus US and Japanese global banks) total tens of billions in clawbacks.

4.15.5 4.13.5 Why cum-ex was structurally possible only in Germany 2007-2012

Three preconditions had to hold: 1. *Settlement-cycle arbitrage*: T+2 settlement created a two-day window where ownership was ambiguous. 2. *Depositary issuance of tax certificates*: the bank holding the stock auto-issued the certificate without central registry. 3. *No reverse-charge mechanism*: tax was withheld at source from the dividend payer, not collected from the certificate-claimer.

The 2012 OGAW IV reform centralized tax certificate issuance through a single registry; the 2016/2018 reforms added reverse-charge logic. By 2020 the fraud was no longer mechanically executable in any major EU jurisdiction.

4.16 4.14 Australian Franking Credits — A Separate Tax Architecture

Australia’s tax system uses *imputation*: when an Australian company pays corporate tax on its profits, it generates “franking credits” that are attached to subsequent dividend payments. Domestic holders use the credits to offset their personal tax on the dividend; foreign holders cannot use them.

4.16.1 4.14.1 The mechanic

A 30%-corporate-tax-paying Australian company earns A\$100 profit.

Pays A\$30 corporate tax to ATO.

Pays A\$70 cash dividend, attaching A\$30 franking credit (fully franked).

Domestic holder receives:

Cash dividend	A\$70.00
Franking credit	A\$30.00
Gross income	A\$100.00 (assessable for income tax)
Tax on A\$100 @ marginal rate (e.g. 32.5%)	= A\$32.50
Less franking credit	= A\$30.00
Net additional tax	= A\$2.50
Net cash retained	= A\$67.50

Foreign holder receives:

Cash dividend	A\$70.00
Franking credit	A\$30.00 (worthless)
WHT on A\$70 (varies by treaty, 0-15%) e.g. A\$10.50 at 15%	
Net cash retained	= A\$59.50

The wedge: A\$67.50 – A\$59.50 = A\$8.00 per A\$100 of original profit, or roughly 11% of the dividend. On large sustained dividend streams, this is structurally enormous.

4.16.2 4.14.2 The trade structure

Foreign holder (US pension fund, global ETF)

Stock loan T-N to T+M across record date

Australian super fund (15% concessional rate, full franking benefit)

Receives dividend + franking credit
Refunds excess franking credit from ATO ("excess imputation")

Foreign holder receives manufactured dividend = approx $0.93 \times$ gross
versus $0.85 \times$ gross holding directly

Pre-2003 this trade was openly executed by every prime broker desk in Sydney. The Aus super fund (typically a domestic pension or non-profit) earned the imputation credit, the prime broker and foreign holder split the gain, and the ATO eventually refunded the franking credit *as cash* to the super fund (because super funds pay tax at 15% but receive the full 30% franking credit — the difference is refundable).

4.16.3 4.14.3 What killed (most of) the franking-credit trade

The ATO introduced the *45-day-at-risk* holding period in 2003 (Income Tax Assessment Act 1936 Section 160APHO and successor provisions). To claim franking credits, the holder must be at risk of price movement on the underlying stock for at least 45 days within ± 90 days of the record date. Stock-loan-like structures fail this test because the loan transfers price risk back to the original holder.

Modern variants survive in two narrow forms: 1. *Long-horizon foreign-via-Australian-trust* structures — the trust legitimately holds the stock for >45 days before/after the record date, taking real price risk in exchange for capturing the credit. 2. *Equity total-return swaps with Australian-domiciled hedge counterparties* — the swap form fragments price exposure across multiple counterparties, sometimes satisfying the at-risk test on a notional basis.

Total Australian franking-credit refunds claimed annually run A\$5-7 billion. Of this, an estimated A\$1-2 billion accrues to structures that economically pass the credit through to foreign beneficial owners.

4.16.4 4.14.4 The 2018 Labor Party proposal

Bill Shorten's Labor Party proposed eliminating excess-franking-credit refunds in 2018 — meaning super funds would still use credits to offset their own 15% tax but would not receive cash refunds for the unused balance. Estimated annual revenue impact: A\$5-6 billion. The proposal was widely seen as a key driver of Labor's 2019 election loss; Australian retiree retail investors were the loudest opponents because their self-managed-super-fund (SMSF) returns were materially dependent on the refund. The policy was abandoned.

4.17 4.15 EU Treaty Shopping and the ATAD Counter

A second class of tax-driven trades exploits *bilateral tax treaties* rather than holding-period mechanics. The structures are:

4.17.1 4.15.1 The conduit holding-company structure

A US-incorporated multinational paying a dividend to its Mexican corporate parent faces a 30% US withholding (no US-Mexico treaty rate exists for corporate-to-corporate dividends without specific treaty status). To reduce this:

US OpCo

Dividend

Netherlands HoldCo (treaty rate to US: 0% on substantial holdings under
US-NL treaty Article 10 with proper qualification)

Onward distribution

Mexican parent (NL-MX treaty, 5-10% withholding)

Total leakage: ~5-10% effective, vs the 30% direct route.

The Netherlands was the dominant conduit for US-bound dividends through the 2010s because of (a) the 0% US-NL substantial-holding rate, (b) the NL participation exemption that exempts onward dividends from NL tax, and (c) the broad NL treaty network. Ireland and Luxembourg play similar roles for different bilateral pairs.

4.17.2 4.15.2 The “Double Irish with a Dutch Sandwich”

The structure that minimized Apple's, Google's, and Facebook's effective tax rates on European earnings to 1-3% in 2010-2015. Mechanism:

Operating profits (EU sales)

Irish OpCo #1 (low Irish corporate rate ~12.5%)

Royalty payment for IP

Dutch HoldCo (0% Dutch tax on royalties to non-EU)

Onward royalty

Irish HoldCo #2 (Bermuda tax-resident - Ireland's pre-2015 rules
allowed this, taxing only on management-and-control
location, which was Bermuda)

Bermuda (0% tax)

Effective tax: 1-3% on European profits, vs the 25-30% statutory rate.

Closed by Ireland's 2015 reform mandating that Irish-incorporated companies be Irish-tax-resident. Phase-out was complete by January 2021.

4.17.3 4.15.3 ATAD I and II — the EU's structural counter

The Anti-Tax-Avoidance Directive (ATAD I, in force 2019; ATAD II, in force 2020) introduced five EU-wide rules:

1. *Interest limitation* — caps the deductibility of interest expense at 30% of EBITDA. Targets debt-pushdown structures used to strip earnings into low-tax jurisdictions.
2. *Exit taxation* — taxes built-in gains when a company moves tax residence out of the EU. Closes the “move IP to Bermuda” exit.
3. *General Anti-Abuse Rule (GAAR)* — disregards arrangements whose principal purpose is tax avoidance.
4. *Controlled Foreign Company (CFC) rules* — taxes the parent on undistributed profits of low-tax foreign subsidiaries.
5. *Hybrid mismatch rules (ATAD II)* — eliminates the double-deduction or deduction/no-inclusion outcomes of cross-border instruments treated differently in two jurisdictions.

The OECD’s Multilateral Instrument (MLI), in force across 99 jurisdictions as of 2024, adds the *Principal Purpose Test* (PPT): treaty benefits are denied for any arrangement whose principal purpose was obtaining the benefit. The PPT is a de facto kill-switch on conduit structures whose only economic purpose was treaty arbitrage.

Empirical effect (Tørsløv, Wier, Zucman 2023): OECD profit-shifting via conduits fell from peak ~40% of multinational profits in 2015 to ~25% by 2022. Still large, but materially smaller.

4.18 4.16 The 2024 EU FASTER Directive

EU Directive 2025/50 (the “FASTER” Directive — *Faster and Safer Tax Excess Relief*) adopted in May 2024 and to take full effect by January 1, 2030, restructures EU withholding-tax refunds. Key provisions:

1. *Common digital tax residence certificate (eTRC)* — a single EU-wide standard for proving treaty-residence eligibility, replacing the 27 distinct national forms.
2. *Two fast-track refund procedures* — “relief at source” and “quick refund” — must be made available by all member states. Maximum 25 days from claim to refund (vs current avg 12-18 months in Italy, France, Germany).
3. *Mandatory registry of certified financial intermediaries* — only registered brokers can submit fast-track refund claims, creating a reputational accountability layer.
4. *Penalties for over-refund* — up to 100% of the over-refund amount, indemnifiable to the certified intermediary.

Predicted effects: - Reduces friction-based arbitrage that protected legitimate trades (most retail/non-fraudulent foreign holders historically forfeited withholding refunds because the cost-to-claim exceeded the refund). FASTER recovers this for them. - Compresses the wedge available to dividend-strip structures because the legitimate refund path becomes faster and cheaper. - Eliminates the “refund timing arbitrage” that some hedge funds operated — claiming a refund 18 months after the trade and discounting that future cash flow back to today. - Net long-term effect on dividend-arb volume: estimated 30-50% reduction in cross-border refund-driven trades, though the absolute number still in the tens of billions annually.

The directive has critics. The European Tax Adviser Federation argues that the certification requirement creates a competitive moat for the largest custodians (BNP Paribas, Citi, BNY Mellon, State Street) and disadvantages smaller asset managers. The compromise: smaller intermediaries can rely on a “passport” system to use a certified peer’s infrastructure.

4.19 4.17 The Tax-vs-Structural Split — Retail Accessibility Matrix

The taxonomy of dividend-related trades resolves into two clean categories. The first depends on a *tax differential* that the retail trader does not have access to. The second depends on *contract mechanics* or *retail-flow patterns* and is theoretically open to anyone with a brokerage account.

4.19.1 4.17.1 Tax-driven trades (institutional only)

Trade	Edge source	Status today	Retail via
Yield enhancement (manufactured dividend / repo wedge)	30% US WHT vs 0% domestic	Active; 871(m) limited	NO
Cum-cum (DE/EU)	Non-resident parking with resident	Largely killed post-2016 (~90% volume); residual in long-horizon variants	NO
Cum-ex	Multiple refund claims on one	CRIMINAL; active prose-	NO (illega

871(m) swap workarounds	dividend via T+2 chain Dividend equivalents not classed as US-source dividends	cutions; mostly historic Closed for delta 0.8; low-delta still works	NO
Australian franking-credit harvesting	Imputation credits worthless to foreign holders	45-day-at-risk rule killed easy version; super-fund variants live	NO
EU treaty shopping (NL/LUX/IRL conduits)	Bilateral treaty rate differentials	ATAD I/II + MLI PPT closed most conduits	NO
2024 FASTER Directive trades	Refund-process arbitrage (timing, friction, multi-jurisdiction)	In force 2030; will shrink ~30-50% of remaining volume	NO

The retail user cannot access ANY of these. The reason is structural: each requires a specific tax wrapper (NRA status with treaty qualification, super-fund or trust domicile, derivative counterparty with delta-one designation, certified financial intermediary status under FASTER) that a US-resident retail trader simply does not have and cannot acquire by signing more brokerage paperwork.

The trades exist in this guide for *understanding*: they explain why dividend pricing in option markets has the structure it does, why deep-pocketed foreign holders often appear willing to lend stock cheaply across record dates, and why dividend-future markets in Europe behave the way they do. Without understanding these flows, a retail trader will misread the signals they create.

4.19.2 4.17.2 Non-tax / structural trades (theoretically retail-accessible)

Trade	Edge source	Edge size	Retail viable?
Special-dividend OCC rare events, adjustment arb	Strikes adjusted only above ~\$0.125/sh threshold; mispriced near boundary	Highly variable (\$0-\$2/share on rare events)	Marginal - requires watching OCC adjustment notices
Dividend-capture (buy cum, sell ex) Put-steal / interest-on-implemented in strike arbitrage	Stock falls < dividend amount on ex-date for retail-heavy DRIP names Retail fails to early-exercise deep-ITM American puts	5-15bps gross Annualized ~10-25% on capital deployed	YES but cost stack k it for typical retail YES - this repo, see [strategies/put_steal]
Dividend-swap implied vs futures access not realized	Eurex SX5E div futures price implied div stream below realized	5-15% on rolling windows	NO - retail-typical
Conversion arb (long stock combo orders not + short call + long put, put-call parity)	Put-call parity break around ex-date	Sub-10bps	NO - atomically fillable Robinhood; needs IBK
Index-rebalance dividend but edge below flow trades	Tracker-fund mechanical buying around index dividend dates	1-5bps	YES - retail commission structure

Of all trades discussed in Chapter 4, **only two are economically viable for retail**:

1. **Put-steal**, already implemented at [strategies/put_steal.py](#). The mathematics, the entry signal, the exit logic, and the historical edge are detailed in Chapter 5.
2. **Special-dividend OCC adjustment plays**, but only opportunistically — these events arrive 3-8 times per year across the entire US listed market and require active monitoring of OCC adjustment memoranda.

Every other trade in this chapter is described to *educate* the reader about the institutional flows that shape the markets they trade in. A retail trader who can read a dividend-future curve, distinguish a special from a regular dividend, and understand what an 871(m) swap is will price options and dividend-paying-stock positions more correctly than one who cannot — even if they never execute a single one of these institutional trades.

4.20 4.18 Conversion Arbitrage Around Ex-Date

A *conversion* (also called “synthetic short stock plus long stock”) combines:

Long 100 shares of stock at S
Short 1 call at strike K
Long 1 put at strike K

By put-call parity, the conversion has zero economic equity exposure: a perfectly synthetic riskless position whose only return source is the cost of carry on the financed stock leg. The fair value of the structure on a non-dividend-paying stock is $S - K e^{-rT}$.

When the underlying pays a dividend D over the holding period, parity becomes:

$$S - D e^{-r \cdot t_D} = (\text{call strike value at } K) - \text{put} + S$$

where t_D is the time to the ex-date. The conversion holder *receives* the dividend on the long stock leg, so the cost of carrying the position is reduced by D . This creates a small but persistent dislocation around ex-dates.

4.20.1 4.18.1 The retail-relevant case

If a retail trader holds 100 shares plus a covered call (long stock + short call), the position is *half a conversion*. Adding a long put at the same strike completes it. On the ex-date:

Pre-ex-date:

Long stock	S
Short call (K)	$-\text{call}(S, K, T, r)$
Long put (K)	$+\text{put}(S, K, T, r)$
Net invested	$K e^{-rT}$ (parity)

Ex-date (ex-dividend):

Long stock	$S - D$	
Short call	$-\text{call}(S - D, K, T, r)$	[option market re-prices]
Long put	$+\text{put}(S - D, K, T, r)$	[option market re-prices]

The dividend D is received by the long-stock leg as cash on payment date. The conversion's NAV is *unchanged* if the option market re-prices consistently with put-call parity. In practice, retail-flow-driven mispricings around the ex-date (small calls bid up by retail dividend-capture flows; small puts underpriced by traders rolling out of dividend-paying names) create occasional small windows where the structure mispriced by 5-15bps.

These trades are not retail-economically-viable in isolation ($15\text{bps} \times \$10\text{k position} = \15 less commissions and bid/ask). They are retail-relevant for *position management* — a covered-call writer rolling into ex-date should understand whether the call is mispriced relative to parity and whether a small adjustment is warranted.

4.20.2 4.18.2 The institutional case

Delta-one desks run conversions and reversals (the symmetric counterpart) as part of their market-making book. They earn the bid-ask on each leg, the net dividend yield enhancement, and the option-adjustment edge around special dividends. Across a \$50-100bn book, the aggregate edge on conversion-and-reversal trading is typically 2-5bps annualized on capital, which compounds to material P&L because the positions are essentially riskless and the leverage is $\sim 10:1$.

The trade is one of the lowest-risk, lowest-edge trades on a delta-one desk's book. It functions as the “bedrock yield” — small but steady — that supports the desk's higher-risk dividend-arb book on top.

4.21 4.19 References

- Fink, A. and Tomio, D. (2024), “The Dividend Disconnect”, *Journal of Financial Economics*.
- Pool, V.K., Stoll, H.R., and Whaley, R.E. (2008), “Failure to Exercise Call Options: An Anomaly and a Trading Game”, *Journal of Financial Markets* 11(1): 1-35.

- Buetow, G.W. and Sellers, M.D. (2002), “Synthetic Equivalence and the Special Dividend Anomaly”, *Journal of Derivatives*.
- Henry, T. and Koski, J. (2017), “Ex-Dividend Profitability and Institutional Trading Skill”, *Journal of Finance* 72(1): 461-494.
- Manley, R. and Mueller-Glissmann, C. (2008), “The Market for Dividends and Related Investment Strategies”, *Financial Analysts Journal* 64(3): 17-29.
- Spiess, K. and Affleck-Graves, J. (1999), “The Long-Run Performance of Stock Returns Following Debt Offerings”, *Journal of Financial Economics* 54(1): 45-73.
- Cooper, I.A. and Mello, A.S. (2006), “Corporate Hedging: The Relevance of Contract Specifications and Banking Relationships”, *Review of Finance* 10(2): 231-267.

5 Chapter 5 — Put-Steal and Early-Exercise Interest Arbitrage

This chapter is the deepest dive in the guide because the underlying mathematics are simple, the institutional mechanism is precise, and the strategy has a *retail-executable variant* implemented in this repo at [strategies/put_steal.py](#).

The chapter’s structure: it begins with the textbook setup (American vs European puts), but spends the bulk of its pages on intuition. It walks through the put-call parity derivation of the Net Interest Income formula on the page rather than citing it. It traces the history of the trade across three rate regimes — ZIRP, normalization, peak rates — with worked numerical examples for each. It explains, drawing on the behavioral-finance literature, exactly *why* retail systematically fails to exercise even when the math screams that they should. It walks through this repo’s [strategies/put_steal.py](#) implementation line by line, connecting the production code to the academic literature it references. It cites and explains four core papers — Diz-Finucane (1993), Pool-Stoll-Whaley (2008), Barraclough-Whaley (2011), and Battalio-Schultz (2006) — describing what each paper contributed to the modern understanding of the trade. Finally, it lists the failure modes and the rules of thumb that a retail quant can apply on TastyTrade or IBKR right now.

The reader should leave this chapter able to: (a) compute NII for any given put and rate environment from first principles, without reaching for a textbook; (b) explain why the trade has been profitable for institutional desks since the 1980s and why the profitability has scaled with rates; (c) identify when a retail position should be early-exercised; (d) understand the structure of the bull put credit spread variant in this repo and why it is a defined-risk version of the institutional arb; and (e) recognize the trade’s failure modes and how to hedge against them. This is more material than most retail traders ever encounter on a single topic. The compensation is that the topic is *the* clean retail edge in equity options and is worth understanding deeply.

5.1 The Setup — American vs European Puts

A European put can only be exercised at expiry. An American put can be exercised on any business day before expiry. For a put, early exercise can be optimal *even before expiry* because exercising the put delivers the strike X in cash today, which can be invested at the risk-free rate r over the remaining time T .

The decision rule, derived from the put-call parity bound, is:

```
Exercise immediately if:
    (X - S) > P_european(S, X, T, r, ) + early_premium
which simplifies (using put-call parity) to:
    NII = X(1 - e^{-rT}) - C(S, X, T, r, ) > 0
```

Here C is the *Black-Scholes call* price at the same strike. NII is the “Net Interest Income” — the value of exercising now versus continuing to hold. The expression is exact for non-dividend-paying stocks.

The intuition: when you exercise a deep ITM put, you give up the right to wait. The right to wait, expressed in continuous-time call/put parity, is exactly worth a European call with the same parameters. So:

- Cash benefit of exercising now: $X(1 - e^{-rT})$. This is the present value of interest you’d earn on X over T .
- Cost of exercising now: the call value C — the option to wait.
- Net: NII.

When $NII > 0$, exercise is strictly optimal. When the underlying is deep enough ITM ($S \ll X$) and rates are positive, NII is reliably positive.

5.2 5.1b Deriving NII From First Principles

The NII formula above is correct but presented without derivation. A reader who has not seen the put-call parity argument before will not understand why it is true. The derivation is short and worth doing on the page, because each step illuminates a different piece of the trade's intuition.

Start with continuous-time put-call parity, which holds for European options on a non-dividend-paying stock:

$$C(S, X, T) - P(S, X, T) = S - X * e^{-rT} \quad \dots(1)$$

Equation (1) says: the difference between a European call's value and a European put's value is the spot price minus the present value of the strike. The intuition is that holding a long call and a short put together produces a payoff at expiry of exactly $(S_T - X)$, which is the same payoff as buying the stock today and borrowing X (to be repaid as X at expiry, present-valued at e^{-rT} today). So the call-minus-put position must cost the same as the stock-minus-PV-strike position, by no-arbitrage.

Rearrange (1) to isolate the put price:

$$P(S, X, T) = X * e^{-rT} - S + C(S, X, T) \quad \dots(2)$$

Now consider the value of an American put. The American put has the additional right of early exercise. If exercised today, the holder receives $(X - S)$ immediately. If not exercised, the holder retains the European put value $P(S, X, T)$. The American put price $A_P(S, X, T)$ satisfies:

$$A_P(S, X, T) \geq \max(X - S, P(S, X, T)) \quad \dots(3)$$

The interesting question for our trade: when is $(X - S)$ — the immediate-exercise value — strictly greater than the European put value $P(S, X, T)$? Substituting equation (2):

$$\begin{aligned} (X - S) - P(S, X, T) &= (X - S) - (X * e^{-rT} - S + C(S, X, T)) \\ &= X - S - X * e^{-rT} + S - C(S, X, T) \\ &= X - X * e^{-rT} - C(S, X, T) \\ &= X * (1 - e^{-rT}) - C(S, X, T) \\ &= NII \end{aligned} \quad \dots(4)$$

There it is. The Net Interest Income is exactly the gap between the immediate-exercise value $(X - S)$ and the value of waiting (the European put). When $NII > 0$, exercising now is strictly better than holding the European-equivalent option.

Why does the formula take this particular shape? The first term, $X * (1 - e^{-rT})$, is the present value of the *interest* you would earn on X if you exercised now (received X in cash today and invested it at the risk-free rate r). To see this: if you hold X today, by time T you will have $X * e^{rT}$. The amount above the original X is $X * (e^{rT} - 1)$. $X * r * T$ for small rT . The present value of this “extra” interest, discounted back to today at rate r , is approximately $X * r * T * e^{-rT}$, which Taylor-expands to $X * (1 - e^{-rT})$ for small rT . So this term measures the dollar interest gain from accelerating the receipt of X .

The second term, $C(S, X, T)$, is the European call value at the same parameters. This is the option-theoretic value of *waiting* — the right to keep the put alive in case the stock moves further. By put-call parity, the European call's value is exactly the option-value embedded in the European put once you strip out the immediate intrinsic. So when you exercise the American put early, you give up exactly $C(S, X, T)$ of waiting-value in exchange for $X * (1 - e^{-rT})$ of accelerated interest.

NII is the net. When the rate is high ($X * (1 - e^{-rT})$ is large) and the stock is deep ITM (the call value C is small, because the call is far out-of-the-money), NII is positive. When the rate is low (interest-acceleration is small) or the stock is near the strike (call value is large, because there's still meaningful chance of recovery), NII is negative or zero.

This is a complete derivation in five lines of math. The reader should pause and verify each step before moving on. The trade described in this chapter rests entirely on equation (4); every numerical example in subsequent sections uses it.

5.3 5.1c A Brief History of the Trade

The put-steal is not new. It has been understood by academic finance since the 1970s and operationalized by professional market makers since at least the 1980s.

The earliest formal treatment was Roll (1977), who derived the optimal early-exercise boundary for an American option on a dividend-paying stock. Merton (1973) had earlier shown that, for non-dividend-paying stocks, early exercise of an American *call* is never optimal. The asymmetry — early exercise on the put can be optimal but not on the call — is a consequence of the fact that puts benefit from accelerating the receipt of cash (the strike), while calls would prefer to defer the payment of the strike.

The 1980s saw the first systematic exploitation of the trade by the major equity-options market makers — particularly Spear Leeds & Kellogg (which later became part of Goldman) and Susquehanna International Group. These firms ran specialist books on the CBOE and the AMEX with thousands of open option positions. As part of their daily routine, they would monitor each deep-ITM American put for early-exercise opportunities. When NII went positive, they would either exercise their long puts (capturing the NII directly) or, more commonly, identify positions where they were short and bet that the long counterparty would *not* exercise (capturing the NII by writing puts at slightly less than parity).

The 1990s brought the Diz-Finucane (1993) paper that put a formal framework around what specialists were already doing. Studying S&P 100 index options 1985-1989, Diz and Finucane found that puts were systematically not exercised when they should have been, and that the unrealized NII summed across the market into millions of dollars per year. They were the first to use the term “irrational” to describe the failure-to-exercise pattern, and the first to suggest the asymmetry was structural rather than transient.

Battalio and Schultz (2006) extended the analysis to single-name equities during the dot-com bubble (1998-2000), finding that retail-heavy underlyings showed even more pronounced non-exercise patterns than indexes. Pool, Stoll, and Whaley (2008) — building on their colleagues’ work — quantified the non-exercise pattern on calls, complementing the put research, and showed that the pattern persisted in modern times.

The seminal paper for our purposes is Barraclough and Whaley (2011). Examining 13 years of CBOE options data (January 1996 through September 2008), they found that retail put holders failed to exercise when NII was positive in approximately 96.3% of cases. They estimated the dollar value of this non-exercise — the amount transferred from retail long-put holders to short-put holders (predominantly market makers and proprietary trading firms) — at \$1.9 billion over the sample period. Per year: roughly \$145m. Per market participant: small. In aggregate: enormous, and structural, and persistent.

Post-2011, the trade has continued. The 2016-2024 period has been particularly profitable because of the Fed’s normalization cycle. NII scales linearly with r , and r went from 0% (2014-2015) to 5.30% (2023-2024). The structural revenue available to capture grew tenfold over that decade.

The regulatory landscape has been quiet on this trade. Unlike the dividend-arb trade (closed by 871(m)) or cum-ex (criminalized), the put-steal has not attracted regulatory attention. The reason: it is technically a transfer between two consenting market participants, not an external tax wedge or fraud. The retail user who fails to exercise is, in some sense, voluntarily giving up the NII. There is nothing illegal about a market maker capturing a wedge that the counterparty leaves on the table. The closest the regulatory machinery has come is broker disclosure: FINRA Rule 4210 requires brokers to clearly explain early-exercise mechanics to retail customers, and some brokers (notably IBKR) provide an “auto-exercise” feature that surfaces NII-positive positions. But the underlying trade is unregulated and ongoing.

The Dodd-Frank Act (2010) had no direct effect on the put-steal. Some indirect effects: the Volcker Rule’s restriction on bank prop trading shifted some of the institutional capture from bank trading desks to non-bank market makers (Citadel Securities, Susquehanna, Optiver, IMC). The capture moved across firms but the total industry capture grew.

5.4 5.2 Why Retail Doesn’t Exercise

Barraclough and Whaley (2011), examining 13 years of options data (1996-2008), found that retail consistently *fails* to exercise even when NII is dollar-positive on every share. The reasons:

1. **Awareness.** Retail platforms historically did not display NII; the user sees only the put price.
2. **Behavioral.** Closing a winning trade by exercising rather than selling feels different.

3. **Friction perception.** Retail conflates “exercise” with “extra commission” — often false.
4. **Path dependence.** A retail user who bought the put at \$2 and is now sitting at intrinsic \$15 sees a 7.5x and may not want to do anything that risks the gain.

The dollar value of this non-exercise, summed across the market, was **\$1.9 billion** over the sample period. Of this:

Captured by:	Share
Market makers / specialists	47.2%
Proprietary trading firms	40.8%
Hedge funds	~9%
Retail (correctly exercising)	3.0%

5.5 5.2b The Behavioral Anatomy of Non-Exercise

The four reasons listed above are accurate but mechanical. To understand why retail systematically fails to exercise even on platforms that *do* display NII, the analysis needs to be more granular. The behavioral finance literature on the put-steal is illuminating because it identifies four distinct cognitive failures, each operating independently and reinforcing each other.

The first cognitive failure is *loss aversion in disposition*. Tversky and Kahneman’s prospect theory predicts that investors weight losses more heavily than equivalent gains. Translated to the put-steal context: a retail user holding a deep-ITM put that has appreciated from \$2 to \$15 sees a \$13 paper gain. Exercising the put converts the unrealized gain to a realized one and starts a new position (long stock at the strike, short stock at the market — or just receiving cash if they were short the underlying). Each step in this conversion feels like a chance to *lose* the gain through some unexpected mechanism — a failed exercise notification, a brokerage error, a tax surprise. The user reasons (incorrectly) that the unrealized gain is “safer” than a realized gain, so they let it ride. This is the disposition effect applied to derivatives.

The second cognitive failure is *complexity avoidance*. The retail user knows how to buy and sell options. The exercise mechanic is different — it requires affirmatively notifying the broker that the user wants to exercise (or, equivalently, allowing the OCC to auto-exercise at expiry, which only happens at expiry, not optimally before). The mental model of “exercise” is murky. Most users have never done it. The default behavior is to *sell the option* rather than exercise, which is operationally familiar. By selling the option for its market value, the user captures less than the intrinsic + NII, but they capture *something familiar*.

The third cognitive failure is *interface design*. Retail brokerages have UX patterns that subtly discourage early exercise. On TastyTrade, the option-position screen prominently displays “Sell to Close” as the primary action for an open long put position; “Exercise” is buried in a sub-menu and accompanied by warnings about settlement and potential overnight risk. On Robinhood, the exercise mechanic is even more obscure — pre-2022, exercising required calling customer support; post-2022, it is buried two screens deep with multiple warning dialogs. The default path is “sell the option for its market price,” which the broker prefers because it generates a commission (until 2019) and clean book-keeping. The exercise path, which is unambiguously better for the user when NII is positive, is the harder path.

The fourth cognitive failure is *peer effect of inactive holding*. Retail traders observe what other retail traders do, and most retail traders do not exercise. The peer-effect feedback loop reinforces the inactive default. When a retail user sees forum discussions of “deep ITM puts” and finds that essentially nobody talks about exercising, the user assumes that not exercising must be the right choice. The consensus is loud and wrong.

These four failures combine into the pattern Barraclough and Whaley observed. The 96.3% non-exercise rate is not driven by a single behavioral bias; it is the conjunction of four independent biases that all push the same direction. Even an occasional retail user who notices NII and considers exercising will often be overridden by one of the other three biases (loss aversion, complexity avoidance, or interface friction).

5.6 5.2c TastyTrade and Robinhood UX Traps — Specific Examples

To make the interface critique concrete: examine the actual UX flows that retail traders encounter. The following examples are based on the platforms’ interfaces as of late 2024.

TastyTrade. A user has an open long put position at \$50 strike, underlying at \$40. The position is displayed as: “−1 XYZ Mar 21 \$50 P / −\$10.05” (the negative sign indicates a long position, the price is the current market

value). Clicking the position opens a panel with two prominent buttons: “Close Position” (large, blue) and “Roll” (medium, gray). To exercise, the user must click “...” (a small ellipsis) to expand options, then select “Exercise” from a dropdown. The exercise option is accompanied by a warning: “Early exercise is final and may result in unexpected stock positions. Are you sure?” Most users dismiss the dialog and choose “Close Position” instead — which sells the put at its market price.

The market price for a deep ITM put is typically intrinsic plus a small time-value remainder. If NII is positive, the market price is *slightly less than intrinsic*, because the market is pricing in the small probability of correct exercise by some holder somewhere. By selling rather than exercising, the user captures intrinsic minus the small time value premium that the market is paying for the call’s option-value. They do not capture the NII. The NII flows to whoever bought the put from them at this price (typically a market maker who will exercise it correctly).

Robinhood. The deep-ITM-put-position view on Robinhood is even more spartan. A long put displays only its market price and the simple “Sell” button. Exercise is not available from the position screen at all. To exercise, the user must navigate to the “Help & Support” section, then “Options Trading,” then “Exercise an Option Early,” then click through three educational screens, and finally submit an exercise request. The request is then queued and processed by Robinhood operations — typically within 1-2 business hours during market hours. The friction is so substantial that the exercise option is almost never used by retail.

The retail user who is even a little curious — who pulls up the options chain for their put, sees “\$10.05 last trade,” sees the spot at \$40, mentally computes “intrinsic = \$50 - \$40 = \$10, so the put is trading at \$0.05 of time value” — has all the information needed to decide on early exercise. They typically do not. The combination of friction and behavioral bias is a moat the retail platforms have built (intentionally or not) that holds NII firmly in the hands of market makers.

IBKR. The exception. Interactive Brokers’ interface, designed for sophisticated retail and small-institutional users, prominently displays NII (computed by IBKR’s own pricing engine) on each open ITM American option position. The display shows the dollar value of the NII along with a recommendation: “Early exercise may be optimal for this position.” A “Exercise” button is on the same screen as “Sell.” The friction is roughly equal between the two paths.

IBKR users — who are systematically more sophisticated than the average retail trader — exercise their deep-ITM American puts at substantially higher rates than TastyTrade or Robinhood users. Empirical comparisons (Battalio-Engel-Mehran 2017) show IBKR retail exercising NII-positive puts ~25-30% of the time, vs ~3-5% for the broader retail base. The platform difference is a behavioral nudge, and the nudge moves capital.

This is the world the bull put credit spread strategy in this repo operates in. The strategy captures the NII that retail-on-easy-platforms forfeits, packaged into a defined-risk structure that retail can run on the same easy platforms.

5.7 5.3 The Short Stock Interest Arbitrage Game

Market makers run a precise book to harvest non-exercise:

1. They buy a deep ITM put (long).
2. They simultaneously sell an equivalent quantity of the same put (short).
3. Net delta is approximately zero.
4. Each day, the OCC randomly assigns exercises across all open shorts on that contract.
5. If a long holder somewhere in the market exercises, *some* shorts get assigned. If the assignment lands on the market maker’s short:
 - They are forced to buy stock at X (or deliver short stock).
 - They receive X in cash.
 - They earn one day’s interest on X.
 - The next morning they re-establish the position by buying back the short and re-shortening at the new lower deep-ITM strike.
6. If no long exercises that day, all shorts (including the market maker’s) hold the position with no assignment.

The expected value calculation: the market maker’s *short* leg is the position that *captures* the NII when retail mis-exercises elsewhere. Specifically, the market maker is short a put that should be exercised; when it isn’t, the market maker is overpaid by the time-value premium that should have been captured by the long. This value flows back to the market maker through the option’s pricing.

A more direct way to see the harvest: the market maker writes deep ITM puts at slightly less than parity (because they price in the small probability of *correct* retail exercise). When 96.3% of long retail holders fail to exercise, the

maker pockets that 3.7% pricing wedge times notional times days held — the full \$1.9bn over the sample.

5.8 5.4 Numerical Examples Across Rate Regimes

The NII formula's first-order driver is r . The dependence on rate is approximately linear for small rT , so doubling the rate roughly doubles the NII. This makes the trade's profitability highly regime-dependent. Across the past 15 years, the US has cycled through three distinct rate regimes; the next decade will likely include at least one more (a normalization back down). Each regime produces a different put-steal economy.

Regime A — ZIRP (2009-2015, 2020-2022). $r = 0\%$. NII = 0 regardless of how deep ITM. The trade does not exist. Market makers don't run the book. Retail mis-exercise is harmless.

To make this concrete: in March 2020, the Fed cut rates to 0-0.25% in response to COVID. SOFR dropped to ~5 bps. A market maker holding a deep ITM put position at this point computes $NII = X * (1 - e^{-0.0005 * T}) - C(S, X, T, r, \sigma)$. For $X = \$100$, $T = 30$ days, $r = 0.05\%$: the first term is $\$100 * (1 - e^{-0.000411}) = \$100 * 0.000411 = \$0.041$ per share. The call value (the cost of waiting) is positive — for a deep ITM put with 30 DTE, the corresponding call has time value of perhaps \$0.10-0.30. So $NII = \$0.004 - \$0.20 = -\$0.196$. *Negative*. The market maker should hold, not exercise. Retail correctly should also hold. There is no edge.

This is why the put-steal trade essentially disappeared from market-maker books in 2020-2021. There was no NII to capture. Citadel Securities, Susquehanna, and the other major liquidity providers were running the put-steal book at minimum activity during this regime, and the institutional flow that used to harvest the trade was redirected into other margin businesses (volatility selling, dispersion trading, etc.).

Regime B — Normal (2016-2019, 2024-present). $r = 4\text{-}5\%$. NII is positive on any reasonably ITM put with reasonable T . The trade is in season. This is the current environment as of 2025-2026.

In May 2024, the Fed funds rate is 5.25-5.50%, SOFR is ~5.30%. A market maker computes NII for the same setup: $X = \$100$, $T = 30$ days, $r = 5.30\%$, $\sigma = 25\%$, $S = \$80$ (20% ITM). First term: $\$100 * (1 - e^{-0.0530 * 0.0822}) = \$100 * (1 - 0.99566) = \$0.434$ per share. Call value at $S=\$80$, $X=\$100$, $T=30$ days, $r=5.3\%$, $\sigma=25\%$: by Black-Scholes, approximately \$0.014. $NII = \$0.434 - \$0.014 = \$0.420$ per share. *Positive*. Strongly positive.

On a single ITM put contract (100 shares), this is \$42 of NII captured by exercising one day vs holding. Over 30 days, the cumulative NII is roughly $\$0.42 * 100 = \42 per contract. For a market maker running a book of, say, 50,000 such contracts on the various deep-ITM dividend-paying names, the daily NII capture is enormous — a single-firm position on a busy day can capture \$50,000-\$200,000 of NII on a single contract series.

Regime C — Hyperinflation (1980s, hypothetical 2026+ tail). $r = 10\%$. NII is dramatic. Market makers run massive deep ITM put books; retail is hammered.

In the 1980s the Volcker Fed raised rates to 14-19%. NII at those levels is enormous. The full S&P 100 deep-ITM put market in 1981 was capturing tens of millions in NII annually for the specialists on the floor of the Pacific Coast Stock Exchange (where S&P 100 options were listed at the time). Diz and Finucane's 1993 paper documented the magnitude.

A modern hypothetical: if inflation pushed Fed funds back to 8% in 2026 (a real-world possibility under various scenarios), the NII trade would be roughly 1.6x more profitable than in 2024. The market maker's daily NII capture on a single contract series might run to \$80-300k.

Regime D — Anticipated normalization down (2026 forward). As of May 2026, the Fed has begun cutting; SOFR is at 4.0-4.5%. NII is positive but compressed. The bull put credit spread variant is still profitable but the credit per spread is lower than 2024 levels (maybe \$2.50 instead of \$3.65 on a comparable position). The retail trader running the trade systematically will see strategy returns decline ~20-30% from peak as the rate regime cools.

Numerical table of NII per share at various stock prices, $X = \$100$ strike, $\sigma = 25\%$, $T = 30$ days:

Spot S	r=0%	r=2%	r=4%	r=5%	r=6%	r=8%
\$60	\$0.00	\$0.16	\$0.33	\$0.41	\$0.49	\$0.66
\$70	\$0.00	\$0.15	\$0.32	\$0.40	\$0.48	\$0.64
\$80	\$0.00	\$0.13	\$0.27	\$0.34	\$0.41	\$0.55
\$90	\$0.00	\$0.07	\$0.16	\$0.20	\$0.24	\$0.32
\$95	\$0.00	\$0.02	\$0.06	\$0.07	\$0.09	\$0.13
\$98	\$0.00	\$0.00	\$0.01	\$0.02	\$0.03	\$0.05

The NII drops rapidly as S approaches X — at $S = X$ the call value (the right to wait) approaches the maximum interest value, neutralizing the NII. **The trade is therefore most reliable when the put is at least 10-15% in-the-money.**

5.9 5.4b Side-by-Side Trade Variants — Put vs Call, Single-Leg vs Spread

The put-steal trade has multiple expressions. Understanding the variants matters because each has different risk, capital, and execution profiles. Five variants are worth comparing.

Variant 1: Long deep-ITM American put, exercise when $NII > 0$. This is the textbook implementation. The retail user buys a deep-ITM put on a non-dividend-paying stock with positive NII, exercises early, captures the strike's interest. Capital required: full intrinsic + small premium. Risk: small (the put is essentially the stock minus a small carrying cost). Expected return: approximately the NII per day held. Practical use: limited because the user must already own the put. Most retail users do not buy deep-ITM puts; they buy near-the-money or out-of-the-money puts for directional exposure.

Variant 2: Short deep-ITM American put, hope long fails to exercise. This is the institutional bread-and-butter. The market maker writes the deep-ITM put at slightly less than parity, hedges with the underlying, and waits. If the long counterparty exercises (correctly), the maker takes assignment, receives X , and continues the position with a fresh hedge. If the long fails to exercise (96.3% of the time per Barraclough-Whaley), the maker captures the NII as a wedge in pricing. Capital required: substantial margin (the short put has theoretically unlimited risk if the stock goes to zero). Risk: high — assignment is not the worst outcome, but the potential for the underlying to gap down through the strike requires careful sizing. Expected return: approximately the NII per day held times the non-exercise rate.

Variant 3: Bull put credit spread. This is the retail-executable variant in `strategies/put_steal.py`. The user sells a slightly-ITM put and buys a further-OTM put as a wing, defining the maximum loss. The credit captured embeds the NII wedge. Capital required: defined (the wing width minus credit, times 100). Risk: defined. Expected return: approximately 20-30% on the spread per cycle when conditions are right (NII positive, AI gating positive, IV and VIX in range). This is the strategy this chapter focuses on.

Variant 4: Deep-OTM call as the symmetric twin. A less-known fact: the put-steal has an analog on the call side, but only on dividend-paying stocks. A deep-ITM American call should be exercised early just before an ex-date if the dividend exceeds the call's remaining time value. Empirically (Pool-Stoll-Whaley 2008), retail systematically fails to exercise these calls. The market maker writing deep-ITM calls on dividend-paying stocks captures a wedge analogous to the put-steal NII. The trade structure mirrors the put trade but is driven by dividend timing rather than rate level.

Variant 5: Conversion arbitrage with NII tilt. The conversion (Section 4.18) plus an NII overlay. The user holds long stock, short call, long put — the synthetic riskless position. The position's normal economics earn the cost of carry; adding an NII tilt by selecting deep-ITM strikes with positive NII captures additional spread. This is the form most institutional desks run. It is largely retail-inaccessible because the put-call combo orders are not atomically fillable on consumer platforms.

The retail trader should focus on Variant 3 (bull put spread) because Variants 1 and 2 require holding deep-ITM puts (which most retail does not), Variant 4 requires dividend-paying stock and call exercise (more complex), and Variant 5 requires institutional execution. The bull put credit spread is the cleanest retail expression of the underlying institutional trade.

5.10 5.5 The Repo Spec Connection

A subtle but important point: the NII formula above assumes the holder can finance the strike at the risk-free rate. For a *bank* desk, the strike X can be invested at SOFR. For a *retail* user, the strike X sits in their margin account earning either nothing (cash) or a bit (interest on cash balance, broker-dependent).

This means retail's NII benefit is *smaller* than the desk's NII benefit by the spread between SOFR and the broker's interest-on-cash rate. At Robinhood Gold, the cash-sweep yield is ~5% (close to SOFR). At a non-Gold retail account, the yield is closer to 0%.

For retail without cash interest:

$$\text{Retail NII} = X \times (1 - e^{-\{r_{\text{broker}} \times T\}}) - C$$

With $r_{\text{broker}} = 0\%$, retail NII is identically negative. Retail correctly should *not* exercise — they have no investment alternative.

This is why the Robinhood Gold tier (and similar at Webull, IBKR Lite vs Pro) materially affects whether the put-steal trade is even worth considering for retail. Without 5% interest on cash, the trade has no edge.

5.11 5.5b Why the Wedge Doesn't Get Arbitraged Away

A natural question: if 96.3% of retail fails to exercise NII-positive puts, why don't more market makers enter the trade until the wedge compresses to zero? The market should "fix" this inefficiency.

The answer has three parts.

First, the wedge *is* compressed by competition, just not to zero. The 1.9 billion that Barraclough-Whaley measured over 1996-2008 represents what was left on the table *after* market makers and prop firms had captured their 88% share. The trade is intensely competitive at the institutional level. New entrants (Citadel Securities post-2010, Jane Street since the late 2000s, Optiver and IMC throughout) have continually pushed the bid-ask wedge tighter.

Second, the trade requires *capital* to scale. A market maker running the put-steal book must hold short put inventory across hundreds or thousands of names, marking each to market and managing the hedges. This consumes balance sheet, margin, and operational capacity. The marginal capital cost of expanding the book is rising, and at the institutional level, the wedge has compressed enough that further expansion is not economically efficient. The trade pays for the capital deployed, but does not pay enough to attract additional new capital.

Third, the trade does not *pay enough relative to the regulatory and operational risk* to attract banks. Post-Volcker, banks cannot run prop trading desks. The trade is therefore concentrated in non-bank market makers who do not face Volcker restrictions. The non-bank market makers (Citadel, Susquehanna, Optiver) have, in aggregate, perhaps \$30-50 billion of capital deployed across all market-making businesses. The put-steal is one of dozens of trades they run. The trade pays approximately a 5-15% annual return on dedicated capital — a respectable but not extraordinary return. The non-bank market makers run it at the level that makes sense for their portfolio; they do not increase exposure beyond that.

The retail user's Bull Put Credit Spread (Variant 3) is therefore *not* directly competing with the institutional desk. The retail trade is capturing a residual: the small portion of the wedge that is left on the table even after institutional competition. The retail edge is structural (not directional) and small (10-25% annualized on the deployed capital), but it is real because the retail trader's capital is a small fraction of the institutional capital and does not push the market.

5.12 5.6 The Bull Put Spread Implementation

The retail-executable version, implemented in `strategies/put_steal.py`, is a **bull put credit spread**. The structure:

Leg	Direction	Strike	Purpose
Short put	Sell	$\text{Spot} \times (1 - \text{itm_pct})$	Captures forfeited NII
Long put	Buy	$\text{Short} \times 0.96$	Defined max loss

The desk's edge of capturing 3.7% of put non-exercise is replicated for retail by *being short* a deep ITM put — collecting the credit that includes the implicit NII premium. When retail counterparties mis-exercise (or fail to), the credit erodes more slowly than fair value would imply, generating the structural alpha.

Concrete example, August 2024.

- Stock: META at \$510.
- Sell put at \$500 strike, 21 DTE: credit \$5.50.
- Buy put at \$480 strike, 21 DTE: pay \$1.85.
- Net credit per spread: \$3.65, or \$365 per spread.
- Max loss: \$20 wing — $\$3.65 = \16.35 , or \$1,635 per spread.
- Risk/reward: 22.3% return on capital if META stays above \$500.

The structural NII edge is small — about 5-15 bps of the strike per month — but it pushes the breakeven slightly favorable to the seller. Combined with the GBM classifier in the strategy file (which filters for crash risk), the combined win rate target is 65-70%.

5.13 5.7 When the Trade Breaks

The trade is structural under stable conditions. It breaks under:

1. **Sudden rate cuts.** A 50bps Fed cut in a single meeting compresses NII immediately. If rates drop from 5% to 4.5% mid-trade, the NII edge halves.
2. **Single-name vol shocks.** A stock that gaps down through the short strike turns the credit spread into a max-loss event. Earnings season is the worst regime.
3. **Dividend reinstatement.** A stock that initiates or increases a dividend during the trade window changes the put-call parity and may reverse the early-exercise calculus for the underlying.
4. **HTB regime shift.** If the underlying enters HTB territory mid-trade, the cost-of-carry rises and put pricing becomes more complex; standard NII calculation breaks down.

The strategy code at [strategies/put_steal.py](#) has explicit gates for $IV \leq 60\%$, $VIX \leq 40$, and AI confidence $\geq 55\%$ to filter most of these regimes. Even so, it is a *structural* not *predictive* edge — it does not beat directional predictions and does not survive every regime.

5.14 5.8 The Mathematics of the Early Exercise Boundary

The optimal exercise boundary for an American put can be computed via a free-boundary PDE. The boundary $S(t)$ is the stock price above* which holding is optimal, and below which exercise is optimal. As $t \rightarrow T$ (expiry approaches), $S(t) \rightarrow X$. As $t \rightarrow 0$ (time-zero), $S(0)$ is the deepest critical level.

For non-dividend-paying stock with constant rate r and vol σ , an approximation (Bjerk Sund-Stensland 2002):

$$S^*(t) = X \times \frac{1}{1 - \frac{1}{2}k} - \sqrt{\left(\frac{1}{2}k - \frac{1}{2}\right)^2 + 2r/\sigma^2}$$

$$\text{where } k = 2(r/\sigma^2) \text{ (no dividend; for dividends } k = 2(r - q)/\sigma^2 \text{)}$$

For $X = \$100$, $r = 5\%$, $\sigma = 25\%$, the approximation gives:

$$\begin{aligned} k &= 2 \times 0.05 / 0.0625 = 1.60 \\ &= (1/2 - 1.60) - \sqrt{((1.60 - 0.5)^2 + 1.60)} \\ &= -1.10 - \sqrt{1.21 + 1.60} \\ &= -1.10 - \sqrt{2.81} \\ &= -1.10 - 1.676 \\ &= -2.776 \end{aligned}$$

$$S^*(0) = 100 \times (-2.776) / (-2.776 - 1) = 100 \times 0.7350 = \$73.50$$

The exercise boundary is at \$73.50 — about 26.5% ITM. Retail holders failing to exercise below this threshold are leaving NII on the table.

The boundary is *time-dependent*: at $t = T/2$, the boundary moves to roughly \$79; at $t = 0.9T$, the boundary is at \$93; at expiry, the boundary is at \$100 (i.e., any ITM put should be exercised).

A retail trader implementing this rule on real positions would: 1. Compute the current $S(t)$ for each ITM put position. 2. Exercise any put where $S < S(t)$. 3. Re-check daily.

This is mechanical, not predictive. The Robinhood mobile app does not surface this calculation. Most retail tools do not. A spreadsheet implementation suffices.

5.15 5.9 Multi-Period NII Analysis

For a position held over multiple days, the cumulative NII captured by the *short* put holder (when the long fails to exercise) approximates:

$$\text{Cumulative NII over } n \text{ days} = X \times (1 - e^{-r \times n/365})$$

For $X = \$100$, $r = 5\%$, $n = 21$ days: $\$0.288/\text{share}$. Over 90 days: $\$1.234/\text{share}$. Over 365 days: $\$4.879/\text{share}$.

A short put position deep ITM, held a full year against a non-exercising long, captures approximately the full year's interest on the strike — multi-percent of notional. The persistence of this revenue is why market makers run the deep-ITM book continuously.

5.16 5.10 Cross-Strike Comparison Table

A useful reference for the put-steal trade. NII per share (in dollars) for a put 30 days from expiry, $\sigma = 25\%$, $r = 5\%$, at various ITM depths:

Stock S	Strike X	Moneyness	NII/share	Notes
\$100	\$100	ATM	\$0.00	No edge
\$95	\$100	5% ITM	\$0.03	Marginal
\$90	\$100	10% ITM	\$0.18	Modest edge
\$85	\$100	15% ITM	\$0.34	Reasonable
\$80	\$100	20% ITM	\$0.41	Strong
\$75	\$100	25% ITM	\$0.41	Plateau (max NII reached)
\$70	\$100	30% ITM	\$0.41	Plateau
\$50	\$100	50% ITM	\$0.41	Plateau

The NII reaches a plateau because once the call value (the right to wait) goes to zero, $NII = X(1 - e^{-rT})$. Further ITM-ness does not add to NII but does reduce *risk* (the put is even safer).

For the bull put credit spread retail variant, a target depth of **10-15% ITM** captures most of the NII edge while keeping the position liquid (deep ITM puts can have wide bid/ask spreads). The strategy in [strategies/put_steal.py](#) uses `itm_pct = 0.05` by default — moderately aggressive on liquidity preference but capturing less of the structural edge.

5.17 5.11 The “Pin and Crush” Variant

A more advanced retail variant of the put-steal strategy: combine the bull put credit spread with a *short call vertical* on the same underlying near the same expiry. This creates an iron condor with a *bias* toward the put-side.

The reasoning: when NII is positive, the underlying is also implicitly being financed by the desk's hedging. The call side typically has *less* premium per unit risk than the put side. By selling the put-side credit (capturing NII) and a small call-side credit (capturing skew premium), the trader builds a position with a tighter expected breakeven and a slightly higher win rate.

This is essentially the mechanism in [strategies/vol_arbitrage.py](#) — the put-side captures NII and skew premium, while the call-side hedges with a smaller credit.

5.18 5.12 Failure Modes — Detailed Analysis

A systematic put-steal book — institutional or retail — has a number of failure modes that the practitioner must understand. Section 5.7 listed four. This subsection walks through each in operational detail.

Failure mode 1: Sudden rate cut. The Fed cuts 50 bps in a single FOMC meeting. NII compresses by approximately 50%. Existing positions are still profitable but the *forward* economics of the trade are halved. The trader running the book must size down (commit less capital to new entries) and may need to close existing positions early if the rate-cut implications are larger than expected.

The August 2024 SEP and the December 2024 FOMC are recent examples where rate paths shifted suddenly. A retail trader running the bull put spread who entered positions in mid-July 2024 (when rates appeared steady) and held through the August unwind found that the credit they had collected was now slightly cheaper to buy back than expected (because the underlying environment had shifted). They closed the positions for partial profit rather than waiting for full expiration.

Failure mode 2: Single-name vol shock. The stock gaps down through the short strike. The bull put spread crystallizes at maximum loss. This is the dominant downside risk. A typical trade has max loss of 4.5x the credit collected; a single failed trade on a 4-cycle holding pattern wipes out the gains from the prior 4 cycles.

The mitigation in `strategies/put_steal.py`: the AI confidence threshold (`confidence_thresh = 0.55` by default) filters out names where the GBM classifier predicts high crash risk. The `iv_max = 0.60` and `vix_max = 40` gates additionally filter out high-volatility entries. These filters cut the trade frequency by roughly half but should reduce the failure-mode-2 incident rate by considerably more than half.

A practical rule for the retail trader: *never* run the put-steal trade through earnings. Even if the IV is below 60% and the VIX is below 40, the binary nature of earnings can produce a 10-20% gap that no other filter can predict. The strategy code does not currently have an explicit earnings filter; the retail user adding one (manually checking earnings dates and skipping trades within 5 trading days of earnings) is a clear improvement.

Failure mode 3: Dividend reinstatement or initiation. A stock that begins paying dividends (or increases an existing dividend) during the trade window changes the put-call parity calculation. The short put becomes more valuable (because the stock is now expected to drop on ex-date), reducing the profitability of the bull put spread. This is a slow-moving risk and rarely fatal but does accumulate.

Mitigation: monitor announced dividends and avoid entering trades with ex-dates within the holding period. The strategy file does not include this filter explicitly; adding a dividend-calendar lookup is straightforward.

Failure mode 4: HTB regime shift. The underlying enters HTB territory mid-trade. The put-call parity calculation changes — synthetic short positions now have a meaningful borrow cost embedded — and the put pricing becomes more complex. The bull put spread’s mark-to-market may move adversely as the HTB premium gets embedded in the spread’s pricing.

Mitigation: avoid names with structural HTB risk (small-cap biotechs, recent IPOs, names with recent convert issuance). The strategy file does not have an HTB filter; the retail user adding a borrow-rate threshold (skip trades on names with borrow > 100 bps) is another practical improvement.

Failure mode 5: Surprise early exercise (the 3.7% case). The 96.3% non-exercise rate has a complement: 3.7% of long put holders *do* exercise correctly. When the holder exercises, the random assignment process at OCC distributes the assignments across all open shorts. If the retail trader’s bull put spread is short the put that gets assigned, they receive an unexpected stock position (long stock at the strike, short their bull put long leg). The trade is no longer a defined-risk options spread; it is a stock position that must be managed.

This is the surprise-early-exercise risk. Probability: roughly 3.7% per ITM contract per day, roughly equal to the implied probability times the time to expiry. For a 21 DTE position, the cumulative assignment risk is approximately $21 * 3.7\% / 365 = 0.21\%$ per day, which over 21 days is roughly 4-5% cumulative. Most retail bull put spread positions never see assignment. But for the few that do, the trader must be prepared.

Mitigation: when assigned early, the trader receives long stock at the strike. The trader should immediately sell the stock at market (closing the long position and realizing intrinsic value). The trader’s net P&L is approximately the intrinsic value at assignment minus the original credit minus the closeout cost. Usually this is a small profit or breakeven, not a disaster.

A subtle mitigation: TastyTrade and IBKR offer “managed accounts” that auto-handle assignment. Robinhood does not — assignment results in a stock position that the user must manually close. Retail users on Robinhood should monitor for assignment notifications during the trade life and respond promptly.

5.19 5.13 Actionable Rules of Thumb for Retail

Given all the above analysis, the following rules of thumb are what a retail quant can actually do, today, on TastyTrade or IBKR. These are not theoretical; they are operational.

Rule 1: Compute NII before every entry. Use the formula $NII = X * (1 - \exp(-r*T)) - C$, where C is the Black-Scholes call value at the same strike. This is one line of code, available in `strategies/put_steal.py` at `_compute_nii`. If $NII < \$0.05/\text{share}$, do not enter.

Rule 2: Use SOFR as r. The relevant rate for the formula is the user’s *opportunity cost of capital* over the trade horizon. For a retail user with cash sweep at 5%, use 5%. For a non-Gold Robinhood user with 0% cash sweep, the institutional NII formula overstates the user’s edge — the user has no investment alternative and the put-steal trade is less attractive. Open a Gold or Pro account first.

Rule 3: Target 10-15% ITM on the short put. This is the sweet spot where NII is near maximum but the option is still liquid. Going deeper ITM (20%+) gets wider bid-ask spreads. Going less ITM (5% ITM) reduces the

NII and the spread's max return.

Rule 4: Use 21-30 DTE. Shorter DTE compresses NII; longer DTE adds variance. The balance is around 21-30 days. The strategy default of 21 is appropriate.

Rule 5: Set wing width to 4-5% below the short strike. This caps max loss. The default `wing_pct = 0.04` is reasonable. Wider wings reduce capital efficiency; tighter wings increase capital efficiency but also increase max-loss / credit ratio.

Rule 6: Never trade through earnings. Manually check the earnings calendar. Skip names with earnings within the trade window (entry to expiry). The risk is binary and not modelable by the GBM classifier.

Rule 7: Skip names with borrow rate > 100 bps. HTB names have option pricing distortions that are hard to model. Stick to general collateral names.

Rule 8: Skip names with VIX > 40 or IV > 60%. The strategy has these gates built in; respect them.

Rule 9: Size positions at 1-2% of capital per trade. The default `position_size_pct = 0.02` is appropriate. Larger sizing concentrates risk; smaller sizing dilutes returns. With 30+ active positions across diverse names, the diversification benefit reduces variance.

Rule 10: Profit-take at 50% of max profit. The default is correct. Holding to expiration captures the last 50% of profit but exposes the position to last-week vol shocks. Closing at 50% maximizes the Sharpe ratio of the trade.

Rule 11: Hard stop at 2x max loss. The default `stop_loss_mult = 2.0` is also appropriate. Letting losses run beyond 2x the credit is rarely productive; the position has likely broken structurally.

Rule 12: Run on 30+ names simultaneously. The trade's edge is small per cycle (maybe \$40-100 of NII captured per spread). The annualized return on capital is in the 10-25% range. Achieving this requires running many trades; concentration on one or two names exposes the user to tail risk that defeats the structural edge.

These twelve rules are the operating manual for a retail put-steal book. The strategy file in this repo encodes most of them; the retail user adding rules 6 (earnings) and 7 (HTB) manually will materially improve performance.

5.20 5.14 Walking Through strategies/put_steal.py

The strategy file in this repo, `strategies/put_steal.py`, is the production implementation of the put-steal trade described in this chapter. The retail quant should read the file before deploying capital. This subsection walks through the key sections to connect the academic content of the chapter to the executable code.

The NII computation. Lines 96-110 implement `_compute_nii`:

```
def _compute_nii(S: float, X: float, T: float, r: float, sigma: float) -> float:
    if T <= 0:
        return 0.0
    interest_income = X * (1.0 - np.exp(-r * T))
    call_value      = _bs_price(S, X, T, r, sigma, "call")
    return float(interest_income - call_value)
```

This is exactly equation (4) from Section 5.1b. The function takes spot, strike, time-to-expiry, rate, and volatility, and returns the NII per share. The function is used both at training time (to build features) and at signal time (to gate entry decisions).

The entry gate. Lines 336-354 implement the live signal generation. The key check:

```
if vix > self.vix_max or sigma > self.iv_max:
    return SignalResult(... "HOLD" ...)
X = S * (1.0 + self.itm_pct)
nii = _compute_nii(S, X, T, r, sigma)
if nii > self.nii_threshold:
    confidence = min(0.90, 0.55 + nii * 3.0)
    return SignalResult(... "SELL" ...)
```

The gate is: $VIX < 40$, $IV < 60\%$, $NII > 5$ cents/share, then return a SELL signal with confidence scaled by NII level. A larger NII corresponds to higher confidence (and thus larger position sizing in the broader system).

The feature set. Lines 113-185 implement `_build_features`, which constructs the 12-feature matrix the GBM classifier uses. The features are:

- NII signals: `nii_level`, `nii_ma5`, `nii_to_cred` (NII relative to put credit), `call_to_put` (proxies for early-exercise leakage).
- Rate: `risk_free_rate` from the auxiliary data.
- Stock momentum: `ret_5d`, `ret_20d`, `dist_from_ma50`, `atr_pct`.
- Vol: `iv_level`, `ivr_20d`, `iv_5d_change`.
- Market: `vix_level`.

The feature set is deliberately broad: NII signals identify the *opportunity*; momentum and vol features identify when the *opportunity is safe to take* (vs likely to crash through the short strike).

The labels. Lines 188-209 implement `_build_labels`. The label is 1 if the stock stays above the short strike for all `n_forward` days (the period over which the bull put spread would survive). Positive rate is around 65-70% historically — i.e., the bull put spread survives about two-thirds of the time when entered randomly. The classifier’s job is to push this success rate higher by gating entries.

The walk-forward training. Lines 484-516 implement the walk-forward simulation. Key features:

- 90-bar warmup (`_WARMUP_BARS = 90`).
- Retrain every 20 bars (`_RETRAIN_EVERY = 20`).
- Walk-forward: each retrain uses only data prior to the current bar, no future leakage.
- GradientBoostingClassifier with 60 trees, depth 3, learning rate 0.05.

This walk-forward methodology is critical for honest backtesting. A backtest that uses future data in feature construction or labeling produces inflated results. The implementation here is careful about this — the labels and features are constructed *over the full series* but the model is *only fit on past data* at each retraining step.

The trade management. Lines 518-568 implement the open-trade management loop. On each bar, the loop checks:

- Current spread value (mark-to-market).
- Profit target: close at 50% of max profit.
- Stop loss: close at 2x max loss.
- DTE exit: close at 5 days before expiry to avoid pin-day vol.

This three-condition exit logic is standard for short premium trading. The 50% profit target is well-supported by tastytrade research on credit spread management.

The entry sizing. Lines 596-606:

```
short_K = spot * (1.0 - itm_p)          # short put = itm_p below spot
long_K  = short_K * (1.0 - wing_p)      # long put  = wing_p below short
credit  = (
    _bs_price(spot, short_K, T, r, iv, "put") -
    _bs_price(spot, long_K, T, r, iv, "put")
)
wing_w  = short_K - long_K
max_loss = max(0.01, (wing_w - credit) * 100)
n_spread = max(1, int((capital * pos_pct) / max_loss))
```

The sizing is risk-based: position size is the floor of $(\text{capital} * \text{position_size_pct}) / \text{max_loss}$. This means each trade risks at most `position_size_pct` of capital, which is 2% by default. With 50 simultaneous trades, this concentrates 100% of capital, which is acceptable given the diversification.

Wait — there is a subtle bug. Line 597 says `short_K = spot * (1.0 - itm_p)`. But the short put strike should be *above* spot for the put to be ITM. With `itm_p = 0.05`, $(1.0 - \text{itm_p}) = 0.95$, so `short_K = spot * 0.95`, which is *below* spot. A put at a strike below spot is OTM, not ITM.

This is a code-level confusion: the variable `itm_pct` is being used as if it were “ITM percentage” but the multiplication is $(1 - \text{itm_p})$ which puts the strike below spot. The strategy is therefore actually selling slightly OTM puts, not slightly ITM puts. The trade is still profitable — selling slightly OTM puts at high enough credit can still capture an NII-related wedge — but the explanation in the docstring is inverted.

The retail user who runs the strategy as-is is therefore running a *slightly OTM* bull put spread, not a slightly ITM one. The code works (it generates returns), but the connection to the put-steal academic literature is via the option-pricing wedge embedded across the chain rather than the direct NII capture on a deep ITM put.

This is a minor discrepancy and does not invalidate the strategy. The retail user can either (a) run the strategy as-is, accepting the slightly OTM interpretation, or (b) modify the line to `short_K = spot * (1.0 + itm_p)` to actually use ITM strikes. The choice depends on the user’s preference for higher-credit (ITM) vs higher-probability-of-survival (OTM) trades.

5.21 5.15 References

- Barraclough, K. and Whaley, R.E. (2011), “Early Exercise of Put Options on Stocks”, *Journal of Finance* 67(4): 1423-1456.
- Diz, F. and Finucane, T.J. (1993), “The Rationality of Early Exercise Decisions: Evidence from the S&P 100 Index Options Market”, *Review of Financial Studies* 6(4): 765-797.
- Battalio, R. and Schultz, P. (2006), “Options and the Bubble”, *Journal of Finance* 61(5): 2071-2102.
- Pool, V.K., Stoll, H.R., and Whaley, R.E. (2008), “Failure to Exercise Call Options: An Anomaly and a Trading Game”, *Journal of Financial Markets* 11(1): 1-35.
- Bjerk Sund, P. and Stensland, G. (2002), “Closed Form Valuation of American Options”, working paper, Norwegian School of Economics.
- Carr, P. and Faguet, D. (1996), “Valuing Finite-Lived Options as Perpetual”, working paper, Cornell University.
- Roll, R. (1977), “An Analytic Valuation Formula for Unprotected American Call Options on Stocks with Known Dividends”, *Journal of Financial Economics* 5(2): 251-258.
- Merton, R.C. (1973), “Theory of Rational Option Pricing”, *Bell Journal of Economics and Management Science* 4(1): 141-183.
- Battalio, R., Engel, S., and Mehran, H. (2017), “The Cost of Short Selling”, working paper.
- Tversky, A. and Kahneman, D. (1979), “Prospect Theory: An Analysis of Decision Under Risk”, *Econometrica* 47(2): 263-291.
- Whaley, R.E. (2003), “Derivatives on Market Volatility: Hedging Tools Long Overdue”, *Journal of Derivatives* 1(1): 71-84.
- Brennan, M.J. and Schwartz, E.S. (1977), “The Valuation of American Put Options”, *Journal of Finance* 32(2): 449-462.
- Geske, R. and Johnson, H.E. (1984), “The American Put Option Valued Analytically”, *Journal of Finance* 39(5): 1511-1524.

5.22 5.16 A Closing Note on Chapter 5

The put-steal trade is, in some sense, the cleanest retail edge described in this guide. It rests on a documented behavioral pattern (96.3% non-exercise rate per Barraclough-Whaley), a transparent mathematical formula ($NII = X(1 - e^{-rT}) - C$), and a structural mechanism that has persisted for decades despite institutional attention. The retail-executable variant (bull put credit spread) translates the institutional trade into a defined-risk wrapper that can be run on TastyTrade, IBKR, or any Level 3 options-enabled retail account.

The chapter has been the longest in the guide because the underlying mathematics deserve the depth, the historical and behavioral context illuminates *why* the trade works, and the connection to the production code in `strategies/put_steal.py` lets the retail quant move from understanding to deployment. The reader who has internalized the 5.1b derivation, the 5.2 behavioral analysis, the 5.4 rate-regime examples, and the 5.13 rules of thumb has the full mental model needed to deploy the trade with discipline.

The reader who skips ahead to deployment without internalizing these pieces will likely run the strategy mechanically and either lose interest after a small drawdown or over-size into a single position when the GBM classifier confidence is unusually high. Both failure modes can be avoided by understanding the trade’s structural foundations.

6 Chapter 6 — SPX Boxes and Synthetic Lending

This is the second flagship chapter. Box spreads are the cleanest, mathematically purest instrument in the options market. They are also one of the few institutional financing tools that retail can execute essentially identically to

institutions — a rare alignment.

6.1 6.1 The Four-Leg Structure

A box spread combines a bull call spread and a bear put spread at the same strikes. The structure:

```
At K1 (lower strike):  
  Buy  call at K1  
  Sell put  at K1
```

```
At K2 (upper strike, K2 > K1):  
  Sell call at K2  
  Buy  put  at K2
```

At expiry, the payoff is:

```
Payoff at expiry, regardless of underlying S_T:  
  Long call(K1) at expiry: max(S_T - K1, 0)  
  Short put(K1) at expiry: -max(K1 - S_T, 0)  
  Short call(K2) at expiry: -max(S_T - K2, 0)  
  Long put(K2) at expiry:  max(K2 - S_T, 0)
```

Sum: $K2 - K1$, deterministic, regardless of S_T

This is independent of the underlying. It pays exactly $K2 - K1$ at expiry. By no-arbitrage, the price today must be:

Box value today = $(K2 - K1) \times e^{-rT}$

If the box trades at $\$X$ today and pays $\$(K2 - K1)$ at expiry T , the implied financing rate is:

$\text{implied_r} = \ln((K2 - K1) / X) / T$

This is *the cost of synthetic lending*. By selling a box (collecting cash today, paying the difference at expiry), the seller is *borrowing* at the implied rate. By buying a box, the buyer is *lending* at the implied rate.

6.2 6.2 SPX vs SPY — The Critical Distinction

SPX boxes (and XSP boxes, the mini SPX) use **European-style** options. They cannot be early-exercised. The payoff at expiry is exactly $K2 - K1$ — no path dependence, no early-exercise risk. This makes SPX boxes the *cleanest* synthetic financing instrument in the world.

SPY boxes (and most equity boxes) use **American-style** options. The short put at $K1$ can be exercised by the counterparty if it goes deep ITM. The short call at $K2$ can be exercised against you if the stock blows past $K2$ and there is a dividend coming. This injects path-dependence and *early-assignment risk* — the box may pay out unexpectedly early, and the borrowed funds may need to be returned at an inopportune moment.

Therefore: **SPX boxes are the institutional-grade synthetic financing instrument**. SPY boxes are not.

A second critical distinction is **tax**. SPX is a Section 1256 contract — 60% long-term capital gain, 40% short-term, regardless of holding period. SPY is regular short-term gain/loss treatment. For the typical box-as-financing trade (held < 1 year), the tax difference is significant:

- SPX box financing: blended tax rate $60\% \times \text{LTCG} + 40\% \times \text{OIT}$.
- SPY box financing: 100% short-term ordinary income tax rate.

For a high-bracket retail user, the SPX 60/40 saves roughly 7-10% of gains — meaningful on multi-year financing trades.

6.3 6.3 The Box Loan: Synthetic Borrowing

Suppose a retail trader wants to borrow \$100,000 for two years at the lowest possible rate. Available channels:

Channel	Effective rate	Notes
Margin loan at IBKR Pro	SOFR + 1.5% (~6.8%)	Variable, tax-deductible only if used for investment
Margin loan at Schwab	SOFR + 5% (~10.3%)	Default retail rate
Personal loan / HELOC	7-9%	Fixed, depends on credit
SPX box (sell)	SOFR + 0.20-0.50%	Locked-in, tax-advantaged

The SPX box implied rate is competitive with the absolute lowest institutional financing rates because the box is fully collateralized — the counterparty’s risk is zero.

Concrete worked example, two-year borrow.

- Today: 2026-05-02
- Target maturity: 2028-05-19 (753 days)
- Desired notional borrow: \$100,000

Construct the box on SPX with $K_2 - K_1 = \$100$. Need $K_2 - K_1$ to give \$100,000 at expiry, so 1000 contract pairs (each contract is 100 multiplier, $K_2 - K_1 = \$100$ means \$10,000 per contract pair, so 10 pairs gives \$100,000).

Wait — recheck: SPX is a \$100 multiplier. If $K_1 = 4000$, $K_2 = 5000$, $K_2 - K_1 = 1000$. Per contract, payout = $\$1000 \times 100 = \$100,000$ per single box. So a **single 1000-point-wide SPX box pays \$100,000 at expiry**.

Trade:

Sell box on SPX, expiry 2028-05-19:

- Buy 4000 call (Long-dated SPX call, deep ITM)
- Sell 4000 put (Long-dated SPX put, deep OTM)
- Sell 5000 call (Long-dated SPX call, OTM)
- Buy 5000 put (Long-dated SPX put, deep ITM)

Selling a box means *collecting* cash today and paying \$100,000 at expiry. The net credit received is roughly:

Box price = $\$100,000 \times e^{-r \times 753/365}$

If the implied financing rate is 4.45% (typical for SPX boxes in a 5.30% SOFR environment with 80-90 bps below SOFR due to massive supply of natural lenders):

Box value = $\$100,000 \times e^{-0.0445 \times 753/365}$
 $= \$100,000 \times e^{-0.0918}$
 $= \$100,000 \times 0.9123$
 $= \$91,230$ received today.

Two years later, the trader pays \$100,000. The net interest cost: \$8,770 over 753 days, or 4.45% annualized. Compare to the \$20,560 cost of a 10.3% Schwab margin loan over the same period — the box saves \$11,790.

Plus the 60/40 tax treatment: \$8,770 of cost gets 60% LTCG and 40% OIT (interest treatment for box is slightly different — see 6.5 below). The effective after-tax cost is roughly 3.3% in a high-bracket account.

6.4 6.4 Margin Treatment of Boxes

A *long* box (buying — lending) has zero risk: it’s fully cash-collateralized at expiry by the $K_2 - K_1$ payoff. Brokers should require zero margin. Reality:

- IBKR: charges full long-option premium as margin. Zero risk-based margin.
- Schwab: variable, sometimes treats it as four separate options.
- Tastytrade: explicit box-spread support, near-zero margin.
- Robinhood: does not support boxes natively (each leg separately).

A *short* box (selling — borrowing) is *not* zero-risk in the path-dependence sense if the legs are American-style. SPX is European, so it is in fact zero path-dependent risk. Brokers ought to charge only the difference between the credit received and the eventual payoff. In practice:

- IBKR Pro: full SPAN-based margin, very efficient. ~\$5-10k initial on a \$100k box.
- Schwab: requires full credit-spread margin on each pair — \$200,000 initial. Inefficient.
- Tastytrade: SPAN-based, similar to IBKR.
- Robinhood: cannot short box.

The variation in broker margin treatment is the *single biggest* friction in the retail box-loan market. The same trade at IBKR is 1/40 the margin requirement of Schwab.

6.5 6.5 The 2018 Robinhood “Box” Incident

In late 2018, a Reddit user (u/ControlTheNarrative) posted that he had used Robinhood’s then-buggy options margin engine to write SPX box spreads with no margin requirement, effectively giving himself an unlimited interest-free loan. His position size eventually reached around \$1m. The post went viral.

Robinhood at the time used a margin engine that: 1. Treated each leg of the box independently for margin purposes. 2. Recognized the four legs as four “credit spreads” with cancelling margin — netting to zero. 3. Did not flag the resulting position as a synthetic loan.

The user used the freed margin to buy lottery-ticket call options. The trade ultimately blew up (the wrong way — the lottery tickets expired worthless, leaving him with the box short).

What happened next:

- Robinhood patched the margin engine within weeks.
- The CFTC, SEC, and FINRA reviewed. No criminal charges (the user did not exploit an undocumented bug; he used the system as it was configured).
- Robinhood disabled multi-leg options entirely for several months while patching.
- The post became a cultural touchstone of “what retail can do when the institutional toolkit becomes available.”

The technical analysis confirmed: SPX boxes are the cleanest synthetic loan available, and at certain brokers with insufficient margin engines they were temporarily *free*. The episode hardened broker margin engines across the industry.

6.6 6.6 Pin Risk and Box Mechanics Near Expiry

A box held to expiry on SPX settles cleanly at $K_2 - K_1$ because SPX is European cash-settled. No pin risk. But on SPY (American physically-settled):

- If S_T pins at K_1 or K_2 on expiration day, options may or may not be exercised.
- If you are long the K_2 put and S_T is exactly K_2 , exercise is borderline.
- This creates a small residual risk on equity boxes near expiry.

For SPX specifically, this is a non-issue. The cash-settlement at the AM open is unambiguous. The only residual risk on an SPX box is *counterparty* risk during the multi-day exercise period, which is essentially zero given OCC clearing.

6.7 6.7 Rolling Boxes

For indefinite financing, traders roll boxes:

Step 1: Sell SPX box, 2-year tenor, receive \$91,230 today.
Step 2: 23 months later, when 1 month to expiry, sell new SPX box, 2-year tenor.
Step 3: Use the \$91,230 from Step 2 to repay the original Step 1 box at the imminent maturity.
Step 4: Continue indefinitely.

The roll cost is the difference in implied financing rates between consecutive boxes. In a steepening yield curve, rolling becomes more expensive over time. In a flat or inverting curve, it can be very cheap.

6.8 6.8 Why Retail Brokers Treat Boxes Differently

The variation comes from regulatory and risk-management posture:

Broker	Risk model	Box treatment
IBKR Pro	SPAN portfolio margin	True risk-based — boxes nearly free margin
Tastytrade	Customized SPAN	True risk-based — boxes nearly free
Charles Schwab	Reg-T strategy-based	Treats each pair as a separate spread — high margin
Fidelity	Reg-T strategy-based	Similar to Schwab
Robinhood	Reg-T plus restrictions	Short boxes typically blocked
E*TRADE	Hybrid	Full margin on shorts

Reg-T (Federal Reserve Board Regulation T) sets minimum margin requirements but does *not* mandate strategy-based vs portfolio-based treatment. Brokers can be more conservative than Reg-T, and most retail brokers are. IBKR Pro’s “Portfolio Margin” requires \$110k+ in account equity but unlocks SPAN-level margining — making it the only mainstream retail broker where SPX boxes are economically attractive at scale.

6.9 6.9 The Tax Treatment of Boxes — A Closer Look

The IRS has had a long, occasionally adversarial relationship with box spreads. Three rulings have shaped current practice:

Revenue Ruling 78-414 (1978): The original IRS position that a box is *economically equivalent to a loan* and should be taxed accordingly. This implied that the “interest” portion of a box should be ordinary income/expense, not capital.

Revenue Ruling 2003-7 and Subsequent Guidance: The IRS clarified that *Section 1256 contracts* (broad-based index options including SPX, NDX, RUT) follow the 60/40 rule regardless of the trader’s intent — including for box spreads.

Practitioner Position (Current): Most tax practitioners and major brokers report SPX box gains/losses as 60/40 capital gains. This is the favorable treatment. The IRS has not actively challenged this in most cases.

The ambiguity matters in scale. A \$1m box held for one year generating \$50k in implied interest:

Treatment A (60/40 capital gain): \$30k LTCG + \$20k STCG
At 23.8% LTCG and 37% STCG rates:
Tax: \$7,140 + \$7,400 = \$14,540
Effective rate: 29.1%

Treatment B (ordinary interest): \$50k ordinary income
At 37% rate:
Tax: \$18,500
Effective rate: 37%

Difference: \$3,960 per \$1m box per year

For a retail user using SPX boxes for ongoing financing, the 60/40 treatment is worth approximately 8% of the implied financing cost annually — a meaningful edge.

6.10 6.10 Why the Box Implied Rate Is Below SOFR

A surprising fact: SPX box implied financing rates have *consistently traded below SOFR* by 30-100 bps over the past decade. The reasons:

1. **Natural lender supply.** Pension funds, insurance companies, and large family offices hold cash that they want to deploy at short tenor with zero risk. Buying boxes is the cleanest, lowest-risk way to lend at fixed terms.

2. **Tax preference.** The 60/40 treatment makes box-as-lending more attractive to taxable lenders than T-bills (100% ordinary).
3. **Regulatory capital.** Banks holding SPX boxes hold them at near-zero RWA (the structure is fully cash-collateralized at expiry). This makes box-buying balance-sheet-efficient for banks.
4. **Liquidity preference of borrowers.** Retail and small institutional borrowers willing to pay a premium for term-locked, mark-to-market-clean financing. They accept a slightly higher rate than they'd theoretically pay for collateralized repo because of the operational simplicity.

The result: SPX box implied rates are an attractive borrowing rate for sellers (typically below SOFR + 50 bps) and an attractive lending rate for buyers (typically above the cleanest T-bill rate). The market clears at a point where both sides win versus their alternatives.

6.11 6.11 Box Liquidity and Bid/Ask

The bid/ask on a box depends on the strikes chosen:

- **Standard 100-point boxes** (e.g., 4500/4600 SPX): tight bid/ask, often \$0.10-\$0.30 wide on a \$100 box. This is the institutional sweet spot.
- **Non-standard widths:** wider markets, sometimes \$1.00+ on a \$50 box.
- **Far-from-spot strikes:** wider markets, especially when the strikes are >20% from spot.
- **Long-dated (>2 years):** tighter the further out, sometimes counter-intuitively, because LEAPS market makers rely on standardized inventory.

Practical implications for retail:

- Use 100-point or 1000-point widths, near-the-money or even ITM. ITM strikes are *cleaner* because the option premiums are larger and bid/asks narrower.
- Avoid 50-point widths unless you can guarantee execution — slippage can eat 10+ bps.
- Use limit orders, not market orders. Box markets are dealer-quoted; market orders get filled at the worst dealer's quote.

A typical execution flow:

1. Identify target tenor (e.g., June 2027 for ~24-month box).
2. Pick strikes $K_1 = 4000$, $K_2 = 5000$ (1000-pt width on SPX = \$100k payoff).
3. Compute fair value:

$$\text{box_fv} = \$100,000 \times e^{-\text{implied_rate} \times T}$$
 At implied rate 4.45%, $T = 24/12 = 2.0$:

$$\text{box_fv} = \$100,000 \times e^{-0.0890} = \$91,476.$$
4. Quote a sell order at \$91,300 (slightly below fair to attract a fill).
5. If filled at \$91,400, your effective borrow rate is:

$$\text{rate} = -\ln(\$91,400 / \$100,000) / 2.0 = 4.49\%$$
6. Plan rollover ~30 days before expiry to avoid pin-day liquidity issues.

6.12 6.12 The “Stress” Box and What Happens in March 2020

During March 2020, SPX option markets temporarily dislocated. Box spread implied rates briefly traded above 7% — well above SOFR (which was 1%). The dislocation reflected:

1. Margin pressure on dealer balance sheets (massive vol increase).
2. Option market makers widening quotes due to bid/ask uncertainty.
3. Speculative demand for cash by hedge funds facing redemptions.

Traders who *bought* boxes during this window earned 7%+ “lending” returns on a fully-collateralized basis for several days. This was a textbook arbitrage opportunity that closed within ~10 trading days as dealers regained capacity.

The lesson: box rates are normally *below* SOFR but can dislocate above SOFR during liquidity crises. A retail trader monitoring box rates can identify these windows and lend (buy boxes) during stress for outsized returns. This requires sufficient capital to deploy on short notice and a disciplined understanding of pin and counterparty risk.

6.13 The Retail Box-As-Loan in Practice — A Detailed Walkthrough

A retail user “TraderA” wants to access \$50k for renovations, terms unclear (maybe 18 months, maybe 36 months). Available options:

Option	Effective rate	Setup cost	Flexibility
HELOC	7.5%	\$500-1500	High; close anytime
Personal loan	9.0%	\$0	Low; fixed term
401(k) loan	6.5% (paid to self)	\$0	Risk if employment ends
Margin (Schwab)	10.3%	\$0	High; rate variable
SPX box (sell)	~4.5%	\$5-15 fee	High; can roll or close

TraderA has \$200k at IBKR Pro (portfolio margin). TraderA decides to sell a 24-month SPX box.

Execution (Day 0):

Trade: Sell SPX June 2028 4500/5000 box

Buy 4500 call: ~\$1190.50 (deep ITM)

Sell 4500 put: ~\$36.20 (deep OTM)

Sell 5000 call: ~\$700.30 (OTM)

Buy 5000 put: ~\$45.10 (deep ITM in inverted sense)

Net credit per box: $1190.50 - 36.20 - 700.30 + 45.10 = (-498.90)$

Wait - let me recheck the legs:

Long 4500 call: pays \$1190.50

Short 4500 put: collects \$36.20

Short 5000 call: collects \$700.30

Long 5000 put: pays \$45.10

Cash flow: $-1190.50 + 36.20 + 700.30 - 45.10 = -499.10$ (net debit)

I made an error. Let me redo: a *short* box is constructed differently to receive credit. The structure is *short* call at K1 and *long* call at K2 — opposite of “long” box.

SHORT BOX = “selling” the box = receiving cash today, paying K2-K1 at expiry.

Legs of a short box:

Sell call at K1 (collect)

Buy put at K1 (pay)

Buy call at K2 (pay)

Sell put at K2 (collect)

For SPX K1=4500, K2=5000, ~24 months:

Short 4500 call: collect ~\$1190

Long 4500 put: pay ~\$36

Long 5000 call: pay ~\$700

Short 5000 put: collect ~\$45

Net credit: $1190 - 36 - 700 + 45 = \$499$ per “unit”

SPX multiplier 100 → \$49,900 cash received per 1 contract.

At expiry, payoff to box-seller: $-(\$K2 - \$K1) = -\$50,000$

Net economic: \$49,900 received today, \$50,000 paid at expiry.

Implied financing rate: $-\ln(\$49,900 / \$50,000) / 2.0 = 0.10\%$ per annum?

This implies a 10 bps rate, which is too low — clearly the prices are unrealistic for illustration. In reality, with implied rate 4.45% and \$50k payoff at 24 months:

Box value today = $\$50,000 \times e^{(-0.0445 \times 2.0)}$

$$\begin{aligned}
&= \$50,000 \times e^{(-0.089)} \\
&= \$50,000 \times 0.9148 \\
&= \$45,742
\end{aligned}$$

So the trader receives \$45,742 today and pays \$50,000 in 24 months.
Total interest cost: \$4,258 over 24 months = \$2,129/year on \$45,742.
Effective rate: 4.65% (roughly matching implied 4.45% net of bid/ask).

Margin (Day 0 to expiry):

At IBKR Pro with portfolio margin, the SPAN initial margin on a short SPX box is approximately *zero* — the structure has no path-dependent risk on European options.

In practice IBKR may require ~\$1k-\$3k of “minimum equity” margin as a buffer. This is roughly 5% of notional, vastly less than Reg-T strategy-based margin of ~\$50k.

Holding period (Day 0 to ~Day 720):

TraderA uses the \$45,742 cash for renovations. The position sits on the books, marking to market daily. Daily P&L fluctuations in SPX move all four legs in offsetting ways — net P&L should be small (driven only by the small change in implied financing rate as macro rates move).

Pre-expiry roll (Day 690, 30 days before expiry):

TraderA decides to roll for another 24 months. Sell new SPX June 2030 4500/5000 box. Use credit to buy back the original.

At final expiry (Day 720):

The original box settles at \$50,000 cash (SPX is European cash-settled). TraderA pays \$50,000.

Net position after roll/expiry:

If TraderA rolls indefinitely, the only cost is the *interest* implicit in each box plus the bid/ask friction on each roll. The position is functionally a permanent ~4.5% loan.

Tax treatment:

The realized gain/loss on the box closeout is reported as Section 1256 — 60% LTCG, 40% STCG. For TraderA in the 32% bracket: - \$4,258 of interest expense over 24 months. - Treated as 60% LTCG (15% rate after high-income surcharge 23.8%) and 40% STCG (32%). - Effective tax on the loss: $60\% \times 23.8\% + 40\% \times 32\% = 27.1\%$ deductible. - After-tax cost of borrowing: $\$4,258 \times (1 - 0.271) = \$3,104$ over 24 months. - After-tax annualized rate: ~3.4%.

Compare to a HELOC at 7.5% (interest fully deductible against home if itemized): after-tax 5.5% in the 32% bracket. The SPX box wins by ~2 percentage points after tax.

This calculation generalizes. For a retail user with portfolio margin access, the SPX box loan is the cheapest unsecured-equivalent financing available.

6.14 6.14 The Risks of Box Lending vs Box Borrowing

Box buyers (lenders) and sellers (borrowers) bear different risks.

Box buyer (lender):

Risks:

1. Pin risk on equity boxes (not SPX) - if S_T pins at K_1 or K_2 , payoff may differ slightly from $K_2 - K_1$.
2. Counterparty risk (zero with OCC clearing).
3. Reinvestment risk if early-exercised against (American boxes only).
4. Opportunity cost if rates rise above the locked-in rate.

Mitigations:

- Use SPX (European) only.
- Hold to expiry - early-close incurs bid/ask.
- Diversify across tenors to mitigate rate-rise opportunity cost.

Box seller (borrower):

Risks:

1. Margin calls if portfolio margin engine misprices the box.
2. Bid/ask drag on rolls (if rolling indefinitely).
3. Tax volatility - Section 1256 mark-to-market at year-end means pre-realization tax events.
4. Pin risk on equity boxes (not SPX).

Mitigations:

- Use SPX with portfolio margin only.
- Roll well before expiry (30-60 days) to avoid pin-day illiquidity.
- Keep a cash buffer for unexpected mark-to-market margin needs.

The pin risk for *SPX is zero* — SPX cash settles unambiguously based on the AM open opening print. This is one of several reasons SPX is the institutional standard for box financing.

6.15 A 2-Year Simulation With Bid-Ask Drag

The previous sections describe the SPX box mechanics in idealized form. This subsection runs a complete 2-year simulation, including realistic bid-ask drag, margin treatment changes during the holding period, and tax events. The goal is to show what the trade actually looks like end-to-end, and where the friction matters.

The setup. A retail user “Sarah” has \$300,000 at IBKR Pro with portfolio margin enabled. She wants to borrow \$100,000 for two years to fund a renovation. She decides to sell a 24-month SPX box.

Day 0 (May 2, 2026): Entry.

Sarah pulls up the SPX option chain for May 19, 2028 expiry. She picks $K1 = 5500$, $K2 = 6000$ (a 500-point width box on a 1000-point notional, which gives \$50,000 payoff per contract). She wants \$100,000, so she will sell 2 contract pairs.

The mid-market box value at entry: $\$50,000 \times e^{-0.0445 \times 2.05} = \$50,000 \times 0.9128 = \$45,640$ per box. Sarah is selling 2 boxes; mid-market credit = \$91,280.

But mid is not what she'll get. The bid-ask on a deep-strike SPX box at 24 DTE 2 years out is roughly \$0.40 wide on the box (split across the four legs). At 100x multiplier and 2 boxes, this is $\$0.40 \times 100 \times 2 = \80 of drag. Not enormous, but real.

Sarah enters a limit order at \$91,200 (slightly below mid to attract a fill). The order fills 11 minutes later at \$91,180. She has now received \$91,180 in cash and committed to paying \$100,000 at expiry.

Implied financing rate: $-\ln(91180/100000) / 2.05 = 4.50\%$ per annum. Slightly above the 4.45% mid because of the bid-ask paid.

Margin: IBKR's portfolio margin engine assigns SPAN margin of approximately \$4,500 to the position. Sarah has \$300,000 in equity, so this margin is comfortably covered.

Day 30: First mark-to-market.

SPX has rallied 4% over the month. The box's mark-to-market value has changed slightly — not because of SPX moves (the box is delta-neutral), but because the implied financing rate has moved. Suppose SOFR has risen 25 bps over the month (a hawkish surprise from the Fed). The implied box rate also moves up roughly in tandem. The new mid-market box value:

$\text{new_mid} = 50,000 * \exp(-0.0470 * (2.05 - 0.083)) = 50,000 * \exp(-0.0925) = 50,000 * 0.9116 = \$45,581$ per box.
For 2 boxes: \$91,162.

Sarah's mark-to-market value of the *short box position*: she received \$91,180 at entry; the position now requires \$91,162 to close. Her unrealized P&L: +\$18. She is mildly profitable on a marked-to-market basis because she locked in a slightly lower rate than the current market.

Day 30 is also the first day she begins paying interest on the implicit loan. The implicit cost is approximately $\$91,180 * 4.50\% / 365 = \11.24 per day, or \$337 over the 30 days. This is amortized into the position; she does not see it as cash, but it accumulates.

Day 90: Tax event.

Year-end approaches. Section 1256 contracts (which include SPX options) are marked to market on December 31. The position must be re-valued for tax purposes.

Suppose on December 31, 2026, the position's mark-to-market value is $\$46,500 \text{ per box} \times 2 = \$93,000$. The position cost basis is $\$91,180$ (the credit received). The mark-to-market unrealized loss: $\$93,000 - \$91,180 = \$1,820$ (it costs more to close now than at entry; this is unrealized loss).

For tax purposes, this $\$1,820$ is recognized as a Section 1256 loss split 60/40: - $\$1,092$ long-term capital loss (60%) - $\$728$ short-term capital loss (40%)

These losses can offset other Section 1256 gains. If Sarah has no offsets, the loss goes against ordinary income (limited to $\$3,000/\text{year}$ in net capital losses, with the rest carried forward).

This is the "tax volatility" risk of Section 1256 contracts. The loss is *unrealized* — Sarah has not actually closed the position — but the tax code treats it as realized at year-end. If she has a large position and the position is deep in the money on a marked-to-market basis, she could owe substantial tax on phantom income.

In Sarah's case, the loss is small ($\$1,820$) and likely offsetable. But for a $\$1\text{m}+$ box position, the year-end mark could swing tens of thousands of dollars in either direction, creating real tax-management complications.

Day 365 (one year mark): Decision point.

Sarah considers rolling the position. The new financing curves:

- Original box: 12 months remaining to expiry. Current mid: $\sim \$95,890$.
- New 24-month box at the same strikes ($K1=5500$, $K2=6000$), expiry May 2029: implied rate around 4.20% (the curve has flattened slightly). New box mid: $50,000 * \exp(-0.0420 * 2.0) = 50,000 * 0.9197 = \$45,985$ per box, or $\$91,970$ for two.

If Sarah rolls now: close the original at $\sim \$95,890$ (cost), open the new for $\$91,970$ (credit). Net cash flow: $-\$3,920$. This is the explicit cost of rolling early; she has effectively borrowed for one extra year at an accelerated rate.

She decides not to roll. The original position will expire in 12 months at $\$100,000$ cost. She'll roll closer to expiry.

Day 720 (just before expiry): The roll.

30 days before expiry, Sarah rolls. Her process:

1. Sell new SPX box, May 2031 5500/6000 strikes. New rate is around 4.10%. New box value: $50,000 * \exp(-0.0410 * 2.0) = 50,000 * 0.9213 = \$46,065$ per box, or $\$92,130$ for two.
2. Buy back the original box. With 30 days to expiry, the original box trades very close to its $\$50,000$ value: approximately $50,000 * \exp(-0.045 * 30/365) = 49,815$ per box, or $\$99,630$ for two.
3. Net cash flow on the roll: $\$92,130$ (received) - $\$99,630$ (paid) = $-\$7,500$ cash out.

Plus the daily bid-ask on each leg (roughly $\$200$ of drag total over the four-leg roll).

Sarah's cumulative cash flow over 24 months: - Day 0: $+\$91,180$ received. - Day 720: $-\$99,630$ paid (to close the original). - Day 720: $+\$92,130$ received (new box). - Net at this point: $+\$83,680$.

She still has 24 more months of "loan" outstanding via the new box. The renovation funds came from the day-0 receipt minus the day-720 payment, totaling $-\$8,450$ over the period that the original box was alive — which corresponds to a 4.50% implied rate on the $\$91,180$ borrowed for 24 months ($\$91,180 * 4.50\% * 2 = \$8,206$, very close).

Day 1440 (4 years out, the new box matures): Final unwind.

Sarah decides to exit the box loan structure entirely (renovation done, no further need). She closes the new box near its $\$100,000$ expiry value. Net cumulative cash flow over 4 years:

Day 0: $+\$91,180$ (sold first box)
Day 720: $-\$99,630$ (bought back first box)
Day 720: $+\$92,130$ (sold second box)
Day 1440: $-\$99,820$ (bought back second box, close to expiry value)

=====

Cumulative: $-\$16,140$ over 48 months

Her total interest cost: \$16,140. Average financing rate over 4 years: roughly 4.40% per annum (slightly below the locked-in rates because of the modest mark-to-market gains during favorable periods).

Compared to a HELOC at 7.5% over the same period: \$30,000 in interest. The box loan saved her roughly \$14,000 over 4 years.

After-tax: the Section 1256 60/40 treatment cuts the after-tax cost. Sarah is in the 32% federal bracket. The net after-tax cost: roughly \$11,900 over 4 years, or 3.2% per annum after-tax.

This is the realistic picture. The implied 4.45% rate from the abstract calculation does not perfectly match reality because of bid-ask drag, mark-to-market timing, and the rolls. But the rate the user actually pays is roughly within 5-10 bps of the implied rate, which is much cheaper than the alternatives.

6.16 The Failure Case — When Boxes Go Wrong

The 2-year simulation above assumed Sarah's plan worked. What does failure look like?

Failure case 1: Forced unwind during stress.

It is March 2030. There is a sudden, unexpected vol spike (the “next March 2020”). VIX has jumped from 15 to 65 overnight. SPX has dropped 12% in two days. Sarah's broker, IBKR, is experiencing portfolio margin stress and has revised its margin model.

Although Sarah's box position is *delta-neutral and approximately rate-neutral*, the broker's margin engine in stress conditions widens the SPAN parameters across all four legs. The position now requires \$25,000 in margin instead of the original \$4,500. Sarah's account, which has \$280,000 in equity (after market losses on her other positions), has \$50,000 of margin already used elsewhere. The new requirement of $\$25,000 + \$50,000 = \$75,000$ can be met, but barely.

If the market move had been larger, or if Sarah had been holding less equity, she could face a margin call and be *forced to close* the box at the worst possible moment. The bid-ask on the box during stress widens dramatically — a normally \$0.20-wide market becomes \$2-3 wide. Closing the box at this point could cost an extra \$400-800 in slippage on top of the normal closing costs.

This is a real failure mode for box financing. The box is fundamentally low-risk, but the combination of mark-to-market margining and stress-period bid-ask widening means the holder must have plenty of buffer capital to weather storms.

Failure case 2: Implied rate inversion.

Suppose during the 2-year holding period, the SPX box implied rate *rises* substantially (from 4.45% at entry to 7.0% at month 12). On a mark-to-market basis, the original box (locked at 4.45%) is now significantly more expensive to close than at entry. Sarah's unrealized loss on the position:

$\text{new_mid} = 50,000 * \exp(-0.07 * 1.05) = 50,000 * 0.9293 = \$46,465$ per box, or \$92,930 for two. Original credit: \$91,180. Unrealized mark-to-market loss: \$1,750.

Year-end Section 1256 mark-to-market means Sarah recognizes a \$1,750 loss on the position even though she hasn't closed it. This loss is offsetable, but if she *needs* to recover the cash (renovation overrun), she'd have to close at a real \$1,750 loss.

This is the “implied rate inversion” failure. The box rate is volatile in real markets; locked-in rates can become disadvantageous if the market moves. The lock is good when rates rise (Sarah keeps her low rate); it is bad when she's marked to market on rising rates.

Failure case 3: Pin risk on equity-style boxes.

Sarah's choice of SPX (European, cash-settled) avoids this. But a retail user who chooses SPY (American, physically settled) instead, perhaps because their broker doesn't support SPX boxes, can face pin risk near expiry. If the underlying is exactly at one of the box strikes on expiration, ambiguous exercise creates a residual risk. This is rare (typical pin probability is <1%) but real, and can cost a few hundred dollars per pin event.

The lesson: SPX boxes are nearly riskless when run with discipline. SPY or other equity boxes have meaningful residual risks that the retail user should avoid. The discipline of using only SPX (or XSP, the mini-SPX) eliminates the pin issue.

6.17 6.17 References

- Hull, J. (2015), *Options, Futures, and Other Derivatives*, 9th edition, Chapter 11 (Box Spreads).
 - IRS Publication 550, Section 1256 Contracts.
 - Reddit r/wallstreetbets, “1R0NYMAN” thread (2018) — primary documentation of the Robinhood box incident.
 - CBOE (2018), “Box Spreads on SPX Index Options”, educational paper.
 - van Binsbergen, J., Diamond, W., and Grotteria, M. (2022), “Risk-Free Interest Rates”, *Journal of Financial Economics* 143(1): 1-29. Foundational paper on the box-implied rate as a clean read on SOFR-equivalent risk-free rate.
 - Longstaff, F.A. (2018), “Valuing Thinly Traded Assets”, *Management Science*.
 - Mehra, R. (2008), “Handbook of the Equity Risk Premium”, North-Holland Handbooks in Finance.
-

7 Chapter 7 — Convertible Bond Arbitrage

7.1 7.1 The Convertible Bond Structure

A convertible bond is a corporate bond with an embedded option to convert into the issuer’s stock. Structurally:

`Convertible bond = Straight corporate bond + Long call on the stock`

The bond pays a coupon (usually below the straight-bond rate, because the embedded call is valuable). The conversion ratio specifies the number of shares the bond converts into. The implicit strike is the bond’s face value divided by the conversion ratio.

Example: a \$1,000 face convertible bond convertible into 25 shares is implicitly long a call at \$40/share (1000/25). If the stock trades above \$40, conversion is in-the-money. The bond’s value is bounded below by the straight-bond floor and rises with the stock when the embedded call has intrinsic value.

7.2 7.2 The Classical Convert Arb Trade

The strategy: buy the convertible bond, hedge by shorting the underlying stock at delta-equivalent quantity. Profit from:

1. **Volatility.** The embedded call benefits from gamma rebalancing in the short stock hedge.
2. **Credit improvement.** If the issuer’s credit improves, the bond’s price rises faster than the hedge predicts.
3. **Dividend capture.** The convertible holder doesn’t receive dividends; the short does — but the stock loan fee is typically much lower than the dividend.
4. **Cheap volatility.** Convertible bonds are often issued at implied volatilities below the listed-options market — sometimes 5-10 vol points below — because issuers can structure them with covenant restrictions that compensate.

A typical convert arb desk runs hundreds of positions, each delta-hedged daily. The blended return historically has been 6-10% per annum with volatility of 5-8% — Sharpe ratios in the 0.8-1.5 range.

7.3 7.3 Where Securities Lending Enters

Convert arb requires *shorting* the underlying stock. The short fee is a direct cost. For high-fee names — names with structural short interest, often the same companies issuing converts because they are stressed credits — the short fee can erase the convert arb edge.

The relationship is precise: a convertible bond’s market price reflects an *implicit short fee* the holder needs to absorb. If the stock is special at 200 bps, the convertible should trade approximately 200 bps × duration cheaper than its frictionless theoretical value. Convertible arb funds price this in:

`Convertible "cheapness" = Theoretical value - Market price
duration × (borrow_fee + credit_spread + transaction_friction)`

The cheaper the convert, the higher the borrow fee implies. This makes convertible bond market prices a *clean read on borrow forecasts*. A convert arb desk uses this signal:

- Convert is 8% cheap to model → borrow fee is implicitly forecast at high level.
- Convert is 0% cheap → borrow fee is benign.

This is one of several “implied borrow” markets traded across asset classes.

7.4 7.4 The Death Spiral Convertible

A particularly dangerous convertible structure: the death spiral. The conversion ratio is *not fixed*; it adjusts based on the stock price. As the stock falls, the conversion ratio rises (more shares delivered), creating dilution that pushes the stock further down. The convertible holder is short volatility and effectively short equity in a self-reinforcing way.

For a delta-one or convert arb desk, the death spiral provides: - High volatility on the underlying. - Massive securities-lending demand from short sellers. - Borrow fees often spiking 50-100% per annum.

The desk’s optimal play: lend the stock to short sellers (collecting massive fees), buy the convertible at a deep discount, and let the structure resolve. Risk: the underlying issuer goes bankrupt and the convert recovery is zero.

7.5 7.5 Worked Example

Issuer: a small biotech, “BIOX”. Issues a 5-year convertible at \$1000 face, conversion ratio 40 (implicit conversion price \$25). Stock at \$20. Coupon 4%. Implied vol on the embedded call (5-year ATM) is roughly 60% per the listed market, but the convert’s pricing implies 50% — 10 vol points cheap.

Trade: - Buy 100 convertible bonds = \$100,000 face value. - Delta of embedded option = roughly 0.45. - Convert delta-equivalent shares = $100 \times 40 \times 0.45 = 1,800$ shares. - Short 1,800 shares of BIOX at \$20.

Daily P&L drivers: - Convertible price changes (function of stock, vol, credit). - Short stock P&L (offsets first-order delta). - Coupon income (4% of \$100k face = \$4k/year). - Short borrow cost (BIOX is special at 8% per annum: $\$20 \times 1800 \times 8\% = \$2,880$ /year). - Short stock dividend (if any — typically \$0 for BIOX-style early biotechs).

Annualized expected return:

+ Convert vol harvest:	\$5,000 (gamma rebalancing on 10-vol-pt edge)
+ Coupon income:	\$4,000
+ Credit improvement:	\$2,000 (if rating moves)
- Short borrow cost:	-\$2,880
= Net expected:	\$8,120 (~8.1% on \$100k notional)

The convert arb trade is *only* profitable above the borrow fee. Hard-to-borrow names are the primary alpha source, but they are also the names where the trade is hardest to size and run.

7.6 7.6 The Convertible Issuance Cycle

A subtle point that connects to the delta-desk’s sec-lending business: when a company issues a convertible bond, hedge funds aggressively buy and immediately short-hedge. This *creates* the demand for stock borrow.

The full cycle:

Day 0: Company announces \$500m convertible offering, conversion premium 30%.
 Day 1: Underwriters market the deal. Convert arb funds indicate interest.
 Day 2: Pricing - bonds priced at par with 4% coupon. Funds receive allocations.
 Day 3: Funds short the underlying stock at ~50% delta hedge.
 This creates immediate, large demand for stock borrow.
 Borrow fees on the issuer's stock spike from <1% to 5-10% in the days following issuance.
 Day 30+: Borrow stabilizes as the convert arb book becomes mature.
 The continuous gamma trading of the convertible against the short maintains the delta hedge at ~50% delta. Volume on stock-loan remains elevated for the bond's life.

For a delta-desk, this is gold: every convertible issuance creates a multi-year, sticky lending business on the issuer’s stock. The desk that runs the underwriting also runs the lending matched-book — an integrated business model.

7.7 7.7 The Death-Spiral Convertible Worked Example

Consider “BIOXYZ”, a small biotech trading at \$5 after a Phase 2 failure. Cash-burn negative, runway ~12 months. Issues \$50m of convertible debt with the following structure:

Convertible bond:

Face: \$50,000 per bond

Maturity: 3 years

Coupon: 5% (paid semi-annually)

Conversion: into shares at the LOWER of:

(a) \$5.00 fixed price, OR

(b) 90% of the trailing 5-day volume-weighted average price (VWAP)

Forced conversion option for issuer if stock > \$7.50 for 20 consecutive days.

The conversion ratio adjustment is the death-spiral trigger. As BIOXYZ’s stock falls, the conversion ratio rises, more shares deliverable per bond, more dilution.

Round 1 (BIOXYZ at \$5). - Bond holders convert: receive $\$50,000 / \$5 = 10,000$ shares per bond. - Bond holders sell shares immediately (typical). - This *short pressure* drives stock to \$4.

Round 2 (BIOXYZ at \$4). - New conversion price: $90\% \times \$4 = \3.60 . - Bond holders convert another tranche: $\$50,000 / \$3.60 = 13,889$ shares. - Sell pressure intensifies. Stock drops to \$3.

Round 3 (BIOXYZ at \$3). - Conversion at \$2.70: 18,519 shares per bond. - Continuing dilution.

End-state: - The original \$50m face becomes hundreds of millions of shares. - The stock approaches zero. - Existing common shareholders are diluted to insignificance.

Delta-desk’s play: - Lending shares to bond holders for short hedging — astronomical fees (50-100% per annum). - Likely refusing to fund death-spiral converts on the bank’s balance sheet (reputation risk). - Trading the equity short itself with conviction (most banks have prop limits but research desks may publish negative views).

Risk to bond holders / delta-desk: - If BIOXYZ has an unexpected positive event (Phase 3 success, acquisition), the stock can spike, the conversion economics inverts (conversion price fixed at \$5 ceiling becomes attractive), and the spiral reverses. - If the company simply runs out of cash and bankrupts, conversion is moot — bond holders are back-of-line creditors.

Worked numerical P&L for a \$1m bond holder over 18 months:

Initial position: \$1m of bonds at par = 20 bonds.

Initial delta hedge: short 100,000 shares at \$5 = \$500k short.

Quarter 1 (BIOXYZ at \$4):

Converted 5 bonds → received $13,889 \times 5 = 69,445$ shares.

Sold immediately at \$4: \$277,780.

Delta-rehedged: re-short 30k shares.

Quarter 2 (BIOXYZ at \$3):

Converted 5 bonds at new conversion price → received 92,594 shares.

Sold immediately: \$277,780.

Delta-rehedged.

Quarter 3-6 (BIOXYZ at \$2 → \$1 → \$0.50):

Continued conversions and selling.

Ultimate cash recovered: ~\$900k of \$1m face.

Plus coupons: $\$25k \times 6 = \$150k$ (to the extent they were paid).

Plus delta-hedge gains: ~\$300k from the falling stock short.

Total recovery: ~\$1.35m on \$1m investment, 35% return over 18 months.

Plus borrow income for lending shares to other shorts: ~\$50k.

Total: 40% over 18 months = ~25% annualized.

Risks borne:

- Coupon non-payment (company in distress).
- Conversion delay or operational issues.
- Sudden upward spike eliminating the short hedge.
- Bankruptcy: recovery to \$0.

The death-spiral is profitable for the convert holder *as long as the spiral plays out predictably*. The risk is the upside spike or sudden bankruptcy.

7.8 7.8 Modern Convertible Markets

In recent years (2020-2024), convertible issuance has been dominated by:

1. **Tech growth companies** (Tesla pre-2020, MicroStrategy 2021-2024, Coinbase, Robinhood). Coupons low (0-1%) because the embedded equity option is very valuable.
2. **Mature dividend payers using converts as financing**. Lower-vol issuers with high credit quality. Convert arb fund interest moderate.
3. **Distressed and SPAC-related issuance**. Higher-vol, more complex.

The MicroStrategy convertibles are particularly notable: MSTR has issued multi-billion-dollar zero-coupon convertibles to fund Bitcoin purchases. The convertibles' embedded call is extremely volatile (BTC vol ~80%). Convert arb on these has been the highest-margin trade in the convertible market in 2023-2024 — borrow on MSTR went to 5-15% per annum from lows of 1-2%.

7.9 7.9 The Historical Arc of Convert Arb — 1990s to Present

Convert arb's history is a parable of financial innovation, regulatory cycle, and capital migration. The strategy is a classic "free money" trade for those who got there first; over four decades it has been domesticated, exploited, blown up, and re-emerged in a new form. Understanding the arc helps explain why the modern convert market has the structure it does and why retail-accessible variants do not exist.

The 1990s — discovery and early profit. Convert arb emerged from the academic recognition that convertible bonds embed a long call option on the issuer's stock. A handful of hedge funds in the early 1990s (Jefferies-affiliated funds, Citadel's earliest convert pool, Lazard's hedge fund) systematically exploited the cheap-call mispricing by buying converts and shorting the underlying stock. Returns in the early years were spectacular: 12-18% per annum with low correlation to traditional equity strategies. The Sharpe ratios on these early books exceeded 2.0.

The trade was not widely understood. Most participants in the convertible bond market — corporate bond funds, insurance companies, retail buyers via mutual funds — did not realize the embedded call's true value. Convert arb funds bought at "100 cents on the dollar" and the implied call was being valued at maybe 60 cents per nominal call. The wedge was substantial, and the funds running the strategy grew rapidly through the mid-1990s.

1998-1999 — the LTCM shock and the early crowding. The Long-Term Capital Management collapse temporarily disrupted convert arb (LTCM had a small convert-arb book and its forced unwind compressed pricing across the segment). But the strategy recovered quickly, and by the late 1990s, convert arb hedge fund AUM had grown to \$20-30bn industry-wide.

The strategy was no longer "secret." More funds entered. The wedge began to compress. By 1999, returns had dropped to 8-12% per annum, still good but no longer the easy money of 1992-1995. The convert arb funds responded by leveraging up — running the strategy at 4-8x leverage to maintain return targets.

2000-2007 — the boom era. Convert arb scaled into a mature hedge fund category. Industry AUM peaked around \$80-100bn by 2007. The strategy was levered, run by 200+ funds, and supported by integrated convert-arb desks at every major investment bank. Returns averaged 6-10% per annum but with thin Sharpe ratios (~1.0-1.4) because the wedge had been bid down.

The structure of the convert market also shifted: corporate issuers, recognizing that hedge funds were natural convert buyers, began structuring convertibles specifically for the hedge fund market. New issuance grew from \$30bn in 1995 to \$200bn in 2007. Convert sponsors (the underwriting banks) earned substantial issuance fees and built relationships with the convert-arb buyer base.

2008 — the catastrophe. The Lehman collapse and its aftermath broke convert arb. Three things happened simultaneously:

1. *Funding crisis*: hedge funds running 4-8x leverage faced margin calls. Their prime brokers (Lehman, Bear, MS, GS) cut leverage drastically. Funds had to delever.
2. *Convert market freeze*: convertible bond prices collapsed. Some converts traded at 50-70 cents on the dollar.
3. *Equity short squeeze*: the SEC's emergency short-sale ban on financials (September 19, 2008) made shorting impossible for some converts. Convert arb funds had short hedges that they couldn't roll, and their long converts were collapsing in price.

Convert arb fund losses in October-November 2008 averaged 40-50%. Some major funds (Blue Mountain, Camulos) wound down or restructured. Industry AUM dropped from \$80-100bn to \$20-30bn within months. The 2008 episode is famous as one of the “death of a strategy” moments; the hedge fund category was widely declared dead.

2009-2015 — the resurrection. As funding markets reopened and convertibles repriced, the surviving convert arb funds saw extraordinary opportunity. Converts that had traded at 50-60 cents in late 2008 could be bought, hedged, and held to par. The 2009-2011 period was, paradoxically, the most profitable in convert arb's history — for the funds that survived 2008 with capital.

2016-2020 — the transformation. Convert issuance shifted dramatically from financial-engineering converts (designed for hedge fund buyers) to *strategic converts* — issued by tech and growth companies as a tax-efficient way to fund operations. Tesla, Carvana, and many SPAC-related issuers used convertibles to raise capital. The embedded call's value was high (because the issuers were high-volatility), and the convert arb book on these names was extraordinarily profitable.

Industry AUM recovered to \$40-50bn. The hedge fund category was back, with a different player base (specialist funds rather than the large multistrats of pre-2008) and different issuers.

2021-2024 — the MicroStrategy era. MicroStrategy (MSTR) became the single most important convert issuer in the convert arb world. Saylor's strategy of issuing zero-coupon convertibles at very high implied vol (because BTC's vol is 80%+) generated multi-billion-dollar issuance pipelines. MSTR alone accounted for 5-10% of US convert issuance in 2023-2024. Convert arb funds running long-MSTR-convert / short-MSTR-stock positions reported the highest single-name returns in the strategy in many years.

The MSTR-driven convert market also changed the borrow market. MSTR borrow rates went from 1-2% pre-2021 to 5-15% in 2023-2024. The convert arb fund's effective return on the trade depended entirely on the borrow rate — a 5% borrow on a 10% expected gain meant the trade was profitable; a 15% borrow meant it wasn't. The borrow market's sensitivity to convert issuance became a critical input for convert arb performance.

2025-2026 — current state. Convert arb is a mature strategy with \$30-40bn industry AUM, a Sharpe ratio of around 1.0-1.5, and modest excess returns over the simpler benchmark of just owning convertible bond ETFs (which provide passive convert exposure). The trade is not retail-executable in any meaningful sense — minimum convertible trade sizes are \$1m+, the underlying short hedging requires institutional borrow access, and the daily delta-hedging is operationally intensive. But the academic and practitioner literature on convert arb continues to inform retail-relevant questions: how convertibles imply borrow rates, how convertible cheapness signals corporate credit stress, and how convert issuance affects equity volatility.

7.10 7.10 References

- Loncarski, I., ter Horst, J., and Veld, C. (2009), “The Rise and Demise of the Convertible Arbitrage Strategy”, *Financial Analysts Journal* 65(5): 35-50.
- Choi, D., Getmansky, M., and Tookes, H. (2009), “Convertible Bond Arbitrage, Liquidity Externalities, and Stock Prices”, *Journal of Financial Economics* 91(2): 227-251.
- Mitchell, M., Pedersen, L.H., and Pulvino, T. (2007), “Slow Moving Capital”, *American Economic Review* 97(2): 215-220.
- Hutchinson, M. and Gallagher, L. (2010), “Convertible Bond Arbitrage: Risk and Return”, *Journal of Business Finance and Accounting*.
- Agarwal, V., Fung, W.H., Loon, Y.C., and Naik, N.Y. (2011), “Risk and Return in Convertible Arbitrage”, *Journal of Empirical Finance* 18(2): 175-194.
- Calamos, J. (2013), *Convertible Arbitrage*, Wiley Finance Series.

8 Chapter 8 — Index-Rebalance and Corporate-Action Microstructure

8.1 8.1 The Index Rebalance Trade

S&P 500, Russell 2000, MSCI EM, Nasdaq 100 — every major index rebalances on a known schedule. Adds and deletes are typically pre-announced. Index funds (which collectively manage trillions in passive AUM) must trade at the rebalance close. This creates a price impact that can be anticipated.

The classical “index buy” trade:

1. Index committee announces XYZ will be added to S&P 500, effective Friday’s close.
2. By Friday’s close, every S&P 500 index fund must own XYZ in proportion to its weight.
3. Demand is large, predictable, and price-insensitive.
4. Specialists, prop traders, and delta-one desks accumulate XYZ before Friday and unload to the index funds at the close.

The expected price impact for an S&P 500 add has historically been 7-10% over the announcement-to-effective window. Post-2010 the impact has compressed to roughly 3-5% as more participants front-run.

A delta-one desk’s role: - Own client flow of S&P 500 funds (the buyers). - Source XYZ from agency lending or proprietary inventory. - Capture both the rebalance bid/ask and the lending fee on XYZ during the holding window.

8.2 8.2 The Russell Reconstitution

Russell rebalances annually (June “Recon”). The list of changes is enormous — typically 200-300 names. The aggregate dollar volume on Russell Recon Friday is one of the largest single-day flows in the market.

Russell typically pre-announces the additions and deletions in May. Between announcement and execution, prop desks and delta-one desks build positions. The price impact is substantial:

Russell Recon historical impact by tier:

Russell 1000 → Russell 2000 (downgrades): typical +3-5% gain
Russell 2000 → Russell 1000 (upgrades): typical -2-3% loss
Pure adds to Russell 3000 (small biotechs, etc.): +5-15%
Pure deletes: -3-7%

The trade is specifically dangerous because reconstitution dates are *the* trading days when the Russell index funds turn over their basket. Flash-crash-like dislocation is possible if liquidity is thin. The June 2020 Russell Recon was particularly volatile due to COVID-era thin liquidity in small-caps.

8.3 8.3 Spin-Offs and When-Issued Markets

When a parent company spins off a subsidiary, both ParentCo and SpinCo trade as separate stocks post-spin. Pre-spin, a “when-issued” (WI) market trades both. Sophisticated traders use the WI market to express views on the post-spin valuation:

- WI trades reflect the market’s forecast of post-spin pricing.
- Price discovery in WI is often inefficient (low volume, dealer-controlled).
- Delta-one desks make markets in WI as a service to clients and capture the bid/ask.

Concrete example: General Electric’s spin of GE HealthCare (GEHC) in January 2023. The WI market for GEHC traded for roughly two weeks before regular-way settlement. Dealer markets were 50-100 bps wide; sophisticated traders bought from dealers and resold to mutual funds at tight spreads, capturing a 30-50 bps margin.

8.4 8.4 M&A Risk Arb From the Desk’s Perspective

Risk arbitrage (merger arb) is its own business, but delta-one desks see it from the *plumbing* angle:

- A cash deal: target stock rises to near-deal-price after announcement, with small “deal spread” reflecting completion risk and time-to-close.
- A stock deal: target stock tracks a fixed exchange ratio of acquirer stock; arb hedge is short acquirer, long target.
- A mixed deal: complex hedge, often requiring options.

The desk's role: locating, lending, and financing the *acquirer's* stock for risk arb funds shorting the acquirer. Acquirer names in deals frequently go special — short interest in pending deals can drive 200-500 bps borrow fees. A delta-one desk sourcing acquirer-stock locates for a risk arb fund books significant fee revenue during the deal pendency.

8.5 8.5 Tender-Offer Arb

A tender offer: a bidder offers to buy shares at a fixed price, conditional on enough shares being tendered. Holders decide whether to tender or hold for the post-deal market.

Sophisticated arb: - Buy stock pre-tender at \$X. - Tender into the offer at \$Y > X. - If the deal closes, profit Y - X. - If the deal fails, hold the stock (now at potentially lower price).

The desk's edge: pricing the probability of deal close. With superior information (research, legal analysis), the desk often makes markets at narrower spreads than retail can access.

8.6 8.6 Corporate Action Worked Example: A Spin-Off

Parent (PRNT): trades at \$80, market cap \$40bn. Announces spin of subsidiary (CHLD) representing 25% of value.

Pre-spin trading: - Post-spin PRNT is forecast at \$60 (75% of pre-spin). - Post-spin CHLD is forecast at \$20 (25% of pre-spin). - Pre-spin investors will receive 1 share of CHLD per share of PRNT.

WI market two weeks pre-spin: - WI CHLD trading at \$19 bid, \$21 ask. - WI PRNT trading at \$61 bid, \$63 ask.

Delta-one desk's market-making P&L: - Buys 1m CHLD WI at \$19. Sells 1m CHLD WI at \$21. Captures $\$2m \times 0.5$ (half realized as midprice) = \$1m gross. - Same on PRNT: $\$2m \times 0.5 = \$1m$. - Plus dealer carries: cost of inventory through the WI period (~\$50k). - Net: \$1.95m on the spin's WI market over two weeks.

This is per spin. A Tier 1 desk sees 30-50 spins per year — multi-tens-of-millions in revenue.

8.7 8.7 The S&P 500 Add Trade — Detailed Mechanics

The S&P Dow Jones Indices announces additions/deletions to the S&P 500 typically late on a Friday, with an effective date the following Friday close. The pre-announcement-to-effective window is 6 trading days.

A typical add: a \$20bn market-cap company replacing one that has been acquired. The aggregate buying demand from index funds:

Total S&P 500 fund AUM tracking the index: ~\$8 trillion (Vanguard, BlackRock, fidelity, others).
The new add will represent ~0.4% of S&P 500 market cap (proportionate weight for \$20bn add).
Aggregate demand: $0.4\% \times \$8 \text{ trillion} = \32 billion in net buying.

The stock's pre-announcement market cap: \$20bn.

The stock's average daily volume: ~\$200m.

\$32bn of demand into a \$200m/day stock: 160 days of normal volume.

Most of this demand concentrates in the 6-day pre-effective window:

$\$32\text{bn} / 6 \text{ days} = \5.3bn/day average - 26x normal volume.

This produces the classic 5-10% pre-effective price spike. Index funds *must* buy at the close on Friday; specialists, prop desks, and delta-one desks accumulate during the week to sell at Friday's close.

A specific historical example: the GameStop S&P 500 add (May 2024).

GameStop wasn't actually added — but the meme-stock anticipation in 2024 around possible additions triggered scrutiny. The 2024 add of CrowdStrike (CRWD) is a cleaner example: - Announced June 21, 2024 close. - Effective June 24 close. - CRWD price: \$385 on announcement day. - Pre-effective spike: ~5% over the weekend and Monday. - Effective day close: \$403.

Index funds bought \$4-6bn of CRWD at the close. Specialists who had accumulated 100,000+ shares between announcement and effective day made \$2-3m on the basis-point spread captured.

8.8 The “Index Fund Front-Running” Question

Critics of the modern index ecosystem note that the *predictability* of index trades creates exploitable front-running opportunities. Vanguard, BlackRock, State Street, and the smaller passive managers must trade at the rebalance regardless of price. Predators (specialists, prop desks) accumulate ahead of them.

The estimated cost to passive investors: 7-10 bps per year on average index funds, or roughly \$5-8bn/year industry-wide. This is the “passive tax” paid to active traders.

Mitigation strategies adopted by major passive managers:

1. **VWAP execution.** Spread the rebalance over the day to avoid concentration in the closing minute.
2. **Pre-rebalance execution.** Begin acquiring within minutes of the announcement to reduce the predictable last-day demand.
3. **Crossing networks.** Use dark pools and crossing platforms to reduce visible footprint.
4. **Sampled tracking.** Don’t replicate the index exactly — hold a sampled portfolio that doesn’t require the most-volatile rebalance trades.

Vanguard reportedly captures ~5 bps of “execution alpha” annually relative to the index by using these techniques. BlackRock’s iShares does similar. The retail buyer of these ETFs benefits.

8.9 Spin-Offs — The When-Issued Market in Detail

When-issued markets are technically a *gray market* — trades are conditional on the corporate action actually occurring. The exchange uses a separate ticker for the WI security (e.g., GEHC.WI for the GE HealthCare spin pre-effective).

The pricing dynamics in WI are unusual:

- **Information asymmetry:** Sophisticated funds run detailed models of post-spin valuation; retail typically does not.
- **Liquidity:** Limited; bid/ask 1-3% wide common.
- **Volatility:** Often higher than the post-spin would imply — uncertainty about deal mechanics, allocation, etc.
- **Exit:** WI trades become “regular way” at the effective date; positions automatically convert.

Worked example: the GE HealthCare spin-off (December 2022 → January 2023).

Day -30 (mid-December 2022):

GE announced 1:1 distribution of GEHC stock to GE shareholders.

GE stock: \$76.

WI markets opened for both GE-pre-spin and GEHC-pre-spin.

Day -20:

GE-WI: \$63.50 ask

GEHC-WI: \$58.50 bid

Implied total: \$122 - meaningful premium to current GE.

Sophisticated traders who saw the post-spin valuation as $\$63 + \$55 = \$118$

could short the WI premium and capture the convergence.

Day 0 (January 4, 2023, effective date):

GE opens at \$62.30 regular-way.

GEHC opens at \$56.10 regular-way.

Combined: \$118.40.

WI premium: ~\$3.60 per share, captured over 20 days.

On a 100,000 share trade: \$360,000 profit.

The delta-one desk’s role: making both sides of the WI market, capturing the 100-300 bp dealer spread, and managing the 10-20 days of inventory exposure. A typical major spin generates \$500k-\$5m in market-making profit per delta-one shop.

8.10 The Tesla S&P 500 Inclusion (December 2020)

The single largest index-inclusion event in the history of US equities was Tesla's addition to the S&P 500 in December 2020. The event is worth its own subsection because it illustrates every dimension of the index rebalance trade — front-running, ETF AP arbitrage, sec lending, dealer market-making, and the limits of the “passive tax” calculation.

The setup. On November 16, 2020, S&P Dow Jones Indices announced that Tesla would be added to the S&P 500 effective December 21, 2020. At the announcement, Tesla had a market cap of approximately \$400 billion, making it the *largest* company ever added to the S&P 500 in a single event. The expected demand from index funds: approximately \$80-100 billion of Tesla buying compressed into a handful of trading sessions.

The pre-announcement positioning. Sophisticated index-rebalance funds and delta-one desks had been speculating about Tesla's inclusion for months before the announcement. The S&P committee's public criteria (positive earnings on a trailing 12-month basis) had been satisfied for the first time in mid-2020. By October, the prediction markets and S&P Index Committee insiders were assigning roughly 70% probability to inclusion in the December review. Hedge fund prop desks accumulated long Tesla positions in anticipation.

The post-announcement spike. On Tuesday November 17 (the trading day after the Friday announcement), Tesla opened up 8% to \$441. Over the following five sessions, Tesla rallied another 12% to \$497 by November 24 — a 20% gain in less than two weeks attributable almost entirely to the index-inclusion buying anticipation.

The rebalance window. The S&P committee delayed the actual inclusion mechanic by allowing a “phased addition” — instead of doing the full \$80bn purchase at the close of December 18 (the rebalance day), the index funds were given guidance to phase in their purchases over the preceding week. This was a procedural innovation specifically because the committee feared a single-day shock.

The phased approach reduced but did not eliminate the price impact. Over December 14-21, Tesla rallied another 10% on top of the prior gains. Total inclusion-related price impact: approximately 25-30% over the announcement-to-effective window.

The delta-one desk's positions. Major delta-one desks had positioned long Tesla ahead of the announcement (because of the high inclusion probability) and continued accumulating through the announcement window. Goldman Sachs alone reportedly held 2-3% of Tesla's float at peak (roughly \$10-12bn long position) — acting as a natural seller into the index funds' demand at the rebalance closes.

The desk's P&L on this trade was extraordinary. Based on public 10-K disclosures and reasonable assumptions about the inclusion-impact share captured by the major desks, Goldman, Morgan Stanley, JPMorgan, and Citadel collectively earned roughly \$1-1.5 billion in index-inclusion alpha on the Tesla addition. The largest single-name index-inclusion P&L in the history of these desks.

The ETF AP arbitrage. Tesla's inclusion affected SPY's basket composition. Beginning December 21, SPY's PCF included Tesla at its index weight (initially around 1.7% of the index). APs adjusted creation/redemption baskets to reflect the new composition. The transition was operationally smooth; the ETF AP teams had been preparing the new baskets for weeks.

Net effect: SPY traded at NAV throughout the rebalance, despite the underlying Tesla volatility. The AP arbitrage worked exactly as designed.

The sec-lending angle. Tesla's sec-lending demand spiked enormously around the inclusion event. Pre-announcement borrow rate: ~3% per annum. By mid-December, it had risen to 8-10% per annum, driven by hedge funds shorting Tesla on the thesis that the inclusion-induced spike would reverse. The lending desks at the Tier 1 PBs earned record fees on Tesla in December 2020 — single-name fees of \$50-100m per major desk for the month.

Post-inclusion, by mid-January 2021, Tesla's borrow rate normalized to 2% per annum as the index-inclusion-related short positioning unwound.

The retail dimension. Tesla retail interest was unprecedented during the inclusion window. Robinhood, Webull, and other retail platforms saw record Tesla trading volume. Retail call options on Tesla were heavily bought (notably by the “Reddit / WSB” community). This retail flow further amplified the inclusion-related demand and created a temporary disconnect between Tesla's option-implied volatility (very high) and its underlying realized volatility.

The retail dimension matters for the analysis because it created a *secondary* arbitrage: the Tesla option market in December 2020 priced extreme volatility that did not actually materialize. Sophisticated traders (delta-one desks, vol arb funds) sold Tesla volatility into the retail demand and captured the realized vs implied gap.

Lessons for the index-rebalance trade. The Tesla event re-validated the basic mechanic — index-inclusion demand creates predictable price impact — at a scale that had never been seen. But it also illustrated three nuances:

1. *Anticipation is now standard.* The market was positioned long Tesla for months before the announcement. The inclusion-day spike was therefore much smaller than it would have been historically because most of the alpha had been front-run.
2. *Phased inclusion reduces dislocations.* The S&P committee’s phased approach prevented a single-day catastrophe. Future large inclusions will likely follow this pattern.
3. *Retail flow can amplify.* Retail option buying on the inclusion name can create vol-pricing dislocations that traders can capture. This is now a standard part of the inclusion-trade playbook.

8.11 The AT&T / Discovery Spin-Off (April 2022)

A different kind of corporate action: a major spin-off where the parent (AT&T) divested its WarnerMedia subsidiary into a new combined entity (Warner Bros. Discovery). The transaction was announced in May 2021 and closed April 8, 2022.

The structure. AT&T shareholders received 0.241917 shares of WBD for each AT&T share held as of the record date (April 5, 2022). Following the spin, AT&T retained its core telecom business; WBD became a separate publicly traded media company.

The when-issued market. Pre-spin, WBD-WI traded on the Nasdaq for approximately two weeks. Mid-WI prices ranged \$24-26 per share, implying a total post-spin valuation (AT&T + WBD) of roughly \$25-27 per pre-spin AT&T share. AT&T pre-spin was trading at \$22-24, suggesting the WI market was pricing in a meaningful spin-off premium.

The desk’s market-making. Delta-one desks at JPM, Goldman, MS provided market-making in both AT&T-WI and WBD-WI. The bid-ask was roughly 30-50 cents on each, with the desk capturing the spread on every trade. Aggregate spin-related market-making revenue across the major banks: estimated \$50-100m for the two-week WI period.

The post-spin price discovery. April 8, 2022 (effective date) saw AT&T open at \$19.62 and WBD open at \$24.78. The combined value: approximately \$25.62 per pre-spin AT&T share. This was slightly below the WI market’s expected price of \$25-26, reflecting the typical “spin-off discount” that emerges when actual price discovery occurs.

The institutional positioning. Sophisticated funds had been positioning long AT&T pre-spin and short WBD-WI in anticipation of the discount. The trade earned about 80 bps over the two-week WI period — modest in percentage terms but enormous in dollars given the size of the spin (the combined AT&T+WBD market cap was roughly \$200bn).

The retail dimension. AT&T was a heavily-held retail dividend stock pre-spin. Many retail holders were surprised by the spin mechanics, particularly the WBD shares they received that did not pay a dividend. Retail selling of WBD in the weeks following the spin contributed to a 30% decline in WBD over the next three months. This retail-driven downside was partially captured by short-WBD positions opened by sophisticated funds during the WI period.

8.12 An M&A Risk-Arb Walkthrough — Microsoft / Activision Blizzard (2022-2023)

The Microsoft acquisition of Activision Blizzard (announced January 2022, closed October 2023) is one of the largest M&A risk-arb opportunities of recent years. The deal had multiple risk-arb cycles as regulatory uncertainty fluctuated.

The deal terms. Microsoft offered \$95 per share in cash for ATVI. Pre-announcement ATVI was trading at \$65; the offer represented a 46% premium.

The risk-arb spread evolution.

Date	Event	ATVI price	Implied prob of close
2022-Jan-18	Deal announced	\$82.31	63%
2022-Mar-15	Initial regulatory review begins	\$80.45	60%
2022-Jul-13	FTC announces review	\$77.45	54%
2022-Dec-08	FTC sues to block	\$75.20	48%

2023-Apr-26	UK CMA blocks deal	\$77.65	52%
2023-Jul-11	FTC drops court challenge	\$90.40	91%
2023-Jul-21	Deal modified for UK approval	\$93.05	96%
2023-Oct-13	Deal closes	\$94.97	100%

The risk-arb trade. A risk-arb fund that bought ATVI at \$82 in January 2022, expecting eventual close at \$95, faced 21 months of uncertainty. The spread was closely tied to regulatory news; specific events (FTC suit, UK CMA block, FTC drop) caused 5-15% intraday moves in ATVI.

The financing structure. Risk-arb funds typically lever 2-3x via prime broker financing. A fund running \$100m of capital could hold \$200-300m of ATVI long. The total return on the deal: $\$95 - \$82 = \$13$ per share = ~16% gross. With 2x leverage and a 21-month hold, the IRR is approximately 18-20% per annum, before borrow and financing costs.

The desk's role. Major prime brokers financed the risk-arb fund positions. They also made markets in the deal-related options (ATVI puts and calls), which were heavily traded by funds expressing views on the regulatory outcome. The PB earned both financing income and option market-making income across the deal life.

The lesson. Risk-arb on a large deal with regulatory uncertainty is a multi-cycle, news-driven trade. The eventual close was favorable for risk-arbs, but the path was not smooth. A retail trader trying to replicate this would face several frictions: option market-making is institutional (the bid-ask in deal options is institutional-grade and not available to retail), the borrow on ATVI during the deal pendency was elevated (10-25% per annum at peak), and the regulatory news flow required real-time interpretation that retail typically cannot do well.

8.13 8.13 References

- Shleifer, A. (1986), “Do Demand Curves for Stocks Slope Down?”, *Journal of Finance* 41(3): 579-590. Original index-effect paper.
- Cai, J., Houge, T., and Marozzi, M. (2011), “Index Rebalancing and the Russell Effect”, working paper.
- Greenwood, R. (2005), “Short- and Long-Term Demand Curves for Stocks: Theory and Evidence on the Dynamics of Arbitrage”, *Journal of Financial Economics* 75(3): 607-649.
- Patel, S. and Weil, J. (2018), “Front-Running and Index Rebalancing”, *Wall Street Journal*.
- Bessembinder, H. (2018), “Do Stocks Outperform Treasury Bills?”, *Journal of Financial Economics* 129(3): 440-457.
- Chakrabarty, B. and Moulton, P.C. (2012), “Earnings Announcements and Attention Constraints: The Role of Market Design”, *Journal of Accounting and Economics* 53(3): 612-634.
- Mitchell, M. and Pulvino, T. (2001), “Characteristics of Risk and Return in Risk Arbitrage”, *Journal of Finance* 56(6): 2135-2175.

9 Chapter 9 — ETF Creation/Redemption Arbitrage

9.1 9.1 The Creation/Redemption Mechanism

ETFs are unique among investment products: shares are created and destroyed in a structured arbitrage process. Authorized Participants (APs), typically a delta-one desk at a major bank, can:

- **Create** ETF shares: deliver the underlying basket of stocks to the ETF issuer in exchange for newly-created ETF shares.
- **Redeem** ETF shares: deliver ETF shares to the issuer in exchange for the underlying basket.

The basket is specified in advance by the issuer (the “PCF” — Portfolio Composition File). For SPY tracking S&P 500, the PCF specifies 500 names with exact share counts.

The arbitrage: - ETF trades at \$0.05 above NAV. AP creates: deliver basket worth NAV, receive ETF shares worth NAV+0.05. Sell ETF shares. Profit \$0.05. - ETF trades at \$0.05 below NAV. AP redeems: deliver ETF shares worth NAV-0.05, receive basket worth NAV. Sell basket. Profit \$0.05.

This keeps ETF prices tight to NAV. The AP's profit is roughly 0.5-2 bps per trade on liquid ETFs (very tight markets) but can be 10-50 bps during stress (March 2020, August 2024 unwind).

9.2 9.2 Cash vs Basket Creates

Two flavors of creation:

Basket create. AP delivers the actual basket of stocks. Most liquid US equity ETFs (SPY, IWM, QQQ) operate this way.

Cash create. AP delivers cash; the issuer’s portfolio manager buys the basket. Used for international ETFs (where the AP can’t easily source the foreign basket), commodity ETFs (where the AP can’t deliver physical commodities), or new/illiquid ETFs.

Cash creates have a wider AP fee (the issuer is bearing transaction friction). On a typical cash create for an EM ETF, the AP may pay 25-50 bps to the issuer to compensate for the issuer’s basket-buying friction.

9.3 9.3 Worked Example: VOO Creation

Vanguard’s S&P 500 ETF (VOO). Suppose VOO trades at \$530.50 with NAV at \$530.45 — a 5-cent (1 bp) premium.

AP’s create: - AP buys the 500-name basket equivalent to one creation unit (typically 50,000 ETF shares for VOO). - Cost: $50,000 \times \$530.45 \text{ (NAV)} = \$26,522,500$. - Delivery: AP gives \$26,522,500 of basket to Vanguard, receives 50,000 VOO shares. - AP sells 50,000 VOO shares at $\$530.50 = \$26,525,000$. - Gross P&L: $\$26,525,000 - \$26,522,500 = \$2,500$. - After hedging frictions and the AP’s basket sourcing cost (typically 0.3-0.5 bps), net P&L: \$1,500.

This is for a single creation unit at a 1 bp premium. At a busy AP, dozens of creation units may transact daily on VOO alone. The aggregate AP P&L on VOO is substantial.

9.4 9.4 Stress Period Dislocations: March 2020

In March 2020, fixed-income ETFs (LQD, HYG, JNK) traded at unprecedented discounts to NAV — 5-8% in some cases. The mechanism:

- Bond market liquidity collapsed.
- Bond mutual funds and ETF NAVs were calculated using stale/dealer-quoted bond prices.
- ETFs traded at the *real-time market clearing price* — which was below the official NAV.
- APs couldn’t easily redeem because the underlying bonds were illiquid; redemption would force selling into a thin market at large losses.

The “discount” was not arbitrageable in the standard sense because the redemption process was operationally broken. The ETF was actually correctly priced; the NAV was the lagging measure.

This is a critical lesson: ETF arbitrage works only when the AP can actually transact in the underlying basket. When the basket is illiquid, the ETF can dislocate from NAV substantially and persistently.

9.5 9.5 The “Cash Drag” and ETF Tracking Error

ETFs hold small cash positions (typically 0.5-2% of NAV) for operational reasons. This creates a small but persistent tracking error vs the index:

- Bull market: ETF lags index by approximately cash-weighted equity return.
- Bear market: ETF outperforms index by approximately cash-weighted equity return (cash held flat vs declining stocks).

Delta-one desks running ETF arbitrage are aware of this and price it into create/redeem decisions. Sophisticated arb may include a small cash leg to optimize.

9.6 9.6 The “Heartbeat” Trade — Tax Efficiency in ETFs

A subtle but important ETF mechanism: the “heartbeat” trade. An AP and an ETF issuer coordinate to use the in-kind redemption process to *purge* low-cost-basis securities from the ETF’s portfolio without realizing gains.

The mechanism:

1. AP creates large new ETF shares with cash (or low-basis stock).
2. Days later, AP redeems an equal quantity of ETF shares.
3. Issuer satisfies redemption with the highest-unrealized-gain stocks from its portfolio.
4. AP receives stocks at the issuer's cost basis (not market value) for tax purposes - but at fair market value economically.
5. AP can sell these in the open market, pay taxes only on small incremental moves since receipt.
6. The ETF has now "purged" low-basis stocks without paying capital gains tax. ETF NAV unchanged. Existing holders saved tax exposure.

The trade is most heavily used by Vanguard mutual fund / ETF “share class” structure (where mutual funds and ETFs share a portfolio) and by major equity ETF issuers managing decade-old stock positions.

Estimated industry tax savings from heartbeat trades: \$10-15bn annually in deferred capital gains. The IRS has periodically threatened to limit the practice, most recently with proposed Treasury Regulations in 2019 that were softened after industry pushback.

For a delta-one desk acting as AP, heartbeat trades are a steady fee business. The AP earns a small fee from the issuer (typically 1-5 bps of the cycled volume) for executing the structure with no market impact.

9.7 9.7 ETF Share Class Structures and Vanguard’s Patent

Vanguard pioneered the “share class” ETF structure: an ETF and a mutual fund are different classes of *the same* underlying portfolio. This allows the ETF to absorb tax efficiency from the mutual fund’s holdings.

Vanguard’s patent on this structure expired in 2023, opening the floodgates for other issuers (DFA, Fidelity, T. Rowe Price) to convert mutual funds into ETF share classes. The total tax-efficiency benefit to investors of these conversions is estimated at \$1-3bn annually.

For delta-one desks, the new ETF share class structures generate additional AP volume — every conversion creates new AP relationships and trading flow.

9.8 9.8 The “Hidden Inventory” of Big-Basket ETFs

A Tier 1 PB’s ETF AP team doesn’t just create/redeem; they hold proprietary inventory of major ETFs. SPY, QQQ, IWM are continuous inventory holdings — the desk is *long* tens of millions of shares of these names overnight, both as inventory for client market-making and as a liquidity reservoir.

This inventory: - Earns securities lending fees (SPY borrows ~10-25 bps). - Earns dividend income. - Generates market-making spread on the bid/ask. - Costs balance sheet (RWA at ~75-100%).

The blended return on a Tier 1 desk’s \$5bn ETF inventory is typically 80-150 bps per annum — a steady, scalable business.

9.9 9.9 The 4:00pm \$50m ETF Redemption — A Detailed Walkthrough

To make the AP arbitrage flow concrete, here is what actually happens when a large ETF redemption arrives at 4:00pm ET. The setting: a Tuesday in October 2024. A multi-strategy hedge fund has decided to liquidate \$50m of its position in iShares MSCI Emerging Markets ETF (EEM). They place the redemption order with their prime broker for end-of-day execution.

3:45pm ET. The hedge fund’s PB sends the redemption notice to BlackRock (EEM’s issuer). BlackRock’s AP team confirms the redemption will be processed at the 4:00pm cash close.

3:50pm ET. The AP team — typically a delta-one desk at one of the major banks (let’s say Morgan Stanley in this scenario) — receives notification of the incoming redemption from BlackRock. The AP team’s role: deliver the cash equivalent to the hedge fund and receive the underlying basket from BlackRock. The basket consists of approximately 800 international stocks held in trust by BlackRock.

3:55pm ET. The AP’s traders begin the arbitrage. Step 1: short-sell the underlying basket through the global execution desks. Step 2: prepare to deliver cash to the hedge fund (the redeem proceeds). Step 3: receive the basket

from BlackRock at 4:00pm and use it to cover the short basket.

The catch: the underlying basket spans 25 countries, multiple time zones, and several trading hours. Many of the EM markets close hours before 4:00pm New York time. The hedges must be approximated using ADRs, futures, or proxies for the closed markets.

3:58pm ET. The MS Equity Trading desk shorts approximately \$20m of the basket through ADRs (the most liquid component) and another \$15m through SPDR EM ETF futures (a hedge for the harder-to-source names). The remaining \$15m is hedged with cash equity sales of the most liquid EM names that are still trading (a few large Mexican and Brazilian names).

4:00pm ET. The redemption settles. BlackRock delivers the basket of 800 EM stocks (worth \$50m) to MS. MS receives the basket. MS delivers \$50m in cash to the hedge fund.

4:01pm ET. The AP's hedge book now consists of: - Long: \$50m of EM basket received from BlackRock. - Short: \$20m ADRs + \$15m SPDR futures + \$15m EM stocks. - Net delta: approximately \$0 (theoretically).

In practice, the basket received from BlackRock and the hedges shorted differ slightly. The basket is the *exact* MSCI EM index composition. The hedges are approximations. The basis between exact and approximate is the AP's residual risk.

4:30pm ET. The AP's basket is processed. The 800 stocks are matched against the short hedges. Where the AP is short an ADR but received the underlying foreign stock, the AP rebalances over the next session. Where the basket and the hedges match exactly, the position is flat. Where they differ, the AP carries a small residual exposure.

Next day — settlement. The basket settlement typically occurs T+2 in the US (some EM markets are T+1). The AP's stock loan and basket sourcing teams reconcile the trades, deliver the borrowed stock, and finalize the operations.

The AP's economics. For this \$50m redemption: - AP fee charged to BlackRock: approximately 5-10 bps of basket value = \$25-50k. - Bid-ask drag on the hedges: approximately 3-5 bps = \$15-25k. - Net AP profit: approximately \$10-25k on a \$50m flow.

This is per-trade. A typical AP team handles dozens of large flows per day, with cumulative daily volume in the billions. The aggregate AP business at a Tier 1 ETF AP team is \$50-150m per year.

The cash drag. The hedge fund that initiated the redemption receives \$50m in cash at the close. This cash sits in their broker account until they redeploy. Typical cash deployment time at a multi-strategy fund: T+1 to T+5. During this period the cash earns the broker's cash-sweep yield (currently ~5%), which is a small drag relative to the equity returns.

9.10 9.10 References

- Madhavan, A. (2014), "Exchange-Traded Funds and the New Dynamics of Investing", Oxford University Press.
- Pan, K. and Zeng, Y. (2019), "ETF Arbitrage Under Liquidity Mismatch", working paper.
- Petajisto, A. (2017), "Inefficiencies in the Pricing of Exchange-Traded Funds", *Financial Analysts Journal* 73(1): 24-54.
- Sushko, V. and Turner, G. (2018), "The Implications of Passive Investing for Securities Markets", *BIS Quarterly Review*.
- Moussawi, R., Shen, K., and Velthuis, R. (2022), "ETF Heartbeat Trades and Their Tax Implications", working paper.
- Ben-David, I., Franzoni, F., and Moussawi, R. (2018), "Do ETFs Increase Volatility?", *Journal of Finance* 73(6): 2471-2535.

10 Chapter 10 — Hard-to-Borrow and Squeezes From the Desk’s Perspective

10.1 10.1 The Other Side of the Squeeze

The retail-facing strategies in this repo include short-squeeze detection and momentum capture. From the *desk’s* side, hard-to-borrow names look entirely different:

- Lending revenue is enormous (recall the 45% of industry SLB revenue that comes from the top 2% of names).
- Recall risk is real but manageable through diversification.
- The *risk* is not the squeeze itself — the desk is long the inventory and benefits as the stock rises — but the *recall cascade* when the squeeze finishes and short demand collapses.

A typical pattern, watched closely by desks:

Phase 1 (Building):	Short interest builds, borrow fees rise from 1% to 10%
Phase 2 (Stable):	Borrow stable at 10-30%, lending revenue large
Phase 3 (Squeeze):	Stock rises, shorts cover, borrow demand spikes to 100%+
Phase 4 (Climax):	All shorts covered, borrow demand collapses
Phase 5 (Recall):	Beneficial owners realize the rate is unrealistic, stop lending Borrow fees normalize back to <1% over weeks

The desk’s revenue profile: maximal in Phase 3, tapering in Phases 4-5. Diversified across many names, the average annual rate is the steady-state “specials” rate — the squeeze episodes are upside, not the bread and butter.

10.2 10.2 GME 2021 Case Study From the Prime Broker Side

Q1 2021. GameStop short interest reached 140% of float. The borrow rate, which had been stable at 5-8% for months, spiked to 25% in mid-January and 80%+ during the squeeze.

Prime brokers’ positions during the event:

Lending side (the desk’s beneficial-owner clients): Earned spectacular fees. A pension fund holding \$50m of GME at 80% per annum was earning \$40m/year on that single position. Some agency lenders booked the highest single-month fees in their history.

Borrowing side (the desk’s hedge fund clients): Catastrophic. Funds short GME (Melvin Capital, Citron, others) faced borrow rates spiking from 5% to 80% while simultaneously losing on the directional move. Borrow cost on a \$1bn short went from \$50m/year to \$800m/year.

Recall and squeeze mechanics: - Some agency lenders recalled GME loans on January 27-28 to capture the high borrow rates by relending elsewhere or to liquidate gains. This forced shorts to cover. - Robinhood, Apex Clearing, and other retail brokers raised margin on GME and other meme names. This was widely misreported as “blocking buying” but was actually a settlement-system adjustment to higher volatility. - The DTCC raised collateral requirements on January 28; this is what triggered the brokers’ restrictions.

Aftermath for desks: - Several Tier 2 hedge funds went bust. - Prime brokers booked large losses on hedge fund counterparty defaults. - Sec lending desks booked record revenue for January 2021 alone.

Net for major prime brokers: lending revenue large positive, hedge fund counterparty losses large negative. JPM, Goldman, Morgan Stanley each booked \$50-200m in losses related to GME-era hedge fund failures. Sec lending revenue offset most of this.

10.3 10.3 The “Locate Theater”

A persistent practice in pre-2008 markets: brokers issuing “locate” assurances to short sellers without actually having located the stock. This enabled “naked shorting” — shorting without the obligation to deliver.

Post-2008 Reg SHO closed most loopholes. Modern brokers must have hard locates (actual borrowed stock) before short orders can be entered. The exception, “options market makers,” retain a narrow exemption for hedging activity, which has been controversial.

The 2021 House Financial Services Committee hearings on the GME event generated extensive testimony on this; no major regulatory changes resulted, but the locate process is now more strictly monitored.

10.4 10.4 The Buy-In Risk

Buy-in: the broker forcibly closes a short position when shares cannot be returned. Triggered by:

- Failed delivery of the borrow.
- Beneficial owner recall with no available re-borrow.
- Regulatory requirement (SHO close-out).

The buy-in is executed at the prevailing market — often the *worst* possible time, since recalls coincide with squeeze conditions. A short fund's buy-in during a squeeze is the most expensive execution they will ever have.

10.5 10.5 Synthetic Short via Options — The Workaround

When direct shorting is unavailable (HTB at >50%) or restricted, hedge funds use synthetic shorts via options:

Synthetic short = long put + short call, same strike, same expiry

This delivers an economic short equivalent to the underlying. The cost: the put-call parity adjusts via the rebate rate. If the borrow rate is 50%, the implied option-pricing adjustment is:

$$P - C = K \times e^{-rT} - S \times e^{-\text{borrow} \times T}$$

For $S = K$ (ATM), borrow = 50%, $T = 30$ days, $r = 5\%$:

$$P - C = K \times (e^{-0.0041} - e^{-0.0411}) \\ K \times 0.0367$$

So the put trades 3.67% above the call at the same strike. Synthetic shorts inherit the borrow cost via option pricing. Trying to dodge the borrow market by going to options is *not* a loophole — the borrow is priced in.

10.6 10.6 The “Locate Pool” and Why It’s Not Infinite

A common retail misconception: “if I can buy 1,000 shares, surely 1,000 shares are available to short.” The reality: the *long* market is much larger than the *lending pool*. Many beneficial owners do not lend their stock:

- IRA holders typically cannot lend (regulatory restrictions).
- Some mutual funds elect not to lend (fiduciary discretion).
- Foreign holders may have tax/jurisdictional constraints.
- Active retail (which trades the stock daily) generally is not in lending programs.

A typical large-cap stock: 70-80% of float is held in lendable form (institutional, custody-bank-held). Of that, maybe 60-70% is actively in lending programs. So the actual lendable pool is 50-60% of float. Most of the time this is enormously sufficient (short interest is rarely above 5% of float). For HTB names, the math goes the other way:

GME, January 2021:

Float: 50m shares

Lendable pool (~50%): 25m shares

Short interest: 70m shares (140% of float)

Net: short interest exceeded the entire lendable pool.

Some shorts had to be re-lent multiple times - each "borrow" satisfied a short, then was re-lent to another short, then another. The locate accounting becomes recursive.

This recursive locate is technically legal but creates fragility — when even one beneficial owner recalls, the whole chain can collapse, forcing multiple covers in sequence.

10.7 10.7 The “Threshold Securities List”

The SEC requires brokers to publish daily a “threshold list” of securities with persistent failures-to-deliver. Stocks on this list are subject to mandatory close-out within 35 days.

Names regularly on the threshold list in 2024-2025: - BBBY (pre-bankruptcy) - AMC (during squeeze periods) - SAVE (during bankruptcy) - Various small-cap biotech with high short interest

For a delta-desk, threshold-list stocks are *operationally expensive* — extra compliance overhead, more frequent margin checks, additional reporting. This is one reason the “specials” rate is so high — the compliance load alone is 20-50 bps per annum on a small-cap special.

10.8 10.8 GME 2021 — A Full Prime Broker Narrative

The GME story is widely told from the retail side (the “Reddit / WSB” narrative) and from the broker-restriction side (Robinhood vs Citadel). The prime broker side is less told but more illuminating about how the institutional infrastructure actually responded. This subsection walks through the prime broker narrative day by day during the critical week.

Pre-event setup (December 2020 to mid-January 2021). GME short interest had been building for months, driven by hedge funds (Melvin Capital, Citron, others) shorting the company on the thesis that retail-store-based gaming would decline structurally. By early January 2021, GME’s short interest was 140% of float — i.e., more shares were sold short than existed in tradeable form. The borrow rate had risen from 3-5% per annum (in 2020) to around 25% by mid-January. A typical prime brokerage borrow on GME was earning 8% per annum to the desk’s beneficial-owner clients (after PB markup and agent split).

Wednesday January 27, 2021. GME has rallied from \$40 a week ago to \$370 intraday. Short interest is forced to cover. Borrow rate spikes to 80% per annum. Prime brokers are receiving recall notices from beneficial owners (some of whom want to capture the high rate by relending elsewhere; others are simply liquidating GME from their portfolios).

The PB’s situation: hedge fund clients (Melvin, Citron, others) are short GME at original borrow rates of 5-8%. They are now facing borrow at 80%+ AND adverse mark-to-market. The PB’s risk team is on the phone constantly with these clients about variation margin.

The PB’s revenue: enormous. A pension fund holding \$50m of GME at 80% per annum is earning \$40m/year on a single position (annualized). A typical PB has several such positions; the daily lending revenue from GME alone hit \$5-10m per day per major PB during the peak.

Thursday January 28, 2021. GME reaches \$483 intraday, then closes at \$193 after Robinhood’s restrictions kick in. The DTCC (Depository Trust & Clearing Corporation) notifies major brokers of increased collateral requirements on meme-stock positions. Robinhood, lacking sufficient capital to post the new collateral, restricts buying.

Behind the scenes at major prime brokers: emergency calls between risk teams and clients. Several hedge funds (Melvin in particular) report multi-billion-dollar losses; their PB relationships re-evaluate margin and credit terms. Citadel reportedly injects \$2.75bn into Melvin to stabilize the fund.

The PB’s position: most major PBs were running secured exposures (i.e., collateralized loans to short funds). The collateral was sufficient to cover Day 1 losses but not enough to absorb additional losses if positions continued to move. Several PBs raised initial margin requirements on GME and other meme names to 100% (i.e., requiring full cash for the short position).

Friday January 29, 2021. GME settles at \$325. Short interest still elevated but coming down rapidly. The borrow rate begins to normalize from peak 80%+.

The PB’s actions: continued margin tightening on remaining GME shorts. Several PBs forced liquidations of clients who couldn’t meet margin calls. The forced liquidations contributed to the further short-cover squeeze.

Monday February 1, 2021. GME is at \$225. Short interest has dropped from 140% to ~50% of float. The squeeze is unwinding.

Industry-wide PB losses for January 2021 from GME-related events: - Hedge fund client losses absorbed by PBs (defaults / dragdowns): estimated \$500m-\$1bn industry-wide. - Sec-lending revenue gain from GME alone: estimated

\$1.5-2.5bn industry-wide for the month. - Net: positive for the PB industry in aggregate, with some specific firms taking large hits.

The Tier 1 perspectives.

Goldman Sachs. Reportedly took \$50-100m in client default-related losses but earned \$300-500m in lending revenue, for a net positive contribution of \$200-400m. The CEO, David Solomon, mentioned the strong sec-lending business in the Q1 2021 earnings call.

Morgan Stanley. Similar profile. Morgan Stanley's prime brokerage business had been growing rapidly through 2019-2020 and had picked up market share from Deutsche Bank (which exited PB). The GME event was, for MS, a stress test that they passed with flying colors.

JPMorgan. JPM's prime brokerage was less aggressive on hedge fund borrowing during 2020 and therefore had less concentration risk during the GME squeeze. JPM earned strong lending revenue and had minimal default issues.

Credit Suisse. This is the firm that was about to face Archegos two months later. CS had a significant equity-swap business with Melvin and other meme-affected funds. The GME event was a warning sign of risk concentration that CS leadership did not fully act on. By March 2021, when Archegos blew up, CS was under-prepared.

The lessons. The GME event illustrated three things about prime brokerage:

1. *The lending business is a positive-feedback in stress.* Higher borrow rates → higher client P&L losses → forced covers → even higher borrow rates → even higher PB lending revenue. The PB benefits when the squeeze runs.
2. *Counterparty risk concentrates around specific names.* The PBs with the most concentrated meme-stock short hedge fund clients took the biggest hits. Diversification of client base matters even more than diversification of inventory.
3. *Systemic risk is centralized at the DTCC.* The DTCC's collateral demands cascaded through the broker chain to retail. The DTCC's risk management — not Citadel's, not Robinhood's — was the actual binding constraint.

The GME event also reshaped the regulatory landscape: SEC Rule 10c-1 (effective 2024) requires reporting of securities lending transactions on a delayed basis, increasing transparency around the kind of borrow-rate dynamics that played out in GME.

10.9 10.9 References

- Diamond, D. and Verrecchia, R. (1987), “Constraints on Short Selling and Asset Price Adjustment”, *Journal of Financial Economics* 18(2): 277-311.
- Boehmer, E., Jones, C.M., and Zhang, X. (2008), “Which Shorts Are Informed?”, *Journal of Finance* 63(2): 491-527.
- US House Financial Services Committee (2021), “Game Stopped: Who Wins and Loses When Short Sellers, Social Media, and Retail Investors Collide”, February-May 2021 hearings.
- Fotak, V., Raman, V., and Yadav, P.K. (2014), “Fails-to-Deliver, Short Selling, and Market Quality”, *Journal of Financial Economics* 114(3): 493-516.
- Allen, F., Haas, M., Nowak, E., and Tenglouf, A. (2021), “Market Efficiency and Limits to Arbitrage: Evidence from the Volkswagen Short Squeeze”, *Journal of Financial Economics* 142(1): 166-194.
- Eaton, G.W., Green, T.C., Roseman, B.S., and Wu, Y. (2022), “Retail Trader Sophistication and Stock Market Quality: Evidence from Brokerage Outages”, *Journal of Financial Economics* 146(2): 502-528.

11 Chapter 11 — Total-Return Swaps, Equity Swaps, and the Archegos Case

11.1 11.1 The Total-Return Swap Mechanism

A total-return swap (TRS) is a bilateral contract:

Party A pays:	Total return on a reference asset (price change + dividends)
Party B pays:	Floating rate (e.g., SOFR + spread) on the notional

Party B obtains synthetic long exposure without owning the asset. Party A obtains hedge income for offering this exposure.

For a delta-one desk operating as Party A: - Hedge by buying the underlying (or basket). - Earn the financing spread. - Earn the securities lending revenue on the underlying. - Earn the dividend wedge. - Bear the counterparty risk (Party B fails).

For Party B (the hedge fund): - Synthetic long exposure with high leverage (initial margin typically 10-25%). - No ownership of underlying — no voting rights, no Schedule 13D filing requirements (a major attraction). - Dividend equivalent payment (subject to 871(m)) for US underlyings.

11.2 11.2 Section 871(m) Detail

Section 871(m) of the Internal Revenue Code, with regulations finalized in 2017:

- Treats “dividend equivalent payments” on certain notional principal contracts as US-source dividends for withholding purposes.
- Applies to “Section 871(m) transactions” — defined broadly to include equity swaps with delta 0.8 to a US underlying.
- Withholding rate: 30% (reduced to 15% by treaty for most foreign holders).

The rule effectively eliminates the dividend tax wedge for synthetic equity exposure on US names. It does not eliminate dividend trades on non-US names (Section 871(m) is US-specific).

11.3 11.3 The Archegos Capital Case Study

Archegos Capital Management, the family office of Bill Hwang, was an aggressive long-equity trader using equity swap leverage with a small number of major prime brokers — Credit Suisse, Nomura, Morgan Stanley, Goldman, UBS, and Mitsubishi UFJ.

The strategy: - Concentrated long positions in a handful of US tech and Chinese ADR names (ViacomCBS, Discovery, Baidu, Tencent Music). - Leveraged via equity swaps at multiple prime brokers simultaneously — none of the PBs knew the full position size at the others. - Total reported notional at peak: roughly \$50-100bn long exposure. - Capital base: roughly \$10-15bn.

The blow-up:

- Late March 2021: ViacomCBS announced a large equity issuance, causing the stock to drop sharply.
- Archegos margin calls hit at multiple PBs simultaneously.
- The fund could not meet calls.
- PBs began liquidating the underlying long positions to recover capital.
- Coordination failed: each PB liquidated independently, accelerating the price collapse.

PB losses:

Credit Suisse:	\$5.5 billion (largest loss)
Nomura:	\$2.9 billion
Morgan Stanley:	\$1.0 billion (managed exit best)
UBS:	\$0.86 billion
Goldman Sachs:	\$0.10 billion (best execution)
Mitsubishi UFJ:	\$0.30 billion
Total industry:	~\$10+ billion

Credit Suisse’s loss was a major contributor to its later collapse and absorption into UBS in 2023.

11.4 11.4 What Banks Should Have Done

Post-mortem analysis identifies specific failures:

1. **Cross-margining.** PBs did not have a coordinated view of Archegos’s total position. Each PB extended margin assuming Archegos was only marginally levered with them. The aggregate position was 5-10x larger than any individual PB recognized.

2. **Concentration limits.** Archegos held >\$10bn of single names like ViacomCBS — making PBs the largest non-corporate holders of these names. Concentration limits were not enforced.
3. **Margin haircuts.** Initial margin on equity swaps for Archegos was reportedly 7-10%. For concentrated single-name positions, 20-25% would have been appropriate.
4. **Liquidation choreography.** When Archegos defaulted, PBs were forced to liquidate. Goldman Sachs and Morgan Stanley exited their positions on March 26 morning. Credit Suisse, Nomura, and others hesitated — and then sold into a collapsing market over the next several days, accumulating the largest losses.

The lesson: in a multi-PB swap-leveraged structure, the fund's collapse is a coordination problem. Whoever moves first wins (sells at the highest price); whoever moves last loses (sells at the lowest). The bank's relationship-management instinct (trying to “work with” the client) is fatal.

11.5 11.5 Post-Archegos Industry Changes

Several industry-wide changes followed:

- US Securities and Exchange Commission proposed Rule 10c-1 (effective 2024) requiring reporting of securities lending transactions on a delayed basis.
- Form PF reporting expanded for hedge funds and family offices.
- Major PBs implemented stricter cross-margining and concentration limits internally.
- Some PBs (UBS, Credit Suisse) significantly reduced their swap-financing books.

11.6 11.6 Detailed Anatomy of an Equity Swap — A Sample Trade

To make the abstract concrete, here is the full structure of a typical hedge fund equity swap with a major bank:

TRADE TICKET (SAMPLE)

=====

Trade Date: March 15, 2025

Effective Date: March 17, 2025

Termination: September 17, 2026 (18 months)

Reference Asset: 100% of MSFT Class A common stock

Notional: \$250,000,000 USD

Performance Leg (paid by Bank to Fund):

Total Return on MSFT, including:

- Stock price appreciation/depreciation
- Cash dividends (subject to 15% withholding per 871(m))
- Special dividends/spinoffs (full economic equivalent)

Reset: Daily mark-to-market

Funding Leg (paid by Fund to Bank):

SOFR + 65 bps spread on the daily MTM notional

Reset: Daily

Initial Margin Required:

\$37,500,000 (15% of notional)

Eligible collateral: cash, T-bills (haircut 0%),
US treasuries < 5y (haircut 2%),
MMF (haircut 0%)

Variation Margin:

Daily, based on previous-day close

MTA: \$1,000,000

Threshold: \$0 (no unsecured exposure permitted)

Termination Events:

- Failure to pay (3 business day cure)
- Bankruptcy
- Cross-default at fund's other PB relationships
- Material adverse change clause

Counterparty Termination Right:

Bank may terminate with 5 business days notice
 if MSFT becomes special (>5% borrow fee per annum)
 or if regulatory capital requirements change materially

The economics for the bank:

- Earn the SOFR + 65 bps spread on \$250m notional: \$1,625,000/year.
- Lend the underlying \$250m of MSFT stock: at GC ~12 bps, \$300,000/year.
- Dividend wedge (Section 871(m) treatment captures 15-30% of the dividend tax wedge): \$200-400k/year.
- Total bank revenue: ~\$2.0-2.3m/year on a ~\$2.5m balance sheet allocation.

The economics for the fund:

- Synthetic long \$250m MSFT.
- Cost: SOFR + 65 bps = 5.95% per annum. On \$250m = \$14.875m/year.
- Compared to outright purchase: \$250m × cost-of-equity-financing ~10% = \$25m/year. The swap saves \$10m/year vs cash purchase. Plus the fund avoids 13D filing requirements (no public disclosure of position). Plus dividend tax efficiency for the foreign client tier.

This is a typical structure. Variations include floored or capped financing, dividend-passthrough vs dividend-equivalent, and basket vs single-name references.

11.7 11.7 The Variance-Swap Cousin

Adjacent to delta-one is the variance swap business. A variance swap pays the difference between realized variance and a strike variance over the swap's life. While not strictly delta-one (the payoff is convex in vol), variance swaps share the same prime brokerage infrastructure.

For a delta-one trader, the relevant detail: variance swaps require *dynamic* delta hedging by the dealer. The dealer holds a static portfolio of options weighted by $1/K^2$ (the variance-swap replication formula). To delta-hedge the residual, the dealer continuously trades the underlying — which means continuous demand for both buying and shorting. Sec lending ties in: shorts require borrow.

A typical \$100m variance swap on SPY generates roughly \$200-500k in delta-hedge churn revenue per quarter for the dealer. Not large per swap, but cumulative across the variance-swap book.

11.8 11.8 Section 871(m) Compliance Detail

The compliance machinery for 871(m) is intricate. A bank holding equity swaps must:

1. **Identify** US-source dividend payments that flow through the swap.
2. **Determine** the deemed dividend equivalent — usually the gross dividend × delta of the swap.
3. **Withhold** at the appropriate rate (30% statutory, treaty rates as applicable).
4. **Report** on Form 1042 to the IRS by March 15 of the following year.
5. **Provide** Form 1042-S to the foreign counterparty.

A miscoded swap (e.g., classified as non-871(m) when it should be) can result in retroactive tax liability years later. Major banks employ tax-compliance teams of 50-100 people specifically for 871(m) work. The 2017 final regulations greatly expanded the compliance load.

11.9 11.9 The Archegos Blow-Up — Week-By-Week

The Archegos collapse in March 2021 is the most consequential prime-brokerage failure since Lehman. The full story is best understood as a sequence of events spread over a single week, with consequences that played out for two more years.

Background (2012-2020). Bill Hwang founded Archegos Capital Management as a “family office” in 2012, after his prior fund (Tiger Asia) was forced to close due to insider-trading charges. Archegos was structurally different from a regular hedge fund: as a family office, it had no external investors, no quarterly reporting requirements, no Form ADV disclosures, and no Form PF reporting (until 2022). The lack of external investors was, in Hwang’s mind, a feature — he could run more concentrated and more leveraged positions than a hedge fund typically would.

By early 2021, Archegos’s strategy had concentrated into a small number of large positions: - ViacomCBS (VIAC, later renamed Paramount): largest single position, \$20-30bn long synthetic exposure - Discovery Communications (DISCA): \$10-15bn long - Baidu (BIDU): \$10-15bn long - Tencent Music Entertainment (TME): \$5-8bn long - Vipshop (VIPS): \$3-5bn long - Several smaller US tech and Chinese ADR positions

Total position size: \$50-100bn long synthetic exposure, on a capital base of approximately \$10-15bn. Leverage: 5-10x, depending on the prime broker.

The structure. Archegos accessed its leverage through equity total return swaps with multiple prime brokers: Credit Suisse, Nomura, Morgan Stanley, Goldman Sachs, UBS, Mitsubishi UFJ. Each PB extended margin assuming Archegos was levered with them at modest size. None of the PBs had visibility into the others’ exposures.

The swap structures gave Archegos several advantages: - Synthetic ownership without 13D filing (no public disclosure of position size). - Tax-efficient delivery of dividends (subject to 871(m) but with treaty benefits). - Initial margin of 7-10% on each PB (compared to 25-50% if positions were held outright).

Friday March 19, 2021 — the trigger. ViacomCBS announced a \$3bn equity offering at a discounted price. ViacomCBS opened down 10% on Monday March 22 and continued falling.

Monday March 22, 2021. ViacomCBS dropped from \$96 to \$86 (-10%). Archegos was long \$20bn+ of VIAC; the day’s loss alone was \$2bn+. Margin calls began at the prime brokers.

Archegos could meet the calls on Day 1 because the position had built up substantial unrealized gains over the prior six months (VIAC had run from \$30 to \$100 in early 2021 before the announcement). The PBs received variation margin and continued.

Tuesday March 23, 2021. VIAC continued lower, closing at \$69 (-19% from Monday’s open). Archegos’s VIAC loss alone was now \$5-7bn cumulatively. Other Archegos positions (DISCA in particular) were also falling.

Margin calls accelerated. Archegos’s capital was being drawn down rapidly. Hwang reportedly told his prime brokers that he could meet the calls but was running out of liquidity.

Wednesday March 24, 2021. VIAC at \$52, down another 25% from Tuesday. DISCA also dropping sharply. Total Archegos paper losses now exceeded \$10bn. The fund’s \$10-15bn capital was approximately wiped out.

Some PBs began considering forced liquidation. Internally at major PBs, risk teams ran scenarios: if Archegos defaults, what is the cost of liquidating each position?

The market dimension: VIAC and DISCA were each among the most-shortest names in the S&P 500 by Wednesday afternoon. Hedge funds, anticipating that Archegos’s positions would be force-sold, established short positions at \$50-55 in VIAC and \$60-65 in DISCA.

Thursday March 25, 2021 — the inflection. Archegos missed margin calls. The PB risk teams convened.

The critical decision: who liquidates first? Each PB could see the writing on the wall. Goldman Sachs and Morgan Stanley reportedly held Thursday afternoon meetings where they decided to sell their Archegos exposures aggressively the next morning. Credit Suisse, Nomura, and others hesitated, hoping for a coordinated exit.

Friday March 26, 2021 — the exit. At market open, Goldman Sachs began selling its Archegos-related VIAC and DISCA positions. Block trades hit the market. Morgan Stanley followed. The two banks each unloaded \$10-15bn of position over the morning.

VIAC traded down to \$40 (a 60% loss from the Friday-before-prior). DISCA similarly. The block trades, despite being executed by sophisticated banks, hit the market with massive impact. Front-running hedge funds (who had built short positions on Wednesday-Thursday) covered their shorts at peak gains.

Credit Suisse and Nomura, having delayed their exits, now faced selling into a market that was already collapsing. They sold over Friday afternoon and the following Monday, accumulating much larger losses than Goldman and Morgan Stanley.

Monday March 29 to Friday April 2, 2021 — the unwind. Markets stabilized at much lower levels for VIAC and DISCA. The remaining Archegos positions at the slower-exiting PBs were liquidated. Final loss tally:

Credit Suisse:	\$5.5bn (largest loss)
Nomura:	\$2.9bn
Morgan Stanley:	\$1.0bn (managed exit relatively well)
UBS:	\$0.86bn
Goldman Sachs:	\$0.10bn (best execution)
Mitsubishi UFJ:	\$0.30bn
Total industry:	~\$10bn+

The Goldman vs Credit Suisse contrast is striking: both had similar nominal exposures, but Goldman lost \$0.1bn while Credit Suisse lost \$5.5bn. The difference: Friday morning’s aggressive selling by Goldman vs the multi-day delayed exit by Credit Suisse.

The Credit Suisse aftermath. The \$5.5bn loss was a major contributor to Credit Suisse’s later collapse. The bank’s reputation and business model never fully recovered. By March 2023 (almost exactly two years after Archegos), Credit Suisse was acquired by UBS in a government-orchestrated rescue. The Archegos loss was not the sole cause of CS’s demise (other factors: Greensill, US tax authority issues, persistent reputation damage), but it was a major one.

The post-mortem. Multiple investigations followed:

1. *The Credit Suisse Special Committee Report* (July 2021): identified governance failures, risk-management lapses, and a culture that prioritized client revenue over risk discipline. The report is a landmark document in modern bank risk management literature.
2. *SEC enforcement.* Hwang and Archegos’s CFO Patrick Halligan were criminally charged in April 2022. Hwang was convicted in July 2024 and sentenced to 18 years in prison.
3. *Industry response.* The major prime brokers implemented stricter cross-margining policies, concentration limits, and Form PF reporting expansion. The SEC adopted Rule 10c-1 (effective 2024) requiring securities lending transaction reporting — partly in response to the Archegos opacity.

The lessons. The Archegos episode is taught in every modern bank risk-management training. The key lessons:

1. *Multi-PB exposure is a coordination problem.* When a fund is levered across multiple PBs, the total exposure can be invisible to any single PB. The first PB to liquidate wins; the last loses. The institutional response: better information sharing on aggregate hedge fund exposure, though this remains imperfect.
2. *Concentration limits matter more than VAR.* Archegos’s positions were concentrated in a small number of names. A 10% move in any one position could wipe out the fund. VAR-based risk metrics did not capture this concentration well.
3. *Family offices have weaker disclosure.* Form PF expansion (effective 2023) now captures family offices over a certain size. The Archegos-specific gap has been closed, but the principle (insufficient disclosure for highly-levered private investors) remains a concern.
4. *Speed of liquidation is the difference.* Goldman’s decision to sell aggressively on Friday morning, while Credit Suisse hesitated, was the difference between \$0.1bn and \$5.5bn in losses. In default scenarios, decisive action wins.

11.10 11.10 References

- Credit Suisse (2021), “Report of the Special Committee of the Board of Directors Relating to Archegos Capital Management”, July.
- Hwang, S.J. (2024), criminal trial verdict, US v. Hwang.
- Office of Financial Research (2022), “Family Offices and Concentration Risk”, working paper.
- Demarzo, P.M. and Duffie, D. (1995), “Corporate Incentives for Hedging and Hedge Accounting”, *Review of Financial Studies* 8(3): 743-771.
- Christoffersen, P., Goyenko, R., Jacobs, K., and Karoui, M. (2018), “Illiquidity Premia in the Equity Options Market”, *Review of Financial Studies* 31(3): 811-851.

- US Senate Banking Committee (2021), “Hearing on Risks Posed by Concentrated Single-Name Equity Exposure”, May 26 transcript.
- Davies, A. and Vault, R. (2021), “The Anatomy of a Family-Office Blowup”, *Risk Magazine*, June.

12 Chapter 11.5 — Historical Case Studies of Delta-Desk and Sec-Lending Episodes

This chapter collects six case studies that illustrate the mechanics of the prior chapters in real-world detail. Each is presented as a narrative timeline with the desk-side and counterparty-side P&L attribution.

12.1 11.5.1 The Volkswagen Squeeze (October 2008)

Setup. Porsche, in a complex multi-year accumulation, had built up cash-settled call options and direct stock positions that — combined with Lower Saxony’s existing 20% stake — left only ~6% of Volkswagen ordinary shares freely floating by mid-2008. Hedge funds, unaware of Porsche’s full exposure, had built short positions of approximately 12% of float.

Trigger. On Sunday, October 26, 2008, Porsche announced it controlled 74.1% of VOW (42.6% direct, 31.5% via cash-settled calls). Combined with Lower Saxony’s 20%, only 5.9% remained tradeable — but 12.8% was short. A cover-or-die situation.

Mechanics from the desk side:

Monday Oct 27 open: VOW gaps from €210 to €520 (+148%)
Monday close: €520
Tuesday Oct 28 open: €700 (+35%)
Tuesday intraday: €1,005 (peak - briefly the world's largest company)
Tuesday close: €945

Borrow fees on VOW had been ~50 bps per annum prior. By Monday afternoon, locates were quoted at 30-50% per annum, and by Tuesday the locate market collapsed entirely — no shares were available at any price. Short funds with no recall protection had to cover into a one-way market.

Estimated industry losses: - Hedge fund short losses: €15-30bn over two days. - Major casualties: Glenview Capital (closed several positions), Greenlight (David Einhorn — partial exposure). Smaller European long/short funds blew up entirely. - Prime broker exposure: large but mostly absorbed by client capital. Several PBs took small losses on margin calls that exceeded posted equity.

What the desks did: - Goldman Sachs Equity desk reportedly made €1bn in trading profit during the spike from prior long positions and front-running cover orders. - Several smaller European delta-one desks lost money on tracking error in DAX-related products (VW’s index weighting briefly went to ~30%, distorting the index). - Securities lending desks saw lifetime-record fee income for Q4 2008 from VW alone.

Aftermath: - Porsche’s strategy backfired; the company’s debt load (which had funded the call accumulation) became unsustainable. VW eventually acquired Porsche’s automotive operations. - BaFin (German regulator) investigated for market manipulation. Settlements but no convictions. - New disclosure requirements for cash-settled equity derivatives in the EU (later codified in Transparency Directive II amendments).

The lesson: the locate market is a *physical* constraint, not a financial one. When the float is too small relative to the short interest, no rate clears. Shorts must cover *at any price*.

The Frankfurt clearing offices on the morning of October 28, 2008 became one of the most stressful environments in modern European finance. By 09:00 local time, the locate market for VW was completely seized. Phones at every Frankfurt and London-based PB rang continuously. Risk managers from Greenlight, Glenview, and a dozen smaller European long-short funds were on the phones explaining their forced-cover positions to senior risk officers. At 11:00 local time, a senior trader at one of the major German banks reportedly told his team: “I have never seen a market like this. I have nothing to quote. We just stand here and watch.” The morning did not get better. By the close, fund-of-fund redemptions across European long-short hedge funds had begun cascading; some funds saw 30-40% of AUM redeemed within 30 days of the squeeze.

For the audit teams reviewing the affected banks' 2008 books, the VW episode became its own chapter in the year-end audit. Auditors at PwC, Deloitte, and KPMG spent late November and early December 2008 reviewing the borrow-rate exposures and the operational responses to the recall cascade. Some auditor reports specifically noted the *opacity* of the locate market: there was no single source of truth on shares-available-to-borrow, and the audit had to be done by reconciliation across multiple internal systems. The 11pm boardroom narrative — auditors and bank risk officers reviewing the day's near-misses — became a pattern repeated across 2008's many crisis events.

12.2 11.5.2 The Long-Term Capital Management Equity Hedging (1998)

Setup. LTCM, the famed Greenwich-based hedge fund, had massive positions in dual-listed stocks (Royal Dutch / Shell pair, Unilever NV / PLC) and merger arbitrage. The strategy required short stock against long stock at the inter-twin spread, hedged delta-neutral.

Mechanics from the desk side: - LTCM's prime brokerage relationship with Bear Stearns required massive sec lending to enable the shorts. - Borrow fees on the names involved were typically 25-100 bps — manageable in calm markets. - During the August 1998 Russian default and subsequent flight-to-quality, the dual-listed spreads *widened* (the opposite of LTCM's bet). Margin calls cascaded. - Bear Stearns and other PBs forced position liquidations.

The sec-lending angle: LTCM's positions were so large that the names involved (Royal Dutch, Shell) became *special* due to LTCM's borrow demand alone. When LTCM was forced to cover, the borrow demand collapsed, and the locked-in PB inventory at high rates suddenly had no demand.

Aftermath: - Federal Reserve coordinated a \$3.6bn private rescue. - Major PBs took small direct losses (most exposure was secured); larger losses were on the bank-level investments in LTCM equity (\$300m collectively for the participating banks). - The long-term lesson: even “fully-hedged” books can blow up when the hedge legs require continuous rolling of borrow.

The auditor at LTCM. PwC was LTCM's auditor at the time of the collapse. The post-rescue audit, conducted in late 1998 and early 1999, was an unusual experience: the auditors were essentially performing a forensic review of a fund that was being recapitalized through a private rescue. The auditors had to certify that the fund's positions could be marked to fair value, that the unwinding had been conducted appropriately, and that no fraud had occurred during the lead-up. The audit took six months and involved more than thirty senior auditors at peak.

The internal report identified what later became standard textbook lessons: model risk in the LTCM positions was severely under-priced (Russian default scenarios had been considered “tail” but were not stress-tested at the magnitude they actually occurred); position sizing relative to capital was extreme (LTCM held \$1.25 trillion notional on \$4.7bn capital); and the firm's Greek-letter sensitivity displays did not show the cross-correlation effects that materialized during the August 1998 panic. Subsequent risk-management literature uses the LTCM case as the canonical example of why VAR-based risk metrics are insufficient — they capture average-case risk but not tail-case risk. The 11pm boardroom narrative around LTCM's resolution shaped a generation of bank risk officers' thinking about position sizing.

12.3 11.5.3 The Bear Stearns Hedge Funds Collapse (June 2007)

Setup. Bear Stearns High-Grade Structured Credit Fund and Enhanced Leverage Fund were highly-leveraged CDO funds. As subprime started cracking, the funds couldn't meet margin calls.

Why this is in a delta-one guide: The funds' counterparty exposure to PBs was the leading edge of the 2008 sec-lending freeze. Once Bear Stearns itself was wobbling (March 2008), the entire prime brokerage industry recognized the *recall risk on Bear's lending pools* — i.e., funds borrowing from Bear's matched-book had the embedded risk that Bear could fail and recall everything.

Mechanics: - Pre-crisis Bear had ~\$140bn in PB balances. - In the week of Bear's collapse (March 10-17, 2008), \$80bn migrated *out* of Bear's PB platform to Goldman, Morgan Stanley, JPM. - The migration overwhelmed legal and operational capacity. - Net effect: temporary 50-100 bps spike in cross-PB borrow rates as positions were re-established at non-Bear PBs.

The sec-lending franchise at Goldman Sachs and Morgan Stanley grew ~30% in 2008 directly from the Bear migration — a windfall for the survivors.

12.4 11.5.4 The Flash Crash (May 6, 2010)

Setup. A retail mutual fund (Waddell & Reed) submitted a large E-mini S&P sell order via algorithmic execution. The algorithm cleared the bid stack faster than HFT could re-fill.

Delta-desk angle: - ETF arbitrage (Chapter 9) broke down for ~20 minutes. SPY traded as low as \$80 (a 9% intraday discount to NAV at that moment). - Some individual stocks traded at \$0.01 due to stub quotes. - APs could not arbitrage because the bid stacks were absent and the basket-side hedges were broken.

Recovery: Within 30 minutes, prices snapped back. Most trades that hit the stub-quote were canceled by the exchanges as “clearly erroneous.” A few hundred million dollars in transactions were left standing — some retail investors lost or won material amounts.

Aftermath: - Single-stock circuit breakers introduced. - Limit-up/limit-down rules implemented. - HFT was widely blamed; analysis later showed HFT was *recovering* the market, not crashing it. The original Waddell & Reed order was the trigger.

The lesson for delta-desks: even fully-functioning AP arbitrage can break for short periods when bid stacks evaporate. ETF NAV discounts can be misleading during stress.

12.5 11.5.5 The 2014 Russian Stock Collapse and the Sberbank Borrow Spike

Setup. In response to Russia’s annexation of Crimea, US and EU sanctions cascaded through 2014. Russian financials (SBER, VTBR) traded at 4-5x P/E by year-end.

Sec-lending angle: - Hedge funds piled in short on Sberbank ADRs and London-listed GDRs. - Borrow rates spiked from 50 bps to 25-40% per annum. - Russian equity inventory in Western lending pools was withdrawn as some institutional holders pulled out.

Mechanics: - A typical short Sberbank GDR position required term borrow at 25-30% per annum to maintain. - Cost of carry exceeded most expected gains — the trade was only profitable for funds with ultra-short timeframes. - Some funds resorted to synthetic shorts via equity swaps (despite the cost), which prime brokers would then warehouse and hedge with what borrow they could find.

Aftermath: - SBER and VTBR rebounded modestly in 2015-2016 before collapsing again in 2022. - The 2014 event was a small-scale rehearsal for the 2022 Russian invasion. Many Western banks had re-built Russian exposure between 2017-2022 and then took massive write-offs in 2022.

12.6 11.5.6 The 2024 Yen Carry Trade Unwind (August 2024)

Setup. Years of zero-rate Japanese policy had created a massive yen carry trade — borrow JPY at near-zero, invest in higher-yielding assets globally. Hedge funds, asset managers, and individual Japanese investors held trillions in this trade.

Trigger. The Bank of Japan’s July 31, 2024 rate hike (to 0.25%, the highest since 2008) and Friday’s weak US jobs report combined to create a perfect storm. The yen rallied violently from 150 to 142 in days.

Delta-desk angle: - Carry-trade unwind required selling US and European equities to repay yen borrowings. - Major hedge funds blew through stops; massive volume on August 5. - The Nikkei dropped 12% in a single day — the largest one-day decline since 1987. - US and European equity sec-lending demand spiked as some funds shorted into the panic.

Mechanics: - Multi-strategy funds reported losses of 5-15% over the week. - Several large multi-strategy funds (Millennium, Citadel) reportedly cut sizing aggressively to limit further damage. - Equity volatility surface re-priced violently — VIX spiked to 65 intraday on Aug 5, the highest since March 2020.

Sec-lending observations: - Names with high carry-trade-fund ownership saw spike in borrow demand. - Specials list expanded by ~30% in the days after the unwind. - Major bank PB books reportedly earned record fees in Aug 2024 from the temporarily expanded specials universe.

Recovery: - Markets recovered most of the losses within 2-3 weeks. - Bank of Japan softened its hawkish tone, calming the carry-trade dynamics. - Many funds re-established carry trades by late September 2024.

The auditor at 11pm. August 5, 2024 was the first major real-time stress event after Archegos. Risk teams across the major banks worked through the night reconciling positions, computing stress-test outputs, and writing memos to senior leadership. At several banks, the CFO joined a 23:00 ET conference call for a status update. At one Tier 1 bank, a senior risk officer reportedly told her CRO at midnight: “I’ve reviewed every client’s exposure and we are within limits. The losses on hedge funds are real but absorbable. The lending revenue uplift will more than offset the counterparty losses.” The boardroom-style narrative texture of the event was that the system had been tested, had been found capable, and had passed. Compare to 2008: the system had been tested, had been found wanting, and had broken. The two decades between Lehman and the August 2024 event saw the build-out of the post-2008 risk machinery, and the system passed its first major test.

12.7 11.5.7 The Bear Stearns Hedge Funds Collapse — The Auditor Narrative (June 2007)

To complete the case-study set with one more “11pm at the boardroom” example: the Bear Stearns hedge fund collapse in June 2007 was the leading-edge tremor of the 2008 crisis, and the auditor narrative around it was instructive about what was about to come.

In June 2007, the Bear Stearns High-Grade Structured Credit Fund and its leveraged sister fund were unable to meet margin calls. The funds’ positions in subprime CDOs had collapsed in value. The funds’ prime broker (Bear’s parent) extended credit to keep them alive temporarily, but by mid-July the situation was untenable.

The auditor narrative for Bear’s Q2 2007 audit was particularly tense. Deloitte (Bear’s auditor) had to certify the firm’s financial statements with the hedge fund subsidiaries’ losses absorbed. The valuation of subprime CDOs in June 2007 was contentious — different pricing services gave different marks, and Bear’s internal models gave still different marks. The auditor had to choose between accepting Bear’s models, accepting third-party marks, or requesting a third-party valuation review.

After several late-night meetings between Bear’s CFO and Deloitte’s audit partners, the agreed-upon valuation was a compromise that recognized about \$1.6bn of writedowns on the hedge fund-related positions. This was the first material recognition of subprime losses by a major Wall Street firm. The disclosure in the Q2 10-Q sent the firm’s stock down 30% over the following month.

The auditor’s letter to Bear’s audit committee — circulated internally — reportedly identified “material weaknesses in internal controls over the valuation of structured credit products.” This was, in retrospect, a clear warning that the firm should have heeded. By March 2008, Bear was the subject of an emergency rescue by the Fed and a forced sale to JPMorgan.

The lesson from the auditor narrative: when valuations are contentious and the firm’s internal models diverge from market prices, the audit becomes a leading indicator of trouble. The retail trader observing financial-press coverage of an audit dispute should treat it as a signal of potential balance-sheet stress.

The persistent feature: the yen carry unwind triggered massive global cross-asset volatility, but the underlying borrow market handled it — supply expanded as some long holders exited, demand expanded from new shorts. The price (rate) found a new equilibrium within days.

The August 5 trading day became, retrospectively, one of the most-studied liquidity events of the post-2008 era. Risk teams across the major banks reportedly spent the following weeks reviewing their hedge effectiveness, their VIX-spike scenarios, and their client-by-client exposure reports. Several banks discovered (in the post-mortem) that they had had higher concentrations of carry-trade-exposed clients than their top-of-book risk dashboards had shown. The fix: better cross-margining views and more granular client-level scenario analysis.

The auditor narrative for August 2024 also had its own chapter. The Q3 2024 earnings season saw multiple bank executives explicitly note the August events in their analyst calls. JPMorgan’s Daniel Pinto referenced the “tactical opportunities” the bank had captured during the unwind. Goldman’s David Solomon noted the strong sec-lending revenue. Citi’s Mark Mason emphasized the operational performance of the firm’s risk infrastructure during the spike. The recurring theme across all the calls: the major US banks had performed well; the dislocations had been profitable rather than damaging; the post-2008 risk discipline had held.

13 Chapter 12 — Risk Management on a Delta One Desk

13.1 12.1 The Core Risks

A delta-one desk's risks are conceptually narrow but operationally complex:

Risk Type	Source	Typical magnitude
Tracking error	Index hedge imperfect	5-20 bps daily
Borrow recall	Lender pulls the loan	Stress event-driven
Reinvestment	Cash collateral investments	Tail event
Counterparty	Swap or repo counterparty defaults	Event-driven
Operational	Failed settlement, mis-booked trade	Daily small losses
Margin	Client cannot post initial/variation margin	Crisis-event
Liquidity	Cannot exit positions in stress	March 2020-style
Tax	Withholding regime change, 871(m) audit	Long-tailed

13.2 12.2 The Central Risk Book

Modern Tier 1 desks operate a *central risk book* — a unified view across all desks (cash equity, derivatives, prime, financing) of the bank's net exposure. Why:

- Delta-one warehouses massive notional. A single hedge fund client's swap may be \$10bn.
- The cash equity desk may also be holding inventory in the same names.
- The derivatives desk may have offsetting Greek exposure.
- Without aggregation, the bank can be unintentionally net long or short by tens of billions across the same names.

The central risk book reconciles all positions hourly (or minute-by-minute at top firms). It identifies natural offsets across desks and routes hedge orders to the most efficient internal source.

13.3 12.3 Stress Testing

Standard stress scenarios:

- **−30% equity, +10 vol points.** A 2008-style crash. Tests basket hedge break, prime broker counterparty losses, margin shortfalls on swap clients.
- **+10% market, +5% USD.** A risk-on scenario. Tests EM swap losses, FX hedging.
- **Single name −50%.** A specific corporate event. Tests concentration limits.
- **HTB rate spike to 100%.** A meme-stock-style event. Tests recall capacity.

Each scenario produces a P&L vector across all books. Desk and risk leadership review weekly.

13.4 12.4 Margin Calls and the 24-Hour Cycle

For a swap client, margin is calculated daily:

Day t close:

Mark client positions to market.

Variation margin = Net P&L on the swap basket.

If client owes: VM call due by Day $t+1$ morning.

Day $t+1$:

Initial margin reviewed: if move was severe, IM may be increased.
Concentration limits checked.

A client failing to meet variation margin triggers default proceedings. The PB has 1-3 business days (per ISDA-CSA terms) to begin liquidation. Speed of liquidation is the difference between Goldman Sachs (lost \$0.1bn on Archegos) and Credit Suisse (lost \$5.5bn).

13.5 12.5 The CSA Hierarchy

Cross-currency Swap Agreement (CSA) terms define collateral:

- **Threshold:** the unsecured exposure permitted before collateral is required.
- **Minimum transfer amount (MTA):** the smallest collateral movement.
- **Independent Amount (IA):** initial margin posted upfront.
- **Eligible collateral:** what can be posted (usually cash, T-bills, sometimes equities).

For a top-tier hedge fund client, the CSA may have: - Threshold: \$50m (no collateral until exposure exceeds this). - IA: \$0 to \$200m (varies by client riskiness). - MTA: \$1m.

For a weaker client (or post-Archegos): - Threshold: \$0. - IA: 10-20% of initial position size. - MTA: \$250k.

The CSA is the legal structure preventing Archegos-style accumulation of unsecured exposure.

13.6 12.6 The Day-of-Failure Playbook

When a hedge fund counterparty defaults, the prime broker's playbook is well-rehearsed. Step-by-step:

T+0 morning:

- Fund fails morning margin call.
- PB's risk team alerts senior management.
- Standstill agreement attempted (negotiate 24-48 hour pause).
- Legal team activates ISDA-CSA default language.

T+0 afternoon:

- If standstill fails, declare event of default.
- Notify other PBs (industry convention; not always followed).
- Begin liquidation planning.

T+1 morning:

- Begin liquidation of derivatives positions. Start with the most liquid, most concentrated names.
- Use TWAP/VWAP algorithms to minimize price impact.
- Coordinate internally to avoid intra-firm front-running.

T+2 to T+5:

- Continue orderly liquidation.
- Liquidate stock positions. Use dark pools where possible.
- Mark losses to P&L; declare to auditors.

T+5 onward:

- Final position closure.
- Total loss assessment.
- Internal post-mortem and process review.

The Goldman Sachs vs Credit Suisse Archegos divergence comes down to T+1 timing. Goldman's risk team began executing aggressively on T+0 afternoon; Credit Suisse delayed in hopes of a negotiated outcome through T+3. The price of those 2-3 days was billions in incremental losses.

13.7 12.7 The CCAR / DFAST Stress Tests

Annual Federal Reserve stress tests (CCAR — Comprehensive Capital Analysis and Review, and DFAST — Dodd-Frank Act Stress Tests) include severe downside scenarios for major bank delta-one books.

The 2024 severe scenario included: - 55% equity decline. - 180bp credit-spread widening. - 30% commercial real estate decline. - High volatility regime.

For a major delta-one desk with \$50bn long inventory, this scenario implies: - Direct equity P&L: ~\$27bn on the long inventory. - Hedging response: hedges should offset most of this; residual ~\$2-5bn. - Counterparty losses on swap clients: estimated \$1-3bn. - Lending book impact: minor (collateralized).

Major banks must demonstrate they can absorb these losses while maintaining required Tier 1 capital ratios. The 2024 results showed all major US banks passing — capital buffers built post-2008 were sufficient.

The Fed's approval/denial of capital distribution plans (dividends, buybacks) hinges on these stress tests. Pre-2017, the test was binary (pass/fail). Post-2017, banks have more flexibility but still face quantitative constraints.

13.8 12.8 The Internal Risk Limits

Beyond regulatory capital, each desk operates under internal risk limits:

Limit Type	Typical Magnitude (Tier 1 desk)
Single-name exposure (long)	\$1-2bn
Single-name exposure (short)	\$500m-1bn
Single-client total exposure	\$5-10bn
Daily VaR (95%)	\$25-50m
Daily P&L stop-loss	-\$100m
Sector concentration	<30% any one GICS sector
Country concentration	<15% any one country (ex-US)
HTB exposure	<5% of total inventory in HTB names

A breach of any limit triggers escalation: first to the desk head, then to the CRO, then to the CEO and board for material breaches. The 2024 stress events (August yen unwind, March 2020 COVID) saw multiple breaches reported and quickly remediated.

13.9 12.9 Operational Risk — The Quiet Killer

Operational risk on a delta-one desk is large and underappreciated. Typical incidents:

- **Mis-booked trades.** A trader books a swap as long instead of short. Hedge is wrong; mark-to-market reveals over weeks. Typical loss per incident: \$100k-\$5m.
- **Failed settlements.** Stock not delivered to a buyer. Triggers buy-in costs and reputation damage. Typical: \$50k-\$500k.
- **System outages.** Trading systems down during volatile periods. Lost arbitrage opportunities; risk-management gaps. Typical: \$1-10m per significant outage.
- **Margin calculation errors.** SPAN engine misprices a swap. Client under-margined or over-margined. Often discovered weeks later. Typical: \$1-20m.
- **Tax-treatment errors.** 871(m) miscoded; dividend tax wedge misapplied. Discovered in audit, often years later. Typical: \$5-50m.

Across a major bank's delta-one book, operational losses typically run \$50-200m/year — material, but small compared to regulatory capital requirements and counterparty risks.

13.10 12.10 The Day in the Life of the Central Risk Book

The Central Risk Book (CRB) is a relatively new construct — most major banks adopted CRBs only in the post-2010 period. Before the CRB, individual desks managed their own risk; after the CRB, the bank's aggregate equity exposure across all desks is centrally tracked and sometimes centrally hedged.

A typical day in the CRB at a Tier 1 bank — let's say JPMorgan, where the CRB is part of “Athena” (the firm's integrated trading platform):

06:30 ET (Tokyo close has happened, London just opening). The CRB risk manager arrives. The first task: review overnight P&L attribution. The CRB's positions accumulated overnight from the Asian and London-open trading. The risk manager checks for any unusual moves.

07:30 ET (London is mid-session). The morning risk meeting. Attendees: CRB risk manager, heads of cash equity, equity derivatives, prime brokerage, financing, central trading. The agenda:

1. Aggregate equity exposure across desks: long \$52bn, short \$48bn, net \$4bn long. Within tolerance.
2. Single-name concentration: largest single-name exposure is \$1.4bn long AAPL across all desks. Below the \$2bn limit.
3. Sector concentration: technology at 28% of long, within the 30% sector cap.
4. Country concentration: 14% China exposure (US-listed China ADRs), within the 15% limit.
5. HTB exposure: 4.2% of long inventory in HTB names, below the 5% limit.

Any flag triggers a discussion. Today: nothing flagged. Meeting ends in 25 minutes.

08:30 ET (NY pre-open). The CRB begins its hedge optimization. Looking at the aggregate exposure, the CRB identifies that the equity derivatives desk is short 200 SPX puts (from client market-making) while the cash equity desk is long \$4bn of S&P 500 names (from inventory). The natural net is short S&P puts + long basket, which has approximately \$0 net delta but positive gamma. The CRB confirms this position is within tolerance and routes any new hedge orders accordingly.

09:30 ET (NY open). Major flow events kick in. The CRB monitors aggregate flow as new orders arrive. A hedge fund placing a \$500m long swap on a single name (let's say NVDA) creates a \$500m short delta exposure for the desk; the CRB confirms this hedge is being naturally offset by other desks' inventories before the desk goes to the external market.

11:00 ET. The CRB runs a stress test. Today's scenarios: -10% S&P 500, +5 vol points, +25 bps rate shock. The CRB simulates the hedge book and confirms losses would be within risk capital limits. Stress test runs are typically multiple times per day and are part of the standard workflow.

14:00 ET. A market event: a major Fed speaker said something hawkish. SPX is down 1.2% in 20 minutes; VIX up 2 points. The CRB's monitors flash red on multiple positions. The risk manager checks: are the hedges working as expected? Is anyone exceeding their VaR limits? Today's answer: the hedges are working; no breaches.

If a breach occurred: the CRB would call the affected desk head, ask for a position-reduction plan, and escalate to the CRO if the response was inadequate. The 2008 culture of “we'll figure it out” has been replaced with “show me the unwind plan within 30 minutes.”

16:30 ET (after close). End-of-day P&L review. The CRB compiles the day's aggregate equity P&L: +\$12m. Decomposition: cash equity +\$5m, derivatives +\$3m, prime brokerage +\$2m, central hedging +\$2m. The day was within normal range.

18:00 ET. End-of-day reporting: the CRB produces a daily summary memo for senior management. The memo covers: aggregate position, day P&L, risk-limit utilization, and any escalations. The memo flows up the chain: CRO → COO → CFO → CEO.

The CRB's role is, in some sense, the “air traffic control” of a modern bank's equity business. Without it, individual desks could create portfolio-level risks that no one sees. With it, the bank has a single pane of glass on its aggregate exposure and can make more efficient hedge decisions.

The Tokyo / London / NY 24-hour cycle is critical. The CRB at JPM has dedicated risk managers in each region: - Tokyo: covers 18:00-02:00 NY time - London: covers 02:00-08:00 NY time - New York: covers 08:00-18:00 NY time

The handoffs at 02:00 NY (Tokyo to London) and 08:00 NY (London to NY) are particularly important. Each handoff is a phone call where the ending region briefs the incoming region on outstanding flows, open hedges, and any flags. The handoff is also where escalations happen — if Tokyo had a stress event during their session, London inherits the cleanup.

This is the operational backbone of modern bank equity risk management. It is largely invisible to outside observers but is the difference between Goldman's 0.1bn Archegos loss and Credit Suisse's 5.5bn loss. The CRB's discipline is

the meta-risk-management; the desk-level decisions matter, but the CRB’s coordination matters more.

13.11 12.11 References

- Counterparty Risk Management Policy Group (2008), “Containing Systemic Risk: The Road to Reform”.
 - Brunnermeier, M.K. and Pedersen, L.H. (2009), “Market Liquidity and Funding Liquidity”, *Review of Financial Studies* 22(6): 2201-2238.
 - Federal Reserve Board (2024), “Comprehensive Capital Analysis and Review 2024 Results”, available at [federalreserve.gov](https://www.federalreserve.gov).
 - Bank for International Settlements (2017), “Basel III: Finalising Post-Crisis Reforms”, December.
 - McKinsey & Company (2022), “The Future of Banking Risk Management: Central Risk Books”.
-

14 Chapter 13 — Microstructure, Tax, and the Retail Lending Program

14.1 13.1 W-9 vs W-8BEN — The Retail Tax Wall

Every US-listed dividend hits the question: is the holder US for tax purposes?

- **W-9:** US person (citizen, green card, US resident). Dividends subject to standard income tax. No withholding (unless backup withholding).
- **W-8BEN:** Non-US individual. Dividends withheld at 30%, reduced to typically 15% by treaty.
- **W-8BEN-E:** Non-US entity (corporation, trust, etc.). Various rates and treaty positions.

For US-resident retail investors, every US dividend is fully received. For non-US-resident retail (the global retail base of Robinhood, IBKR, etc.), the 15-30% withholding is a permanent friction.

14.2 13.2 Qualified vs Ordinary Dividends

For US holders, dividends bifurcate:

- **Qualified dividends:** taxed at long-term capital gains rate (15% or 20% federal). Requires holding period (60+ days around ex-date).
- **Ordinary dividends:** taxed at ordinary income rate (up to 37% federal).

REITs distribute mostly ordinary. Most equity dividends qualify if the holding period test is met.

The critical point for sec-lending borrowers: **the holding period test is interrupted when stock is on loan**. If a US individual lends their stock and it’s recalled across the ex-date, they may lose qualified status on that quarter’s dividend. This is one reason why beneficial owners (mutual funds, pensions) lend more aggressively than direct retail — they are largely tax-exempt and don’t care about the qualification.

14.3 13.3 IRS Section 1058

Section 1058 of the IRC governs securities loans by US holders. Key provisions:

- A loan that meets the requirements is *not* a taxable sale (no gain/loss recognized).
- Substitute payments (dividend equivalents from the borrower) are taxed as ordinary income, *not* as qualified dividends.
- The lender retains the economic exposure.

The substitute-payment-as-ordinary-income rule is the key tax wedge. A retail individual lending their stock receives substitute payments instead of qualified dividends, paying ordinary income tax (up to 37%) instead of qualified dividend tax (15-20%). The wedge for high-bracket retail: 17% of dividends.

This is why most retail lending programs (Robinhood, Fidelity, Schwab) are *unattractive* for high-bracket holders of high-dividend stocks. The lending fee captured (typically 30-50% of the gross fee) often does not compensate for the tax wedge.

14.4 13.4 Per-Share Lending Fee Math

When a retail account lends 1000 shares at a 5% per annum borrow fee on a \$50 stock:

Daily fee revenue (gross): $\$50 \times 1000 \times 5\% / 360 = \$6.94/\text{day}$
\$208.33/month

Of this:

Broker (Robinhood/Schwab/Fidelity) keeps: \$104-156 (50-75%)
Retail receives: \$52-104 (25-50%)

On retail platforms, the disclosed “earn an extra return” lending program often nets the user 30-50% of the gross fee. The broker captures 50-70%. This contrasts sharply with institutional agency lending, where the beneficial owner captures 56-70% of the gross.

14.5 13.5 The Robinhood Stock Lending Program

Robinhood’s stock lending program (launched 2022):

- Default-on for free-tier users (must opt out).
- Splits 50/50 between Robinhood and the user.
- Collateral held in a Robinhood-affiliated entity.

Worked example: a Robinhood user holds 1000 shares of a special name at a 20% per annum borrow fee.

Annualized gross fee: $\$50 \times 1000 \times 20\% = \$10,000$
Robinhood keeps 50%: \$5,000
User receives 50%: \$5,000

Tax treatment for user:

Substitute payments taxed as ordinary income.
At 37% bracket: \$1,850 in tax.
Net to user: \$3,150.

For a non-special name at 1% per annum:

Gross fee: \$500.
Robinhood: \$250.
User: \$250 gross, \$158 net at 37%.

The economics are favorable for high-fee names but marginal for low-fee names. Most retail users earn \$1-50 per year from the program despite holding tens of thousands of dollars in stock.

14.6 13.6 Schwab and Fidelity Programs

Schwab Stock Lending Program: - Opt-in (must affirmatively enroll). - Splits typically 60/40 to Schwab / user. - Cash collateral held in Schwab’s Cash Management Account at Schwab Bank.

Fidelity Fully Paid Lending Program: - Opt-in. - Splits 60/40 to Fidelity / user. - Available only to certain account types (margin accounts only, not IRAs).

Across all major retail brokers, the user-side share is 25-50% of gross fees, versus the 50-70% institutional split. The roughly 20% wedge is the cost of the retail intermediation — broker risk-bearing, credit transformation, operational handling.

14.7 13.7 The Wash-Sale Rule and Sec Lending

Section 1091 of the IRC prohibits a “wash sale” — selling a security at a loss and repurchasing a “substantially identical” security within 30 days. The disallowed loss is carried into the basis of the replacement.

For sec-lending purposes: - Lending out a stock and continuing to hold the economic exposure does *not* trigger wash sale (the holder still owns the asset in legal terms, even if in loaned form). - Selling at a loss while having simultaneously borrowed the same name (e.g., as part of a hedging arrangement) *can* trigger wash sale issues — the

IRS interprets the borrowed position as a substantially identical position. - Cross-broker is also captured — selling at a loss in Account A while holding the same name in Account B.

Practical implication for retail: a user lending stock through their broker’s program can sell at a loss without wash-sale concern *if* the loan is unwound or the lent shares are returned before the sale. If lent shares remain “out” during the sale-and-rebuy window, the IRS interpretation is unclear and conservative practitioners assume wash-sale applies.

14.8 13.8 The 60-Day Holding Period for Qualified Dividends

For US individuals to receive *qualified dividend* tax treatment (15-20% rate vs ordinary income up to 37%):

Holding period requirement:

Must hold the share for more than 60 days during the 121-day period beginning 60 days before the ex-dividend date.

If the share is *on loan* during the ex-date or the holding-period window, the holding-period clock typically *stops*. Most retail SLP programs warn about this — the substitute payment received in lieu of the dividend is *always* ordinary income, regardless of holding period.

For a typical retail user with \$50k of dividend stocks at 2% yield, the impact: - Annual qualified dividend: \$1,000 at 15% = \$150 tax. - If lent through ex-date: \$1,000 substitute payment at 37% = \$370 tax. - Tax cost of lending: \$220.

Most retail SLP programs return only \$30-100 per year in fees on a \$50k position. The lending program is therefore *negative-EV* for high-bracket retail with significant dividends — a fact rarely surfaced clearly in user-facing communications.

14.9 13.9 The Qualified Dividend and Substitute-Payment Distinction

The legal distinction between a *dividend* and a *substitute payment* is hair-splitting but tax-significant:

- A *dividend* is a corporate distribution paid to the legal owner of record on the record date.
- A *substitute payment* is a contractually-required payment from a stock borrower to the lender, designed to make the lender economically whole.

The IRS has consistently held that substitute payments are ordinary income. The legal reasoning: the lender no longer owns the stock during the loan, so they cannot receive a “dividend” — only a contractual substitute.

A retail user reviewing their 1099 in February will see substitute payments listed under “1099-MISC Other Income” rather than “1099-DIV”. This is a flag that the user has been lending stock and lost qualification.

14.10 13.10 The Substitute Dividend Wedge for Foreign Holders

For non-US holders, the substitute payment treatment is even more complex. Pre-2010, the substitute payment was generally *not* subject to US withholding when paid by a US borrower to a foreign lender — creating the Section 871(m) loophole that drove the dividend-arb trade.

Post-2017 final regulations: - Substitute payments on US securities generally ARE subject to 30% withholding. - Treaty rates apply (typically 15%). - The result: most cross-border dividend-arb wedge is closed.

Specific carve-outs and exemptions still create niche opportunities, but the headline trade — synthetic exposure to US dividends at 0% withholding — is gone.

14.11 13.11 The Retail Brokerage Lending Program Comparison Matrix

A side-by-side comparison of the major US retail brokers’ stock-lending programs as of 2025-2026:

Broker	Auto-enrolled?	User share	Tax treatment surfaced?	Min account
Robinhood	Yes (opt out)	50%	Yes (T-2 disclosure)	\$0
Schwab	No (opt in)	40%	Yes (with sign-up)	\$50k
Fidelity	No (opt in)	40%	Yes	Margin only
IBKR	No (opt in via PRO)	Variable, 50-90%	Yes (rate posted daily)	\$0

Broker	Auto-enrolled?	User share	Tax treatment surfaced?	Min account
TD Ameritrade	Discontinued program 2023	N/A	N/A	N/A
E*TRADE	No (opt in)	30%	Limited disclosure	\$25k
Webull	Yes (opt out)	35%	Limited disclosure	\$0
Tastytrade	No (opt in via FIS Astec)	Variable, 50-70%	Yes	\$0

IBKR’s program (formally “Stock Yield Enhancement Program” or SYEP) is the most generous to users — typically 50-90% of fees flow back. The reason: IBKR runs a *transparent* matched-book that publishes daily rates. The user can see exactly what their share is being lent for.

Robinhood’s 50% split sounds attractive but the underlying “fees received” can be opaque — some Robinhood users have reported their ostensibly held special names earning very low daily rates (suggesting Robinhood may be lending at lower rates than the open market or capturing more of the spread than the official 50% disclosure implies).

14.12 13.12 The Voting-Rights Tradeoff

When a stock is on loan, the *borrower* has voting rights, not the lender. For most retail users this is irrelevant — they don’t vote in proxy elections. For institutional lenders (pensions, endowments, mutual funds), this matters significantly:

- Major proxy contests (e.g., Tesla 2018 board, Disney 2024 board) saw active institutional reluctance to lend stock during the contest period.
- The “Big Three” passive managers (BlackRock, Vanguard, State Street) have written policies requiring stock recall for material proxy votes.
- This recall demand creates short-term spikes in borrow rates around contested votes.

For retail users in SLP programs, the contractual fine print specifies that voting rights pass to the borrower. A retail user who wants to vote on a proxy must withdraw from the lending program (recall their shares) before the record date. Few retail users notice or act on this.

14.13 13.13 Retail Tax Foot-Guns — Specific Examples

The tax mechanics described above are abstract until the retail user gets bitten by them. This subsection walks through three specific scenarios where retail traders have lost real money to tax surprises, drawn from public reports and broker-customer-service summaries.

Foot-gun 1: Section 1058 misclassification. A retail user holds 1,000 shares of a high-dividend REIT (e.g., Realty Income, O) in their Robinhood account. They have enrolled in Robinhood’s stock lending program. The REIT pays \$0.265 per share monthly = \$265 per quarter.

The user’s expected tax: they assume “qualified dividend” treatment at 15% LTCG rate. Expected tax: $\$265 * 0.15 = \$39.75/\text{quarter}$.

Actual outcome (when the position is on loan across the ex-date): - The dividend received is a “substitute payment,” not a dividend. - Substitute payments on REITs are typically classified as ordinary income. - Tax: $\$265 * 0.32$ (32% bracket) = $\$84.80/\text{quarter}$.

The user lost \$45.05 per quarter to the substitute-payment classification. Robinhood’s lending program shared 50% of the gross fee (let’s say total fee was \$50/quarter on this position) with the user, returning \$25 to the user. Net outcome: user gained \$25 in lending fees but lost \$45 in tax = $-\$20$ net per quarter, *minus* the further indirect cost of additional ordinary income increasing the user’s federal-tax-bracket liability.

This is a real loss that happened to many Robinhood users in 2022-2023. Robinhood’s disclosure of this risk was technically present (in fine print) but most users did not understand it. Several class-action lawsuits have followed, with settlement discussions ongoing as of 2025.

Foot-gun 2: NRA WHT misapplication. A retail user is a UK resident holding US dividend stocks through IBKR. The user properly filed a W-8BEN claiming the 15% UK-US treaty rate.

The user’s expected tax: 15% withheld on dividends, \$0.85 per dollar of dividend received.

Actual outcome (when the user briefly traveled to the US for >183 days in a calendar year): - The user’s tax residence shifted to “US person for tax purposes.” - IBKR’s compliance flagged the W-8BEN as potentially invalid. - Backup withholding kicked in at 24% on subsequent dividends. - The user’s actual receipt: \$0.76 per dollar of dividend.

The user lost 9% of their dividends to the backup withholding, on top of any UK tax they would have owed in their home jurisdiction. The cleanup required filing US tax returns (Form 1040NR), claiming a refund, and managing the documentation across multiple years.

Foot-gun 3: Broker reclassification. A retail user is a US person but their broker (in this case, a small online broker) misclassifies their account as foreign due to a clerical error during a corporate event (e.g., a SPAC merger that triggered a re-paperwork process).

The user’s expected tax: standard US tax treatment — qualified dividends at 15-20%, ordinary income on substitute payments at marginal rate.

Actual outcome (after misclassification): - The broker began withholding 30% NRA WHT on dividends. - The user did not notice for 3-4 quarters. - By the time they noticed, \$4,500 had been withheld from a \$300k portfolio’s dividend stream.

The cleanup: the user filed a corrected W-9, requested a refund of the over-withheld amount, and waited 6-9 months for the broker to process. They eventually recovered most of the over-withheld amount, but with significant delay and frustration.

The lessons. Tax surprises in retail brokerage are common and small (a few hundred to a few thousand dollars) but cumulative. The categories:

- 1. *Substitute payment vs dividend*: most relevant for high-dividend names held in lending programs.
- 2. *Tax residence shifts*: relevant for users with international travel or moves.
- 3. *Broker reclassification*: relevant for users in less established brokerages or after corporate events.
- 4. *Wash sale rules*: relevant when lending stock around losses.
- 5. *Section 1256 mark-to-market*: relevant for large SPX/futures positions at year-end.

The retail user who is doing significant volume should have an annual conversation with a tax advisor familiar with brokerage operations. This is not an institutional-only need; small-account retail can lose hundreds of dollars per year to these mechanics.

14.14 13.14 References

- IRS Publication 550, Investment Income and Expenses.
- Internal Revenue Code Section 1058 — “Transfers of Securities Under Certain Agreements”.
- Internal Revenue Code Section 1091 — “Loss From Wash Sales of Stock or Securities”.
- Treasury Regulations Section 1.871-15 — “Section 871(m) Transactions”.
- Aggarwal, R., Saffi, P.A.C., and Sturgess, J. (2015), “The Role of Institutional Investors in Voting: Evidence from the Securities Lending Market”, *Journal of Finance* 70(5): 2309-2346.
- IRS Form W-8BEN Instructions (2024).
- Treasury Department (2023), “Foreign Account Tax Compliance Act (FATCA) Guidance”.

15 Chapter 14 — Connecting Back to This Repo

15.1 14.1 Strategy Status Summary

The strategies in this repo that intersect with delta-desk and sec-lending economics:

Strategy file	Status	Retail viable?	Connection to chapter
strategies/put_steal.py	Active (with AI filter)	Yes — bull put spread variant	Chapter 5

Strategy file	Status	Retail viable?	Connection to chapter
strategies/vol_arbitrage.py	Active	Yes — defined-risk skew capture	Chapter 4, 10
strategies/dividend_arb.py	Documented (low-edge)	Marginal — sub-10bps edge	Chapter 4

15.2 14.2 Put Steal — Confirmed Math

The implementation in [strategies/put_steal.py](#) computes NII as:

```
NII = X * (1 - exp(-r*T)) - BlackScholesCall(S, X, T, r, sigma)
```

This matches Equation 3 from Barraclough-Whaley (2011). The strategy’s bull put spread entry condition `nii > 0.01` filters for environments where the structural edge is dollar-positive — exactly matching the institutional logic in Chapter 5.

The retail-executable variant (bull put credit spread) does not directly capture NII via short stock arbitrage but inherits the option-pricing wedge created by retail non-exercise. Empirically the strategy’s win rate target of 65-70% is consistent with the structural-edge literature.

Retail viability: Yes, with caveats. The strategy is Robinhood-eligible (defined risk, no naked legs). It requires Robinhood Gold or similar tier (5% interest on cash) for the desk’s NII analog to apply at retail. At a non-Gold account, the strategy still works directionally but loses the structural NII tilt.

Operating notes for the retail user. The strategy is designed to be run continuously across a watchlist of 30-100 names. Best practices for the retail user:

- Run the strategy once per trading day, scanning the watchlist. The signal generation is deterministic given the inputs, so multiple intraday runs are unnecessary.
- Use the strategy’s `iv_max = 0.60` and `vix_max = 40` gates as hard limits. Entries when these are exceeded are higher-variance and have historically produced larger drawdowns.
- The `confidence_thresh = 0.55` is the GBM classifier’s confidence threshold. Lowering it to 0.50 increases trade frequency but reduces win rate. Raising to 0.60 reduces frequency but improves win rate. The default is reasonable; the user should not tune it without doing their own backtesting.
- The strategy uses `position_size_pct = 0.02` per trade. With 50+ active positions, this aggregates to 100%+ of capital. The broker’s portfolio margin will accommodate this, but the user should verify their account size supports the resulting margin requirements.
- Profit-take at 50% of max profit per the default `profit_target_pct = 0.50`. Holding to expiration captures more total profit on average but exposes more positions to expiration-week vol shocks. The 50% rule is a Sharpe-optimal compromise.
- Stop-loss at 2x max loss per `stop_loss_mult = 2.0`. Letting losses run beyond 2x is rarely productive; the position has likely structurally broken.

The retail user adapting these defaults should run their own backtest with the parameter overrides. The strategy’s `backtest()` method accepts overrides for every key parameter; experimentation is encouraged. The expected behavior is that small parameter changes produce small return differences, and large parameter changes (e.g., halving the NII threshold) produce qualitatively different return profiles.

15.3 14.3 Dividend Arb — Why Inactive

The [strategies/dividend_arb.py](#) strategy is documented (see [dividend_arb.md](#)) but flagged as low-edge. The reasoning, expanded with this guide’s framework:

1. **Structural wedge.** The institutional dividend-arb edge (Chapter 4) is largely closed for US securities by Section 871(m) since 2017. The retail variant (capture quarterly SPY dividend with put hedge) has a maximum edge of ~1.5-1.7% annually, ungeared.
2. **Tax disadvantage.** For a retail user, the dividend is qualified (15-20% tax). The hedge cost (put premium) is not deductible against the dividend in any meaningful way. The after-tax edge falls to ~1.0-1.3% annually.

3. **Friction.** The position requires owning the underlying (capital-intensive), buying puts (commission/spread), and managing the timing precisely. Friction eats most of the remaining edge.
4. **Market regime sensitivity.** The trade only works when $VIX < 20$ (puts are cheap enough to offset dividend). This is roughly 60-65% of trading days historically. Outside the window, the trade is negative-EV.

The strategy is therefore retained as documentation (educational value, occasional opportunistic use) but not run systematically. **Retail viability: Marginal.**

15.4 14.4 Volatility Arbitrage and Sec-Lending Dislocation

`strategies/vol_arbitrage.py` captures put-call IV skew on retail-heavy names. The connection to sec-lending:

- Hard-to-borrow names have *put-skewed* options pricing because synthetic shorts (long put, short call) embed the borrow cost.
- The IV skew premium capture trade therefore *profits* from HTB regime when the desk's borrow rate forecast becomes embedded in option pricing.
- High HTB → wider put-call IV skew → more skew premium to capture.

The trade structure (sell the rich put, buy a cheaper put as wing, possibly delta hedge with a call spread) extracts the same option-pricing wedge that the desk extracts via direct short-stock lending. **Retail viability: Yes.**

15.5 14.5 What Is NOT Retail-Executable

For completeness, the following from this guide are *not* retail-executable:

- **Equity total-return swaps** (Chapter 11). Requires ISDA agreement, typically \$10m+ minimum size, accredited-investor and qualified-purchaser status.
- **Direct securities lending with negotiated rates** (Chapter 2). Retail SLP programs (Chapter 13) capture only 25-50% of fees and have tax disadvantages.
- **Index rebalance front-running** (Chapter 8). Requires institutional settlement, large size, and prime-broker financing.
- **ETF creation/redemption** (Chapter 9). Requires AP status with the issuer (typically \$50m+ creation unit minimums).
- **M&A risk arbitrage with leverage** (Chapter 8). Possible at retail scale but uneconomic without prime financing.
- **Convertible bond arbitrage** (Chapter 7). Possible in principle, but bond market access and short borrowing make it impractical for retail.

15.6 14.6 What IS Retail-Executable

The following are explicitly retail-executable:

- **Bull put credit spread on NII-positive names** (Chapter 5, `strategies/put_steal.py`).
- **SPX box financing** at IBKR Pro or Tastytrade (Chapter 6). Requires \$110k+ for portfolio margin to be efficient.
- **IV skew premium capture** on retail-heavy names (Chapter 4, 10, `strategies/vol_arbitrage.py`).
- **Dividend capture with put spread** in low-VIX environments (Chapter 4, `strategies/dividend_arb.py`). Marginal but occasionally profitable.
- **Stock lending program participation** at Robinhood, Schwab, Fidelity (Chapter 13). Pure passive income, low magnitude unless holding specials.
- **Synthetic shorts via options on HTB names** (Chapter 10). When direct shorting is unavailable.

15.7 14.6b A Repository Tour for the Curious Reader

The strategy files referenced in this chapter are part of a larger codebase. For the reader interested in inspecting the implementations and adapting them, a brief map of the relevant directories:

```
strategies/
  put_steal.py          - The bull put credit spread variant of the
                        put-steal trade (Chapter 5).
```

<code>vol_arbitrage.py</code>	- IV skew premium capture using five-leg option structures (Chapter 4, 10).
<code>dividend_arb.py</code>	- Dividend capture with put hedge variant (Chapter 4).
<code>short_squeeze_detector.py</code>	- Pattern-based squeeze identification on HTB names (Chapter 10).
<code>tail_risk_long_put.py</code>	- Long out-of-the-money put strategy as portfolio tail hedge.
<code>earnings_pin_risk.py</code>	- Volatility crush capture around earnings announcements.
<code>opex_max_pain.py</code>	- Options expiration max-pain pinning trade.
<code>news_sentiment.py</code>	- NLP-driven news-flow signal.
<code>hmm_regime.py</code>	- Hidden Markov Model regime classifier.

<code>dash_app/guide_articles/</code>	
<code>delta_desk_sec_lending.md</code>	- This guide.
<code>put_steal.md</code>	- Strategy-specific guide for put-steal.
<code>vol_arbitrage.md</code>	- Strategy-specific guide for vol arb.
<code>dividend_arb.md</code>	- Dividend arb strategy guide.
... (other strategy-specific guides)	

<code>docs/guide/</code>	
<code>pdf/</code>	- Compiled PDF exports.

Each strategy file follows a consistent structure: docstring with thesis and references, helper functions (Black-Scholes pricing, NII computation, etc.), feature/label construction, walk-forward training loop, signal generation, and a backtest function with parameter overrides. The retail user adapting the strategies should start with the docstring and the `generate_signal` method, which together describe the trade's logic.

The connection between this guide and the strategy files is intended to be direct. Where the guide references a paper (e.g., Barraclough-Whaley 2011), the strategy file's docstring cites the same paper. Where the guide describes a parameter regime (e.g., NII threshold), the strategy file uses the same default. The retail user should be able to read the guide, then read the strategy file, and have a unified mental model of what the strategy is doing and why.

15.8 14.7 The Big Picture for the Retail Quant

The institutional delta-desk ecosystem generates roughly \$20-30bn in annual revenue across the major banks. Of this, perhaps 5-10% flows back to retail in some form — primarily through ETF tight tracking, secondarily through securities lending program payouts, and tertiarily through option pricing being efficient enough that retail can construct synthetic positions at near-fair-value.

The retail trader who understands this ecosystem has three advantages over the retail trader who doesn't:

1. **Better pricing intuition.** Knowing that an option's price embeds borrow rates and dividend forecasts allows the user to trade options aware of these dimensions.
2. **Better execution.** Recognizing that dealer markets are made by delta-one desks pricing in financing and lending costs allows the user to time entries when these costs are favorable.
3. **Direct edge.** A small subset of trades (put-steal, IV skew, occasionally SPX boxes) provide a structural alpha that survives retail-grade execution.

The retail trader's job is not to replicate the institutional desk. It is to recognize when the institutional desk is offering a price that contains an exploitable wedge — and to act on those windows with discipline.

15.9 Appendix A — Glossary

Term	Meaning
AP	Authorized Participant — entity entitled to create/redeem ETF shares
Box spread	Four-leg options structure paying K2–K1 at expiry
Caput	Black-Scholes call value used in NII calculation
CSA	Credit Support Annex — collateral terms in ISDA agreements
Cum-ex	German tax fraud scheme exploiting double-claimed dividend credits
Delta one	Linear payoff product (delta = 1.0)
GC	General Collateral — easy-to-borrow names
HTB	Hard-to-borrow
ISDA	International Swaps and Derivatives Association
LCR	Liquidity Coverage Ratio (Basel III)
NII	Net Interest Income (in put-steal context)
NPC	Notional Principal Contract
OCC	Options Clearing Corporation
PCF	Portfolio Composition File (ETF basket spec)
PB	Prime Broker
Rebate	Rate paid by stock lender to short seller on cash collateral
Repo	Repurchase agreement
RWA	Risk-Weighted Assets
SBL / SLB	Securities-Based Lending / Securities Lending Business
SOFR	Secured Overnight Financing Rate
TRS	Total Return Swap
VM	Variation Margin
WI	When-Issued (pre-spin trading)

15.10 Appendix B — Symbols and Notation

S	= Spot price
K, X	= Strike price
T	= Time to expiry (years)
r	= Risk-free rate
	= Volatility (annualized)
D	= Dividend (cash)
NII	= Net Interest Income
C	= Call price
P	= Put price
	= Delta
Γ	= Gamma
θ	= Theta
	= Vega

15.11 Appendix C — Worked Examples Compendium

This appendix gathers the major worked numerical examples from the body of the guide for quick reference. Each example is annotated with key parameters and the chapter where the full discussion appears.

15.11.1 C.1 The \$200m Index Swap (Chapter 1.3)

Notional: \$200m on Russell 2000

Tenor: 6 months

Funding: SOFR + 60 bps

Internal funding cost: SOFR + 25 bps

Bank revenue streams:

Net financing P&L:	\$0.35m
Sec-lending share:	\$0.245m (35 bps blended × 70% × 0.5 yr)
Dividend wedge:	\$0.30m
Execution friction:	\$0.20m

Total:	\$1.10m on \$200m / 6 months = ~110 bps annualized

15.11.2 C.2 The \$1m DRUG Short at 8% Borrow Fee (Chapter 3.2)

Position: short \$1m DRUG, 30 days, borrow fee 8%, SOFR 5.30%

Daily cost: $\$1m \times 8\% / 360 = \$222/\text{day}$ (approximate)

30-day total: \$6,800

PB share (~25%): \$1,700

Agent (~25%): \$1,275

Beneficial owner: \$3,825

15.11.3 C.3 The MEDX Lifecycle (Chapter 2.9)

50,000 MEDX short at \$42, 22.5% borrow fee, 30 days

Day 1-7: Daily cost \$1,313 (\$9,188 cumulative)

Day 8-13: Daily cost \$1,434 (after VM call increased collateral)

Day 14: Forced cover 2,000 shares due to recall (lost \$4k vs original)

Day 14-30: Daily cost \$1,280 (at 24% blended rate)

Total:

Gross trade P&L: +\$88,000

Borrow cost: -\$39,554

Net P&L: +\$48,446

15.11.4 C.4 The MSFT Pre-Ex-Date Term Repo (Chapter 4.11)

Foreign holder: \$200m MSFT, \$0.83 quarterly dividend

30% withholding without trade: net \$0.581/share = \$116k after withholding

With repo trade: net \$0.747/share = \$149k

Improvement to client: \$33,200 per quarter

Bank's spread: ~\$280k per trade (after hedging costs)

15.11.5 C.5 Put-Steal NII Across Rate Regimes (Chapter 5.4)

$X = \$100$, $r = 25\%$, $T = 30$ days

NII per share at $S = \$80$:

$r = 0\%$: \$0.00

$r = 4\%$: \$0.27

$r = 5\%$: \$0.34

$r = 8\%$: \$0.55

NII per share at $S = \$90$ (less ITM):

$r = 5\%$: \$0.20

NII per share at $S = \$98$ (near ATM):

$r = 5\%$: \$0.02

15.11.6 C.6 The Bull Put Credit Spread on META (Chapter 5.6)

META at \$510, August 2024

Sell put \$500, 21 DTE: credit \$5.50
Buy put \$480, 21 DTE: pay \$1.85

Net credit per spread: \$3.65 (\$365)
Max loss: \$1,635
Risk/reward: 22.3% return on capital if META > \$500 at expiry

15.11.7 C.7 The 24-Month SPX Box Loan (Chapter 6.3, 6.13)

Sell SPX June 2028 4500/5000 box (1000-pt width)
Receive today: \$91,476 (at 4.45% implied rate)
Pay at expiry: \$100,000

Cost: \$8,524 over 24 months = 4.45% annualized
After Section 1256 tax (60/40):
Effective rate: ~3.4% in 32% bracket

Comparison to alternatives:
HELOC: 7.5% (5.5% after-tax)
Schwab margin: 10.3%
Personal loan: 9%

15.11.8 C.8 The BIOX Convert Arb (Chapter 7.5)

\$100k face convertible at 4% coupon, conversion 40 shares
Stock at \$20, embedded option 10 vol points cheap
Borrow fee on BIOX: 8% per annum

Annual P&L:
Vol harvest: +\$5,000
Coupon: +\$4,000
Credit improvement: +\$2,000
Borrow cost: -\$2,880

Net: +\$8,120 (8.1% on \$100k)

15.11.9 C.9 The CrowdStrike S&P 500 Add (Chapter 8.7)

CRWD added June 24, 2024
Pre-announcement price: \$385
Effective day close: \$403 (+4.7% over weekend + Monday)

Index fund buying: \$4-6bn at the close
Specialist accumulation profit: \$2-3m on a 100k share book

15.11.10 C.10 The VOO Creation (Chapter 9.3)

VOO trades \$530.50, NAV \$530.45 (1 bp premium)
AP creates one creation unit (50,000 shares):
Cost of basket: \$26,522,500
ETF shares received and sold: \$26,525,000
Gross profit: \$2,500
Net (after frictions): ~\$1,500

15.11.11 C.11 The Equity Total-Return Swap (Chapter 11.6)

\$250m MSFT TRS, 18 months
Funding: SOFR + 65 bps

Initial margin: 15% (\$37.5m)

Bank revenue:

Financing spread: \$1,625k/year
Sec lending: \$300k/year (12 bps)
Dividend wedge: \$200-400k/year
Total: ~\$2.0-2.3m/year

Fund cost:

\$14.875m financing/year (vs \$25m if cash purchase)
Net savings: \$10m/year + 13D anonymity

15.11.12 C.12 The Death-Spiral BIOXYZ (Chapter 7.7)

\$1m bond holder over 18 months in death-spiral structure
Multiple conversions as stock falls \$5 → \$4 → \$3 → \$2 → \$1 → \$0.50
Cash recovered: ~\$900k
Plus coupons: \$150k
Plus delta-hedge gains: \$300k
Plus borrow income: \$50k
Total: \$1.35m on \$1m = 35% over 18 months

15.12 Appendix D — Quick-Reference Formulas

PUT-CALL PARITY (NON-DIVIDEND)

$$C - P = S - K \times e^{-rT}$$

PUT-CALL PARITY (WITH DIVIDEND)

$$C - P = S \times e^{-qT} - K \times e^{-rT}$$

where q = continuous dividend yield

NET INTEREST INCOME (PUT-STEAL)

$$NII = X \times (1 - e^{-rT}) - C(S, X, T, r,)$$

BOX SPREAD VALUE

$$\text{Box} = (K2 - K1) \times e^{-rT}$$

BOX-IMPLIED FINANCING RATE

$$r = -(1/T) \times \ln(\text{box_price} / (K2 - K1))$$

BORROW FEE FROM REBATE

$$\text{borrow_fee} = \text{SOFR} - \text{rebate}$$

IMPLIED BORROW FROM PUT-CALL PARITY

$$q = -(1/T) \times \ln((C - P + K \times e^{-rT}) / S)$$

ETF NAV ARBITRAGE THRESHOLD

Profitable AP create: $\text{ETF_price} > \text{NAV} + \text{AP_fee} + \text{transaction_costs}$
Profitable AP redeem: $\text{ETF_price} < \text{NAV} - \text{AP_fee} - \text{transaction_costs}$

SUBSTITUTE PAYMENT TAX WEDGE (US RETAIL)

Substitute payment tax = ordinary income rate (up to 37%)
Qualified dividend tax = LTCG rate (15-20%)
Maximum wedge: $37\% - 15\% = 22\%$ of dividend

CONVERTIBLE BOND DELTA APPROXIMATION

$$\Delta_{\text{convert}} = \Delta_{\text{call at implicit strike}} \times (C_{\text{value}} / \text{Bond price})$$

15.13 Appendix E — Common Trade Comparison Table

Trade	Retail viable?	Edge	Capital	Risk	Notes
Index swap arb	No	N/A	Massive	Low	Bank-only
SPX box (sell/borrow)	Yes (IBKR Pro)	4.5%	\$50-200k	V Low	Best retail
SPX box (buy/lend)	Yes (IBKR Pro)	4.5%	\$50-200k	V Low	Tax-eff
Put-steal credit sprd	Yes	Marg	\$1k-50k	Med-Hi	See repo
Skew premium capture	Yes	Marg	\$1k-50k	Med-Hi	See repo
Dividend capture	Yes (marginal)	Thin	\$50k+	Med	Not sys
Stock lending program	Yes	Tiny	\$0	V Low	Tax neg
Convert arb	No	~8-15%	\$1m+	Med	Bond mkt
ETF AP creation	No	<1bp	\$50m+	V Low	AP-only
Cross-border div arb	No (closed)	N/A	N/A	Reg	871(m) closed
M&A risk arb (lev)	Marginal	~5-10%	\$100k+	Med	Pre-broker
Synthetic short	Yes	Mid	Var	Med	Vs HTB
SX5E div future	No	5-15%	€100k+	Med	Eurex only
Index reconst trade	No (margin)	~3-5%	\$5m+	Low	Front-run
Term borrow lending	No	Var	\$50m+	Low	Lend pool
Heartbeat (AP-side)	No	~1bp	Massive	Low	ETF tax

15.14 Appendix F — A Note on Sources

This guide synthesizes:

- Academic literature in the *Journal of Finance*, *Review of Financial Studies*, *Journal of Financial Economics*, and *Journal of Financial Markets* from 1985 to 2024.
- Industry reports from S&P Global Securities Finance, IHS Markit, EquiLend, and the Risk Management Association.
- Regulatory documents from the SEC, CFTC, IRS Treasury Regulations, and the Federal Reserve Board.
- Public testimony from the 2021 House Financial Services Committee hearings on GameStop.
- The ISDA Master Agreement and Credit Support Annex standard documents.
- Public bank disclosures (10-Ks, prime brokerage business segment reports) for Goldman Sachs, Morgan Stanley, JPMorgan, Credit Suisse, and others.

Specific paper citations are inline with each chapter. For comprehensive industry data, the IHS Markit *Securities Finance Year in Review* and the *Risk Management Association Securities Lending Quarterly Aggregate Report* are the standard references.

15.15 14.8 A Frequently Asked Question — “Why Do The Strategies in This Repo Look Different From Institutional Versions?”

The retail-executable strategies in this repo (put-steal credit spread, skew premium capture, dividend arb with put hedge) are *cousins* of the institutional trades, not direct copies. The differences are dictated by retail constraints:

Constraint	Retail Limitation	Institutional Equivalent
Cash investment yield	~0% non-Gold; ~5% Gold	SOFR ~5.30% always
Short stock	Limited locate, extra fees	Full PB matched-book
Margin efficiency	Reg-T strategy-based	SPAN portfolio margin
Tax — sub payments	Ordinary income	Pension/foreign tax-eff
Tax — Section 1256	Available for SPX only	All swaps qualified

Constraint	Retail Limitation	Institutional Equivalent
Bid/ask	Retail-grade	Tier 1 dealer flow
Counterparty	Broker	OCC/ISDA
Reporting	None pre-trade	Full Form PF, 13F

The strategies file at [strategies/put_steal.py](#) explicitly uses a defined-risk credit spread structure (bull put with wing) rather than the institutional “long-and-short same put” arbitrage. The institutional structure requires unlimited risk on the short leg and continuous monitoring; the retail structure caps risk and can be entered/exited normally.

The strategy file at [strategies/vol_arbitrage.py](#) similarly uses a five-leg structure to capture skew premium with full delta hedging via spread structures, all defined-risk and Robinhood Level 3 compliant.

[strategies/dividend_arb.py](#) implements the put-hedged dividend capture (Chapter 4.5/4.6) — the only retail-accessible variant of the broader dividend trading franchise. As noted in 14.3, the edge is structurally thin and the strategy is documented but not run systematically.

15.16 14.9 What’s Missing — Future Work

Several major delta-desk activities are not implemented as live strategies in this repo, either because they are not retail-viable or because additional infrastructure is needed:

1. **SPX box financing strategy.** A “synthetic margin loan via SPX box” implementation would be retail-accessible (with portfolio margin) but requires brokerage integration and is more of a *cash management* function than a *trading* function. Future work item.
2. **HTB synthetic short.** Implementing synthetic shorts via long-put + short-call on names where direct shorting is unavailable. Useful as a portfolio hedging tool. Future work item.
3. **Index rebalance front-running.** Possible at retail scale for small additions, but requires real-time index-membership change detection. Could be a useful event-study strategy. Future work item.
4. **ETF NAV-discount monitoring.** Real-time monitoring of ETF NAV vs market price for stress-period dislocation. Requires NAV feed integration. Useful as a stress-detection signal. Future work item.
5. **Convertible cheapness scanner.** Scanning the convertible market for “cheap” issues (theoretical value > market) as an indirect borrow-rate forecast. Requires convertible bond data feed. Lower priority — convertible markets are illiquid for retail.
6. **Borrow rate alerts.** Monitoring the broker’s HTB list for daily rate changes and alerting on extreme moves (squeeze candidates). Low-priority; data is broker-specific.

These are documented here for clarity on what the framework is *not* yet capturing. None are critical gaps; they are extensions.

15.17 14.10 The Final Word — How to Read Industry News After This Guide

Having reached this point, the reader is equipped to read any institutional delta-desk or sec-lending news with structural insight. A few examples of how a sophisticated reading should look:

News headline: “Hedge fund shorts surge on biotech XYZ as borrow fees hit 50%.”

Sophisticated reading: A 50% borrow fee implies the locate market is very tight on XYZ. Shorts are paying \$14k/year per \$100k notional just to maintain. Either the directional thesis is very high-conviction, or this is a small sub-position. A squeeze risk is now meaningfully elevated — at this rate, even a modest positive news event could trigger panic covering.

News headline: “Bank reports record Q1 prime brokerage revenue.”

Sophisticated reading: Specials concentration likely high — the meme stock universe or the Russian sanction names. The bank’s hedge-fund-counterparty risk has likely also grown. Watch for any Q2/Q3 default events.

News headline: “ETF X trades at 8% discount to NAV during selloff.”

Sophisticated reading: Authorized participant arbitrage has broken. Either the underlying basket is illiquid (could be EM, junk bonds, etc.), or the AP balance sheet is constrained. The discount may persist for hours or days; opportunistic buying at deep discount is profitable *if* the basket eventually trades at NAV. But this requires basket-level analysis to confirm the discount isn't itself the correct price.

News headline: “Major hedge fund wound down due to margin calls on equity swaps.”

Sophisticated reading: Archegos pattern — possibly multi-PB exposure. Watch for cascading liquidations over 1-3 days. Underlying names with rumored exposure may have outsized flow into PB liquidations. PB earnings will reflect counterparty losses.

News headline: “SEC proposes new rules for securities lending transaction reporting.”

Sophisticated reading: Likely Rule 10c-1 or similar. Will reduce dealer information advantages over time. Specials business may face transparency-driven margin compression. Beneficiaries: hedge funds (better information), retail (eventually).

This kind of structural reading is the prize. Whether or not the trades themselves are executable, the understanding equips the trader to interpret market events at a much higher level than is possible from generic financial press coverage.

15.18 A Final Mental Model — The Retail Quant's Synthesis

After fourteen chapters and a dozen case studies, the retail quant should leave with a synthesized mental model. The model has five elements.

Element 1: The institutional desk is a fee factory, not a directional bettor. Every chapter of this guide reinforces that delta-one and prime-brokerage businesses run on fees and arbitrage residuals, not on directional bets. The desk that “loses money” loses it because of operational failures or counterparty defaults, not because they got the market wrong. The desk's profitability is therefore reasonably stable across market regimes — it scales with assets under management, not with market timing skill.

Element 2: The retail trader is a price-taker on the products the desks make. Every option price retail sees, every ETF NAV retail trades against, every dividend retail receives — these prices are made by delta-one desks pricing in their own financing, sec-lending, and tax economics. Understanding these economics lets retail recognize when the desk's price contains a wedge that retail can profit from.

Element 3: Most desk trades are not retail-executable. The list of trades retail cannot do is long: institutional sec lending, cross-border dividend wedges, equity TRS, ETF AP, M&A risk arb at scale, etc. The retail user accepts this and focuses on the small subset of trades that *are* retail-executable.

Element 4: The retail-executable subset is small but real. Three trades are clearly retail-executable: bull put credit spreads on NII-positive names (Chapter 5), SPX boxes for synthetic financing (Chapter 6), and IV skew premium capture on retail-heavy names (mentioned in Chapter 4 and 10). Each has a structural edge that survives modest retail-grade execution. None is a get-rich-quick scheme; they are 10-25% annualized on disciplined deployment.

Element 5: The reading discipline matters as much as the trading discipline. A retail user who reads industry news with the structural framework of this guide sees three or four levels deeper than a retail user who reads the same news without the framework. The trader who recognizes “this hedge fund is unwinding via PB block trades” knows to watch for cascading liquidations; the trader who reads “hedge fund X reported losses” without the framework misses the cascade entirely.

This synthesis is what the guide aims to leave the reader with. The numbers, the formulas, the case studies, the regulatory references — all of these support the synthesis but do not constitute it. The synthesis is the meta-skill: reading the institutional flow correctly, recognizing when retail can play, and acting with discipline when the conditions align.

15.19 Closing Note

The delta one and securities-lending franchise is, in dollar terms, one of the largest businesses on Wall Street. Most retail traders will never see inside it. This guide has attempted to make it visible: the products, the economics, the math, the regulatory architecture, and the small set of places where retail can profitably participate.

The reader who has reached this point should now be able to read industry coverage of any delta-one or sec-lending event — Archegos, GME 2021, the next inevitable crisis — and recognize the mechanics rather than the headlines. That recognition is the prize.

The retail trader who additionally implements the put-steal, skew-premium, or SPX box trades described herein will find that the structural edges are real, modest in magnitude, and survivable in expectation but not without disciplined risk management. The trades work because the institutional infrastructure that creates them exists, not because of anything special the retail trader does. The trader's job is to recognize the conditions and execute cleanly.

Beyond the retail-executable trades, the value of this guide is the *intuition* it provides for every option price, every ETF spread, every borrow rate, and every dividend cycle. These prices are not random; they are set by the daily operation of a multi-trillion-dollar funding and lending machine. Understanding the machine is the point.